



Introduction to Bambu

Tutorial @ FPL Conference 2017 – Ghent – Belgium

Fabrizio Ferrandi

Politecnico di Milano
Dipartimento di Elettronica, Informazione e Bioingegneria
fabrizio.ferrandi@polimi.it

- ❑ Introduction
- ❑ Bambu HLS flow
- ❑ Synthesis C subset supported by Bambu
- ❑ Interfacing Bambu with GCC
- ❑ Component generation
- ❑ Synthesis of function pointers
- ❑ Integration of hand-written components
- ❑ Extended language support
 - ▶ C++
 - ▶ Fortran
 - ▶ OpenMP simd

- ❑ HLS tool developed at Politecnico di Milano (Italy) within the PandA framework
 - ▶ Available under GPL v3 at
 - <http://panda.dei.polimi.it/>
 - <https://github.com/ferrandi/PandA-bambu>
- ❑ Example features
 - ▶ Front-end Input: interfacing with GCC for parsing C code
 - Complete support for ANSI C (except for recursion)
 - Support for pointers, user-defined data types, built-in C functions, etc..
 - Source code optimizations
 - may alias analysis, dead-code elimination, hoisting, loop optimizations, etc...
 - ▶ Target-aware synthesis
 - Characterization of the technology library based on target device
 - ▶ Verification
 - Integrated testbench generation and simulation
 - automated interaction with Iverilog, Verilator, Xilinx Isim, Xilinx Xsim, Mentor Modelsim
 - ▶ Back-end: Automated interaction with commercial synthesis tools
 - FPGA: Xilinx ISE, Xilinx Vivado, Altera Quartus, Lattice Diamond
 - ASIC: Synopsis Design Compiler

- ❑ Complexity of embedded systems is constantly growing
 - ▶ Designing for FPGAs is mostly an engineering discipline carried out by highly-trained specialists
 - ▶ Difficult for a hardware designer when working at RTL (or lower) level(s)
 - ▶ Moving to higher levels of abstractions is becoming more and more important

- ❑ Example: 1M gates design (study from NEC)
 - ▶ RTL description: ~300k lines of code (avg)
 - ▶ Behavioral description: ~30k lines of code (avg)

- ❑ Writing behavioral code is much “easier”
 - ▶ e.g. C/C++ compared to Verilog/VHDL

- ❑ HLS raises the abstraction level and is going to increase the designers productivity
- ❑ It is starting to be an option for software engineers to accelerate key application kernels
- ❑ Quality of design produced by HLS?
 - ▶ Worse than hand-made RTL design
 - ▶ In several cases better than pure software

Not true anymore!

❑ From LegUp 1.0:

Well, it turns out that none of the existing high-level synthesis tools have source code available for researchers. **GAUT** claims to be open-source but the code is not available for download. **xPilot** from UCLA is an advanced research tool but only the binary is available and it hasn't been updated since 2007. **ROCCC** provides an open source eclipse plugin based on SUIF and LLVM but only supports small C programs. Standard C code must be rewritten to work with ROCCC because all function parameters must be structs. **Trident** uses a very old version of LLVM to interface with an extensive amount of Java code, but unfortunately no longer compiles with the latest version of LLVM.

- ❑ PandA framework development started on 2004 as a support research infrastructure for PoliMi in the context of ICODES – FP6-IST EU-funded project
 - ▶ Parsing and analysis of TLM 2.0 SystemC descriptions (gcc v.3.5)
- ❑ In the hArtes EU-funded project (2006-2010), it was used to
 - ▶ Analyzing generic C-based application annotated with pragmas (OpenMP)
 - ▶ Extracting parallel tasks
 - ▶ Estimating performance of embedded app
 - ▶ C-to-c rewriting
- ❑ Later, in Synaptic (2009-2013) and in Faster (2011-2014) EU-funded projects, logic- and high-level synthesis has been extended
 - ▶ Bambu (HLS tool) was first released in March 2012.
- ❑ ESA funded many research on code predictability analysis, performance analysis, and integration of HLS in model-based design flows.

1. Are we classical computer hackers, skilled in SW development and with 100% control of our PCs and workstations?
 - ▶ Not too surprisingly, that is the perspective from which RMS wrote the GPL.
 - ▶ RMS defined **freedom** as the ability of people like himself to get read-write access to code.
 - ▶ How does this definition map to today's world?
 - ▶ How much transparent are today's systems?
 - HW vs SW
2. Funding of PandA comes mainly from public institutions
3. ... because it's fun

- ❑ Leon Stok, "Data path synthesis", Integration VLSI Journal, 18(1): 1-71 (1994).
- ❑ Jason Cong, Zhiru Zhang, "An efficient and versatile scheduling algorithm based on SDC formulation", DAC 2006, 433-438.
- ❑ Florent de Dinechin, Bogdan Pasca: "Designing Custom Arithmetic Data Paths with FloPoCo". IEEE Design & Test of Computers 28(4), 18-27 (2011).
- ❑ Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, F. Kenneth Zadeck, "Efficiently Computing Static Single Assignment Form and the Control Dependence Graph", ACM Trans. Program. Lang. Syst. 13(4), 451-490 (1991).
- ❑ Preston Briggs, Keith D. Cooper, Timothy J. Harvey, L. Taylor Simpson, "Practical Improvements to the Construction and Destruction of Static Single Assignment Form", Softw., Pract. Exper. 28(8), 859-881 (1998).
- ❑ Benoit Boissinot, Florian Brandner, Alain Darte, Benoît Dupont de Dinechin, Fabrice Rastello, "A Non-iterative Data-Flow Algorithm for Computing Liveness Sets in Strict SSA Programs". APLAS 2011, 137-154.

- ❑ one component per function
 - ▶ function interface
 - ▶ start and done
 - ▶ parameter passing
 - wires
 - memory interaction
- ❑ hierarchy based on call graph
 - ▶ no-recursion
 - ▶ proxy
- ❑ option
 - ▶ `-fwhole-program`

Example: Multiply & Accumulate

11

Fmin=100Mhz
Device=Xilinx Zynq

Constraints

```
int mac(int a, int b, int c)
{
    return a+(b*c);
}
```

C code



Verilog

```
`timescale 1ns / 1ps
module mac(clock, reset, start_port, done_port, a, b, c, return_port);
    // IN
    input clock;
    input reset;
    input start_port;
    input signed [31:0] a;
    input signed [31:0] b;
    input signed [31:0] c;
    // OUT
    output done_port;
    output signed [31:0] return_port;
    // Component and signal declarations
    wire wrenable_reg_0;

    controller_mac Controller_i (.done_port(done_port), .wrenable_reg_0(wrenable_reg_0),
                                .clock(clock), .reset(reset), .start_port(start_port));
    datapath_mac Datapath_i (.return_port(return_port), .clock(clock), .reset(reset),
                             .in_port_a(a), .in_port_b(b), .in_port_c(c), .wrenable_reg_0(wrenable_reg_0));
endmodule
```

Example: (FSM) Controller


12

```
Fmin=100Mhz  
Device=Xilinx Zynq
```

Constraints

```
int mac(int a, int b, int c)  
{  
    return a+(b*c);  
}
```

C code



```
module controller_mac(done_port, wrenable_reg_0,  
                      clock, reset, start_port);  
  
    . . .  
    // IN,OUT signals declaration;  
    // regs/parameters declaration and init  
    . . .  
  
    always @(_present_state or start_port)  
    begin  
        _next_state = S_0;  
        done_port = 1'b0;  
        wrenable_reg_0 = 1'b0;  
        case (_present_state)  
            S_0 :  
                if(~start_port)  
                begin  
                    _next_state = S_0;  
                end  
                else  
                begin  
                    wrenable_reg_0 = 1'b1;  
                    _next_state = S_1;  
                end  
            S_1 :  
                begin  
                    _next_state = S_0;  
                    done_port = 1'b1;  
                end  
            default :  
                begin  
                    done_port = 1'b0;  
                    wrenable_reg_0 = 1'b0;  
                end  
        endcase  
    end  
endmodule
```

controller_mac

Example: Datapath


13

Fmin=100Mhz
Device=Xilinx Zynq

Constraints

```
int mac(int a, int b, int c)
{
    return a+(b*c);
}
```

C code



```
`timescale 1ns / 1ps
module datapath_mac(clock, reset, in_port_a,
in_port_b, in_port_c, return_port, wrenable_reg_0);
// IN
input clock;
input reset;
input signed [31:0] in_port_a;
input signed [31:0] in_port_b;
input signed [31:0] in_port_c;
input wrenable_reg_0;
// OUT
output signed [31:0] return_port;
// Component and signal declarations
wire signed [31:0] out_mult_expr_FU_fu_mac_1448_1472;
wire signed [31:0] out_plus_expr_FU_fu_mac_1448_1473;
wire [31:0] out_reg_0_reg_0;

mult_expr_FU #(BITSIZE(32), PIPE_PARAMETER(0)) fu_mac_1448_1472
(.out1(out_mult_expr_fu_mac_1448_1472),
.clock(clock), .in1(in_port_c), .in2(in_port_b));
plus_expr_FU #(BITSIZE(32)) fu_mac_1448_1473
(.out1(out_plus_expr_fu_mac_1448_1473),
.in1(out_mult_expr_fu_mac_1448_1472),
.in2(in_port_a));

register_SARSE #(BITSIZE_in1(32), BITSIZE_out1(32))
reg_0 (.out1(out_reg_0_reg_0), .clock(clock), .reset(reset),
.in1(out_plus_expr_fu_mac_1448_1473), .wenable(wrenable_reg_0));
// io-signal post fix
assign return_port = out_reg_0_reg_0;

endmodule
```

datapath_mac



Example: Datapath, 200 Mhz

14

Fmin=200Mhz
Device=Xilinx Zynq
Constraints

```
int mac(int a, int b, int c)
{
    return a+(b*c);
}
```

C code



```
`timescale 1ns / 1ps
module datapath_mac(clock, reset, in_port_a, in_port_b,
in_port_c, return_port, wrenable_reg_1, wrenable_reg_0);

// IN
input clock, reset;
input signed [31:0] in_port_a, in_port_b, in_port_c;
input wrenable_reg_1, wrenable_reg_0;
// OUT
output signed [31:0] return_port;
// Component and signal declarations
wire signed [31:0] out_mult_expr_fu_mac_1448_1472;
wire signed [31:0] out_plus_expr_fu_mac_1448_1473;
wire [31:0] out_reg_0;
wire [31:0] out_reg_1;

mult_expr_FU #(.BITSIZE(32), .PIPE_PARAMETER(4)) fu_mac_1448_1472
(.out1(out_mult_expr_fu_mac_1448_1472),
.clock(clock), .in1(in_port_c), .in2(in_port_b));
plus_expr_FU #(.BITSIZE(32)) fu_mac_1448_1473
(.out1(out_plus_expr_fu_mac_1448_1473),
.in1(out_reg_0), .in2(in_port_a));
register_SARSE #(.BITSIZE_in1(32), .BITSIZE_out1(32)) reg_0
(.out1(out_reg_0), .clock(clock), .reset(reset),
.in1(out_mult_expr_FU_fu_mac_1448_1472), .wenable(wrenable_reg_0));
register_SARSE #(.BITSIZE_in1(32), .BITSIZE_out1(32)) reg_1
(.out1(out_reg_1_reg_1), .clock(clock), .reset(reset),
.in1(out_plus_expr_fu_mac_1448_1473), .wenable(wrenable_reg_1));
// io-signal post fix
assign return_port = out_reg_1_reg_1;

endmodule
```

datapath_mac

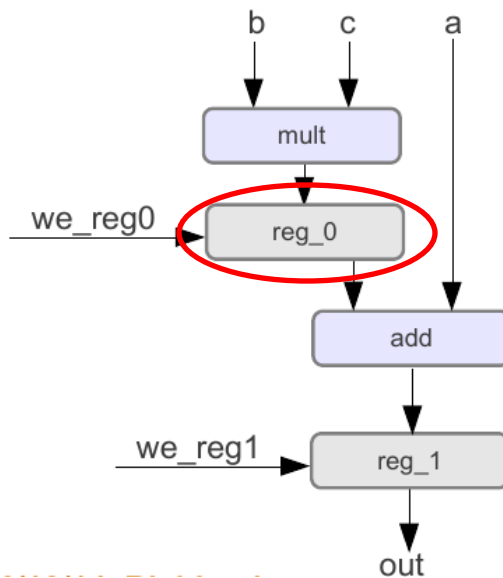

Example: Datapath, 200 Mhz

15

Fmin=200Mhz
Device=Xilinx Zynq
Constraints

```
int mac(int a, int b, int c)
{
    return a+(b*c);
}
```

C code



```
`timescale 1ns / 1ps
module datapath_mac(clock, reset, in_port_a, in_port_b,
    in_port_c, return_port, wrenable_reg_1, wrenable_reg_0);

    // IN
    input clock, reset;
    input signed [31:0] in_port_a, in_port_b, in_port_c;
    input wrenable_reg_1, wrenable_reg_0;
    // OUT
    output signed [31:0] return_port;
    // Component and signal declarations
    wire signed [31:0] out_mult_expr_fu_mac_1448_1472;
    wire signed [31:0] out_plus_expr_fu_mac_1448_1473;
    wire [31:0] out_reg_0;
    wire [31:0] out_reg_1;

    mult_expr_FU #(.BITSIZE(32), .PIPE_PARAMETER(4)) fu_mac_1448_1472
        (.out1(out_mult_expr_fu_mac_1448_1472),
         .clock(clock), .in1(in_port_c), .in2(in_port_b));
    plus_expr_FU #(.BITSIZE(32)) fu_mac_1448_1473
        (.out1(out_plus_expr_fu_mac_1448_1473),
         .in1(out_reg_0), .in2(in_port_a));
    register_SARSE #(.BITSIZE_in1(32), .BITSIZE_out1(32)) reg_0
        (.out1(out_reg_0), .clock(clock), .reset(reset),
         .in1(out_mult_expr_fu_mac_1448_1472), .wenable(wrenable_reg_0));
    register_SARSE #(.BITSIZE_in1(32), .BITSIZE_out1(32)) reg_1
        (.out1(out_reg_1_reg_1), .clock(clock), .reset(reset),
         .in1(out_plus_expr_fu_mac_1448_1473), .wenable(wrenable_reg_1));
    // io-signal post fix
    assign return_port = out_reg_1_reg_1;

endmodule
```

datapath_mac

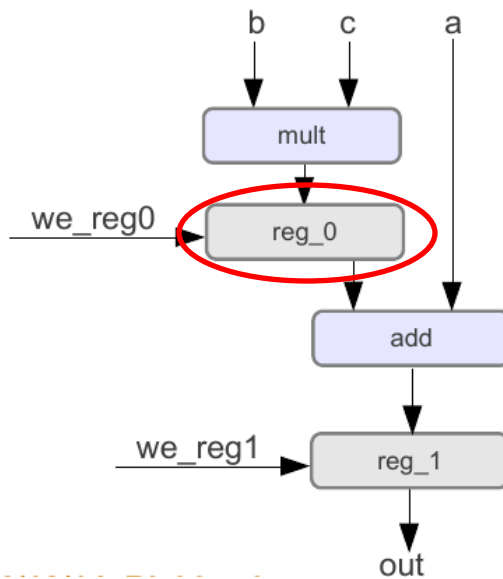

Example: Datapath, 200 Mhz

16

Fmin=200Mhz
Device=Xilinx Zynq
Constraints

```
int mac(int a, int b, int c)
{
    return a+(b*c);
}
```

C code



```
`timescale 1ns / 1ps
module datapath_mac(clock, reset, in_port_a, in_port_b,
    in_port_c, return_port, wrenable_reg_1, wrenable_reg_0);

    // IN
    input clock, reset;
    input signed [31:0] in_port_a, in_port_b, in_port_c;
    input wrenable_reg_1, wrenable_reg_0;
    // OUT
    output signed [31:0] return_port;
    // Component and signal declarations
    wire signed [31:0] out_mult_expr_fu_mac_1448_1472;
    wire signed [31:0] out_plus_expr_fu_mac_1448_1473;
    wire [31:0] out_reg_0;
    wire [31:0] out_reg_1;

    mult_expr_FU #(.BITSIZE(32), .PIPE_PARAMETER(4)) fu_mac_1448_1472
        (.out1(out_mult_expr_fu_mac_1448_1472),
         .clock(clock), .in1(in_port_c), .in2(in_port_b));
    plus_expr_FU #(.BITSIZE(32)) fu_mac_1448_1473
        (.out1(out_plus_expr_fu_mac_1448_1473),
         .in1(out_reg_0), .in2(in_port_a));
    register_SARSE #(.BITSIZE_in1(32), .BITSIZE_out1(32)) reg_0
        (.out1(out_reg_0), .clock(clock), .reset(reset),
         .in1(out_mult_expr_fu_mac_1448_1472), .wenable(wrenable_reg_0));
    register_SARSE #(.BITSIZE_in1(32), .BITSIZE_out1(32)) reg_1
        (.out1(out_reg_1_reg_1), .clock(clock), .reset(reset),
         .in1(out_plus_expr_fu_mac_1448_1473), .wenable(wrenable_reg_1));
    // io-signal post fix
    assign return_port = out_reg_1_reg_1;

endmodule
```

datapath_mac

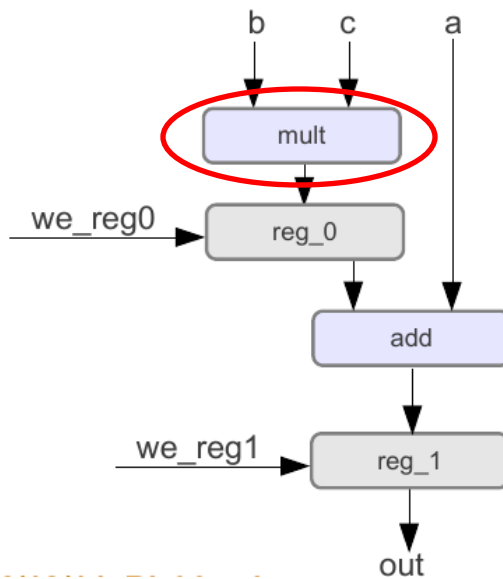

Example: Datapath, 200 Mhz

17

Fmin=200Mhz
Device=Xilinx Zynq
Constraints

```
int mac(int a, int b, int c)
{
    return a+(b*c);
}
```

C code



```
`timescale 1ns / 1ps
module datapath_mac(clock, reset, in_port_a, in_port_b,
in_port_c, return_port, wrenable_reg_1, wrenable_reg_0);

// IN
input clock, reset;
input signed [31:0] in_port_a, in_port_b, in_port_c;
input wrenable_reg_1, wrenable_reg_0;
// OUT
output signed [31:0] return_port;
// Component and signal declarations
wire signed [31:0] out_mult_expr_fu_mac_1448_1472;
wire signed [31:0] out_plus_expr_fu_mac_1448_1473;
wire [31:0] out_reg_0;
wire [31:0] out_reg_1;

mult_expr_FU #(.BITSIZE(32), .PIPE_PARAMETER(4)) fu_mac_1448_1472
(.out1(out_mult_expr_fu_mac_1448_1472),
.clock(clock), .in1(in_port_c), .in2(in_port_b));
plus_expr_FU #(.BITSIZE(32)) fu_mac_1448_1473
(.out1(out_plus_expr_fu_mac_1448_1473),
.in1(out_reg_0), .in2(in_port_a));
register_SARSE #(.BITSIZE_in1(32), .BITSIZE_out1(32)) reg_0
(.out1(out_reg_0), .clock(clock), .reset(reset),
.in1(out_mult_expr_FU_fu_mac_1448_1472), .wenable(wrenable_reg_0));
register_SARSE #(.BITSIZE_in1(32), .BITSIZE_out1(32)) reg_1
(.out1(out_reg_1_reg_1), .clock(clock), .reset(reset),
.in1(out_plus_expr_fu_mac_1448_1473), .wenable(wrenable_reg_1));
// io-signal post fix
assign return_port = out_reg_1_reg_1;

endmodule
```

datapath_mac

Example: (FSM) Controller, 200 Mhz


18

Fmin=200Mhz
Device=Xilinx Zynq

Constraints

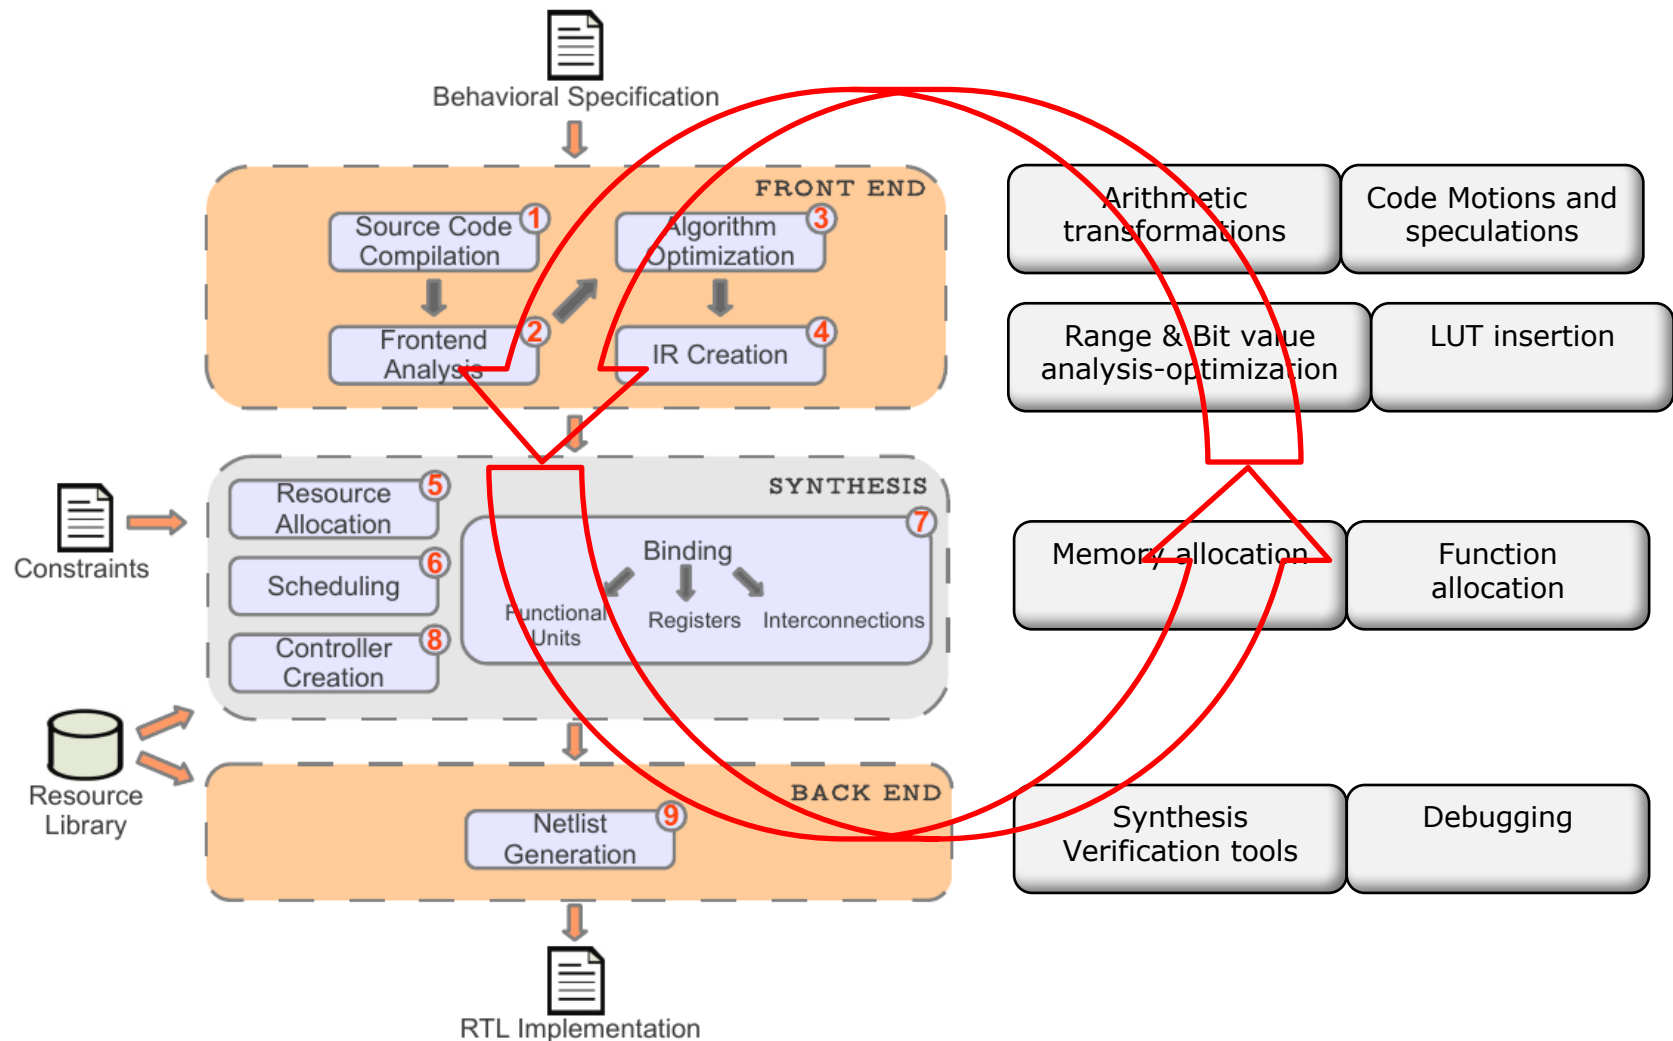
```
int mac(int a, int b, int c)
{
    return a+(b*c);
}
```

C code



```
module controller_mac(done_port, wrenable_reg_0,
    wrenable_reg_1, clock, reset, start_port);
    . . .
    // IN,OUT signals declaration;
    // regs/parameters declaration and init
    . . .
    always @(_present_state or start_port)
    begin
        _next_state = S_0;
        done_port = 1'b0;
        wrenable_reg_1 = 1'b0;
        wrenable_reg_0 = 1'b0;
        case (_present_state)
            S_0 :
                if(~start_port)
                    _next_state = S_0;
                else
                    _next_state = S_1;
            S_1 :
                _next_state = S_2;
            S_2 :
                _next_state = S_3;
            S_3 :
                _next_state = S_4;
            S_4 :
                wrenable_reg_0 = 1'b1;
                _next_state = S_5;
            S_5 :
                wrenable_reg_1 = 1'b1;
                _next_state = S_6;
            S_6 :
                _next_state = S_0;
                done_port = 1'b1;
        end
    endcase
end
endmodule
```

controller_mac



Modular Framework based on the specialization of HLS_step

❑ Minimal command

▶ `$ bambu filename.c`

❑ Controlling the clock period (100Mhz)

▶ `$ bambu filename.c --clock-period=10`

❑ Select the device

▶ `$ bambu filename.c -device-name=xc7z020,-1,clg484,VVD`

- ❑ We support the standard GCC support
- ❑ Supported features:
 - ▶ Expressions of any kind: arithmetic, logical, bitwise, relational, conditional, comma-based expressions.
 - ▶ Types: integers, single- and double-precision floating point, `_Bool` and `_Complex`, struct-or-union, bitfields, enum, typedef, pointers and arrays, type qualifiers.
 - ▶ Variable declarations, initialization, storage-specifiers
 - ▶ Functions definition and declaration, extern or static, pointer to functions, parameters passed by copy or reference, tail recursive functions.
 - ▶ Statements and blocks: labeled (`case`), compound, expression, selection (`if`, `switch`), iteration (`while`, `do`, `for`), jump (`goto`, `continue`, `break`, `return`)
 - ▶ All preprocessor directives
- ▶ Unaligned memory accesses and dynamic pointers resolution
- ▶ GCC vectorization

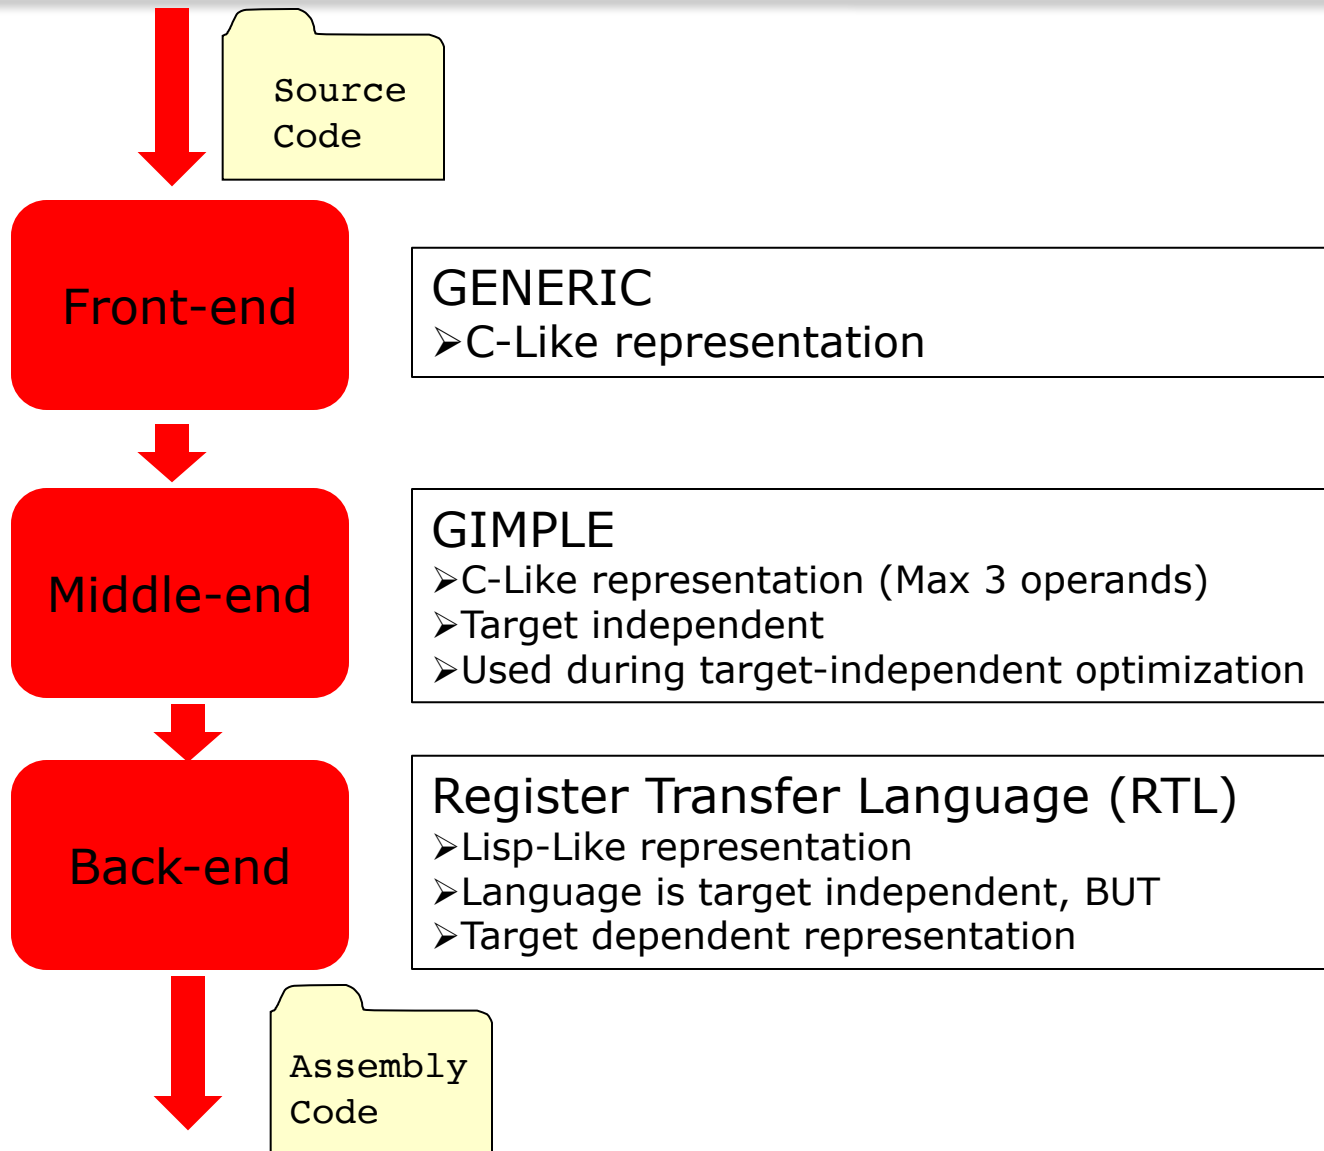
Subset not supported

22

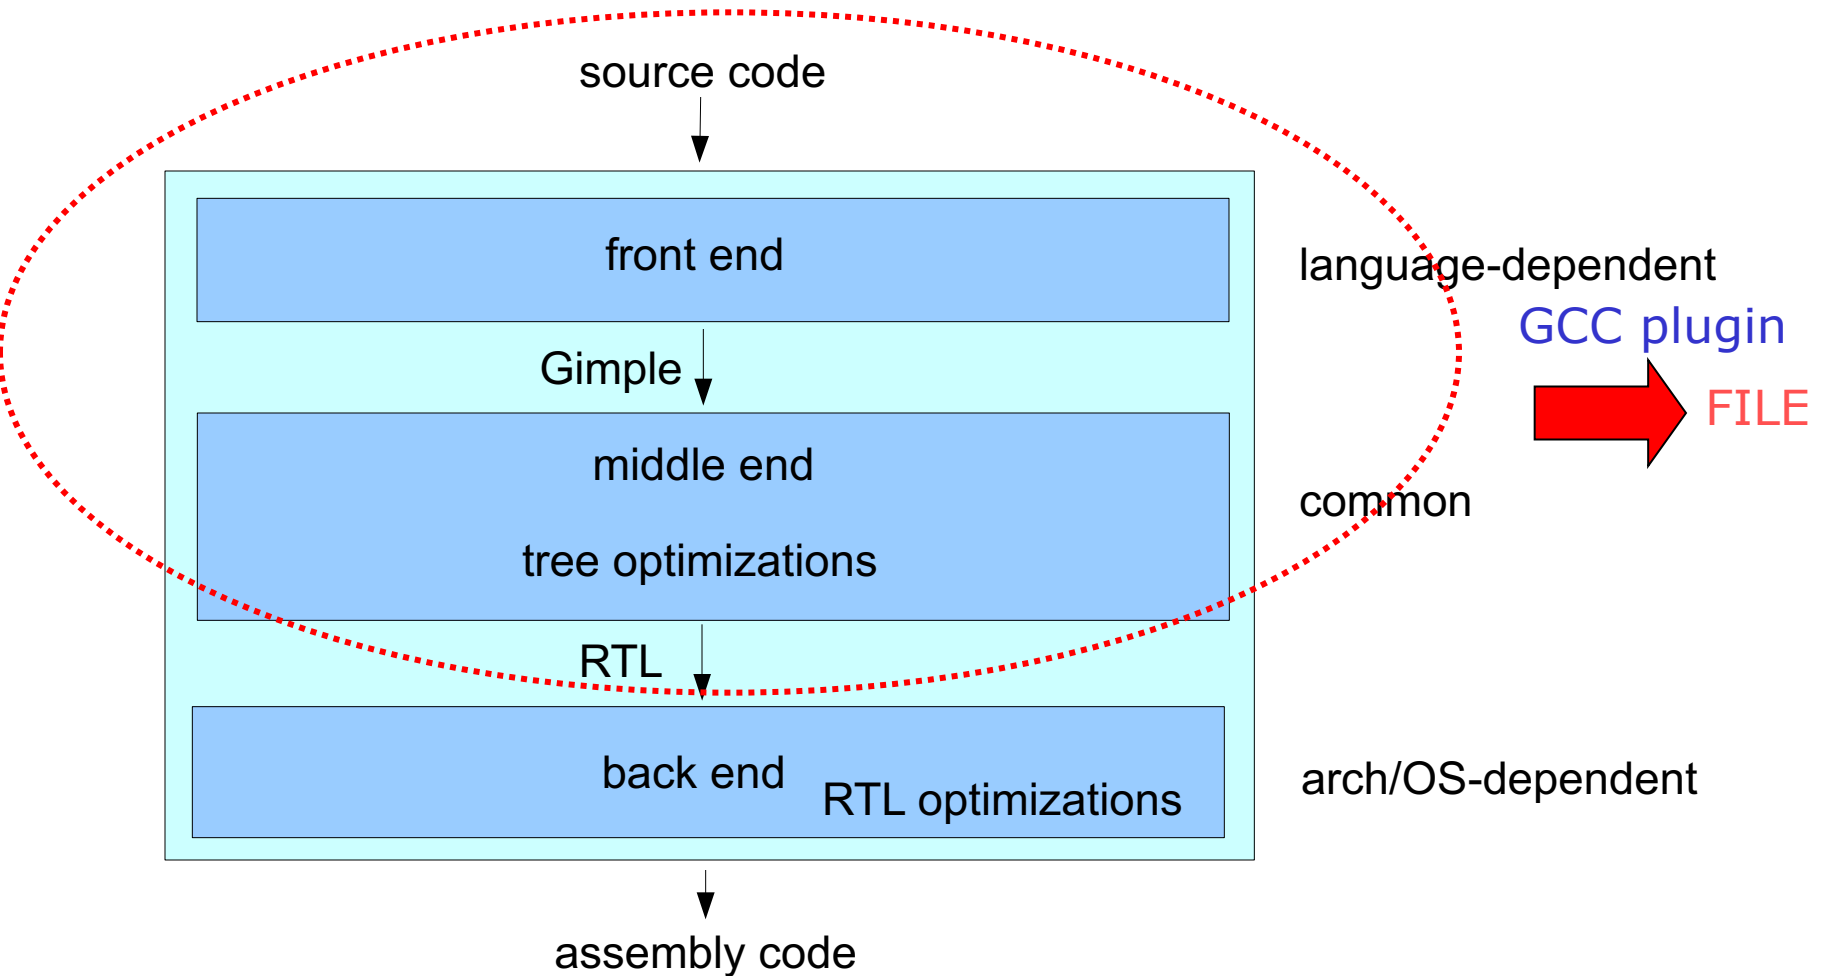
- ❑ struct returned by copy
- ❑ Non-tailing recursive functions

- ❑ `assert`, `puts`, `putchar`, `read`, `open`, `close`, `write`, `printf`, `exit`, `abort`
- ❑ libc functions: `bswap32`, `memcmp`, `memcpy`, `memmove`, `memset`, `malloc`, `free`, `memalign`, `alloca_with_align`, `calloc`, `bcopy`, `bzero`, `memchr`, `mempcpy`, `memrchr`, `rawmemchr`, `stpcpy`, `stpncpy`, `strcasecmp`, `strcasestr`, `strcat`, `strchr`, `strchrnul`, `strcmp`, `strcpy`, `strcspn`, `strdup`, `strlen`, `strncasecmp`, `strncat`, `strncmp`, `strncpy`, `strndup`, `strnlen`, `strpbrk`, `strrchr`, `strsep`, `strspn`, `strstr`, `strtok`
- ❑ libm functions: `acos`, `acosh`, `asin`, `asinh`, `atan`, `atan2`, `atanh`, `cbrt`, `ceil`, `cexp`, `cexpi`, `copysign`, `cos`, `cosh`, `drem`, `erf`, `exp`, `exp10`, `expm1`, `fabs`, `fdim`, `finite`, `floor`, `lffloor`, `fma`, `fmax`, `fmin`, `fmod`, `fpclassify`, `frexp`, `gamma`, `lgamma`, `tgamma`, `hypot`, `ilogb`, `infinity`, `isinf`, `isnan`, `j0`, `j1`, `jn`, `ldexp`, `log`, `log2`, `log10`, `loglp`, `modf`, `nan`, `nearbyint`, `nextafter`, `pow`, `pow10`, `remainder`, `remquo`, `rint`, `lrint`, `llrint`, `round`, `lround`, `llround`, `scalb`, `scalbln`, `scalbn`, `signbit`, `significand`, `sin`, `sincos`, `sinh`, `sqrt`, `tan`, `tanh`, `trunc`.

- ❑ Efficient, extensible well maintained open source project
- ❑ Provide a front-end for a lot of programming language: C,C++,fortran, etc
- ❑ From GCC version 3.5 it has an unified intermediate representation
- ❑ From version 4.1 it supports some parallelizing techniques: auto-vectorization
- ❑ From version 4.5 it supports custom plugins to integrate external code transformation/analysis
- ❑ Not easy to deal with anyway...



- ❑ building the SSA representation
- ❑ expanding vector to scalar
- ❑ may-alias information
- ❑ dead code elimination
- ❑ partial redundancy elimination
- ❑ full redundancy elimination
- ❑ forward propagation of expressions
- ❑ SSA copy renaming
- ❑ global value numbering
- ❑ sparse conditional constant propagation
- ❑ scalar replacement of aggregates
- ❑ value range propagation
- ❑ loop optimization
- ❑ common subexpression elimination
- ❑ reassociation pass
- ❑ dead store elimination



- ❑ We start from the standard output GCC creates with the `-fdump-tree-all-raw`
 - ▶ We add information we needed
 - ▶ We make it unambiguous
 - ▶ We add a linker starting from this format
 - ▶ We add a new command wrapping GCC:
 - `tree-panda-gcc`

Details in these directories:

- [src/tree](#)
- [src/parser/treegcc](#)
- [src/wrapper/treegcc](#)

```
GCC_VERSION: "4.9.3"
PLUGIN_VERSION: "0.11"
@1 function_decl name: @2 type: @3 scpe: @4 srcp: "/panda/etc/libbambu/assert.c":6:6 uid: 1392
arg: @5 arg: @6 arg: @7 body: @8
@2 identifier_node strg: "__assert" lngt: 8
@3 function_type size: @9 algn: 8 retn: @10 prms: @11
@4 translation_unit_decl srcp: "<built-in>":0:0 uid: 1404
@5 parm_decl name: @12 type: @13 scpe: @1 srcp: "/panda/etc/libbambu/assert.c":6:27 uid: 1389
argt: @13 size: @14 algn: 32 used: 1
@6 parm_decl name: @15 type: @13 scpe: @1 srcp: "/panda/etc/libbambu/assert.c":6:52 uid: 1390
argt: @13 size: @14 algn: 32 used: 1
@7 parm_decl name: @16 type: @17 scpe: @1 srcp: "/panda/etc/libbambu/assert.c":6:64 uid: 1391
argt: @17 size: @14 algn: 32 used: 1
@8 statement_list bloc: 2 loop_id: 0 pred: ENTRY stmt: @18 stmt: @19
```

- ❑ `tree-panda-gcc` could be used to create a custom library that could be deployed by `bambu` by passing `-l` and `-L` options.
- ❑ Example:

```
$ tree-panda-gcc -O3 -o librho_a-rho.o rho.c
$ ar cru librho.a librho_a-rho.o
...
$ tree-panda-gcc -O3 -o keccak keccak-keccak_coproc.o
chi/libchi.a iota/libiota.a pi/libpi.a rho/librho.a
theta/libtheta.a
$ bambu --use-raw keccak
```

Autotools based project described in directory
[examples/crypto_designs/multi-keccak](#)

- ❑ bambu uses the text-based IR and the builtin linker to build a single in-memory representation and perform HLS

▶ `$ bambu file1.c file2.c --top-fname=fname`

- ❑ In case the option `--top-fname` is not used, bambu automatically identifies which function can act as the top one.

- ❑ Based on C description
 - ▶ Global variables
 - ▶ Top function parameters
- ❑ No restriction on top function signature
- ❑ C semantic
 - ▶ Scalar parameters: input
 - ▶ struct/union parameters: input
 - ▶ array/pointer: input/output
 - ▶ return value: output
 - ▶ Global variables: input/output

```
int global;  
  
float my_top_function(int * first, struct s second, int  
third[10], float * fourth)  
{...}
```

❑ Input parameters:

- ▶ Second

❑ Input/output parameters:

- ▶ First, third, fourth

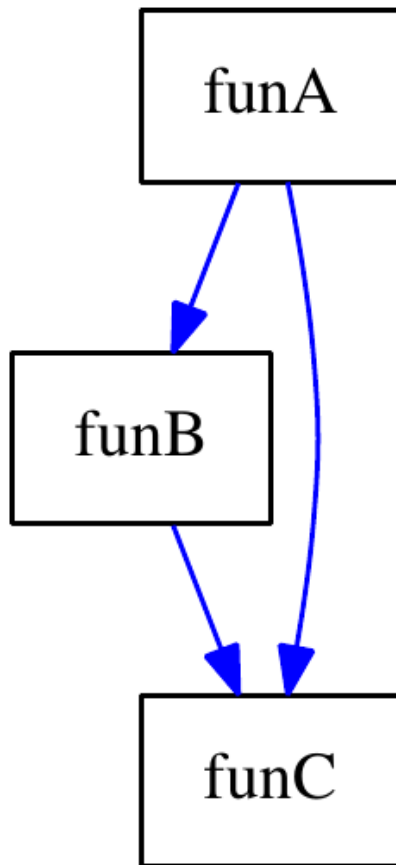
❑ Output parameters:

- ▶ Return value

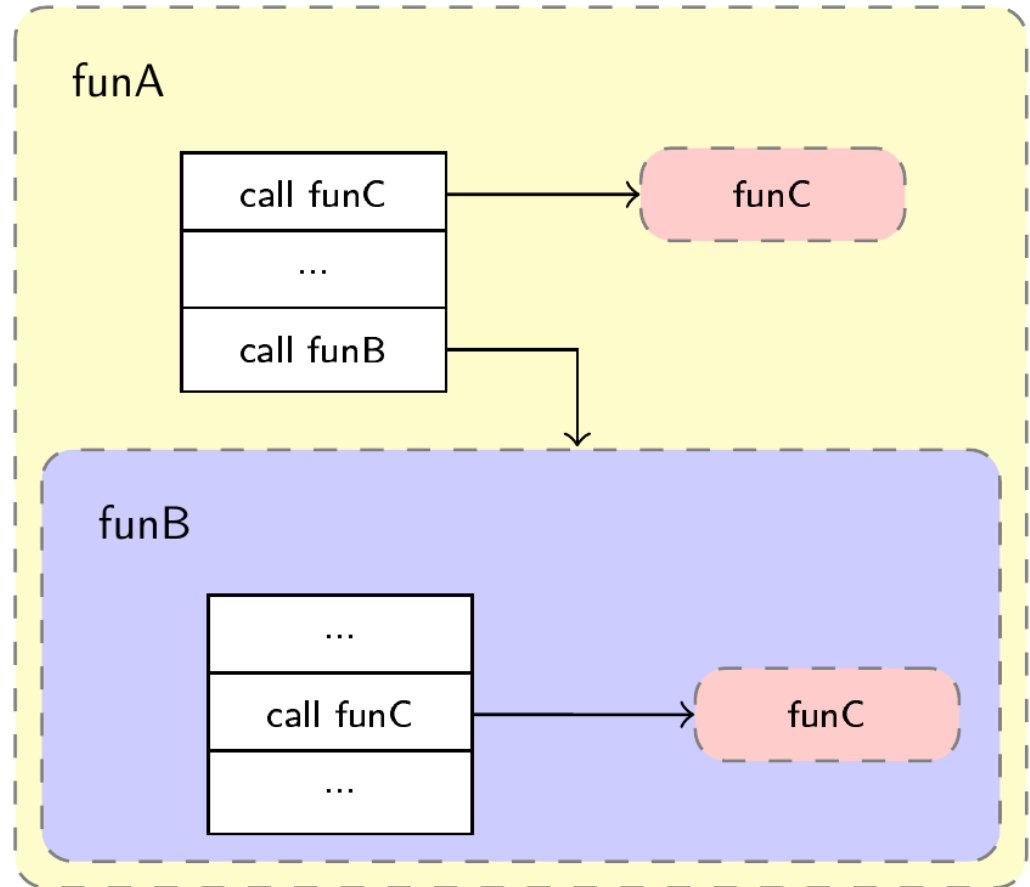
❑ Shared memory:

- ▶ global

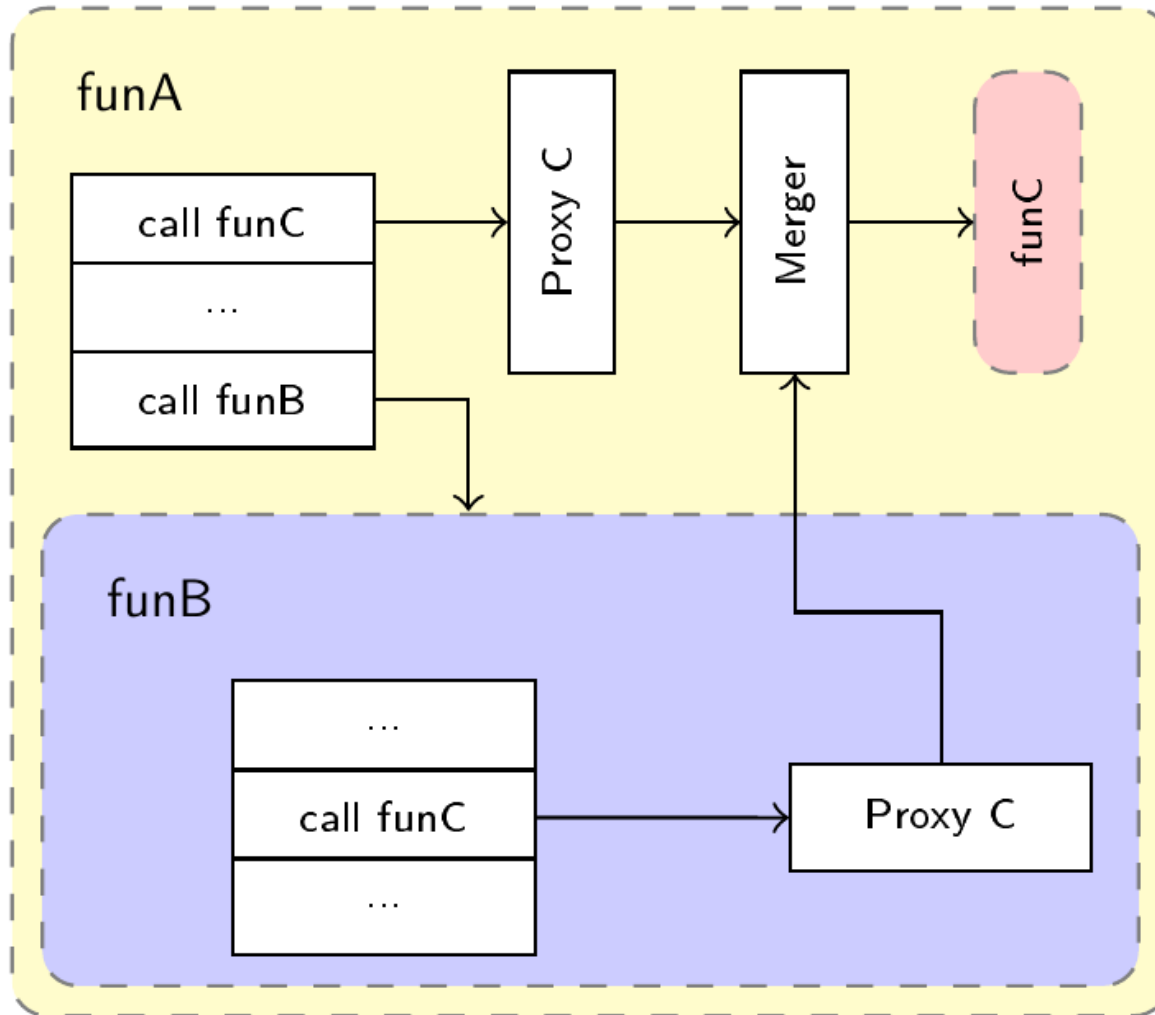
- ❑ One component per function



call graph



RTL hierarchy



- ❑ bambu assumes that all global variables could be accessed by external CPU/accelerators
- ❑ It is possible to change this default with option:
 - ▶ `--do-not-expose-globals`
 - ▶ All global variables are considered local to the compilation units as it they are declared `static`.

❑ From GCC man:

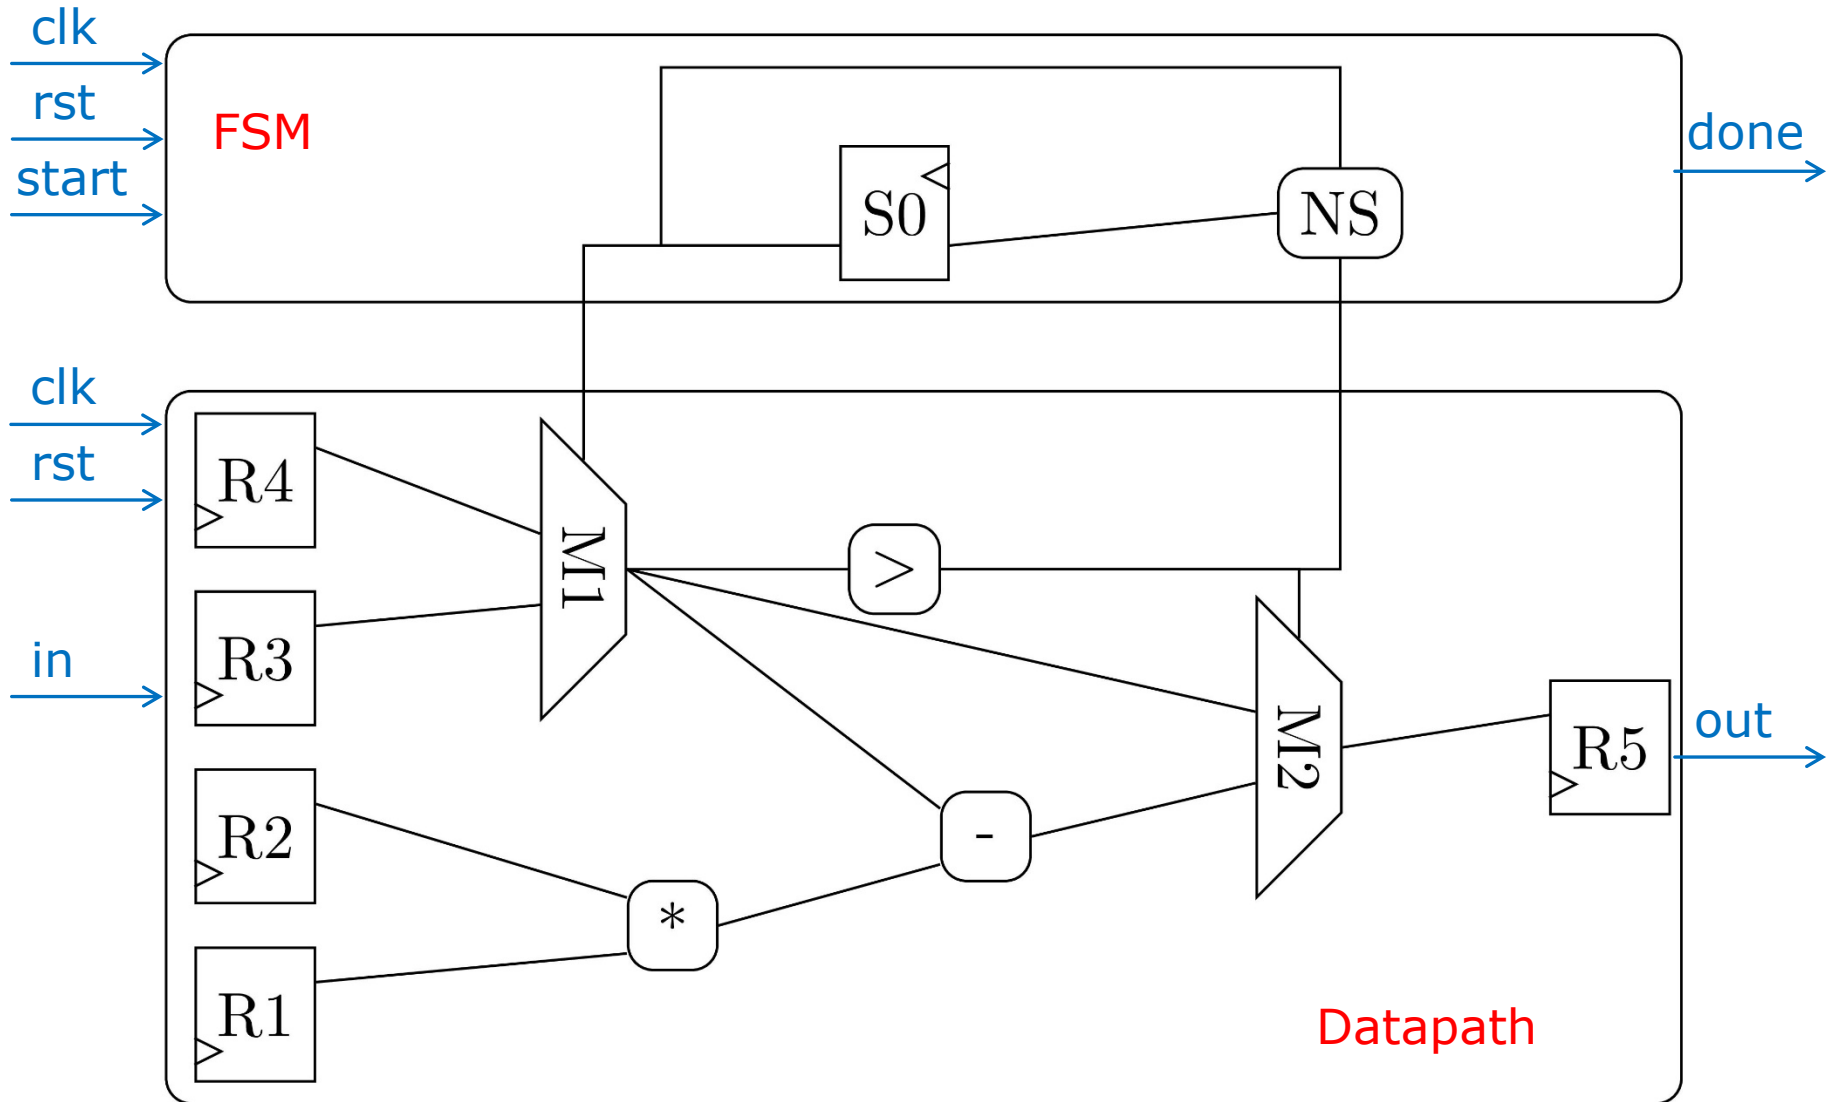
- ▶ Assume that the current compilation unit represents the whole program being compiled. All public functions and variables with the exception of "main" and those merged by attribute "externally_visible" become static functions and in effect are optimized more aggressively by interprocedural optimizers.



- ### ❑ You can apply this optimization when the top function is `main` or when you specify an alternative top function (`--top-fname`) provided that the option `--do-not-expose-globals` is also given (currently supported only with GCC 4.9)

- FSM with datapath (FSMD)
 - set of storage variables VAR
 - set of expressions $EXP = \{\Delta(x, y, z, \dots) \mid x, y, z, \dots \in VAR\}$
 - set of assignments $ASG = \{X = e \mid X \in VAR, e \in EXP\}$
 - Set of status variables $STAT = \{(a, b) \mid a, b \in EXP\}$
- An FSMD is defined by the following tuple
 - $\langle S, I \cup B, O \cup A, \Delta, \Lambda \rangle$
- where S, Δ, Λ are the standard set of states, the output function and the next state function:
 - B = set of some status variables ($B \subseteq STAT$)
 - A = set of storage variable assignments ($A \subseteq ASG$)


Synthesis per function: Architectural template


38



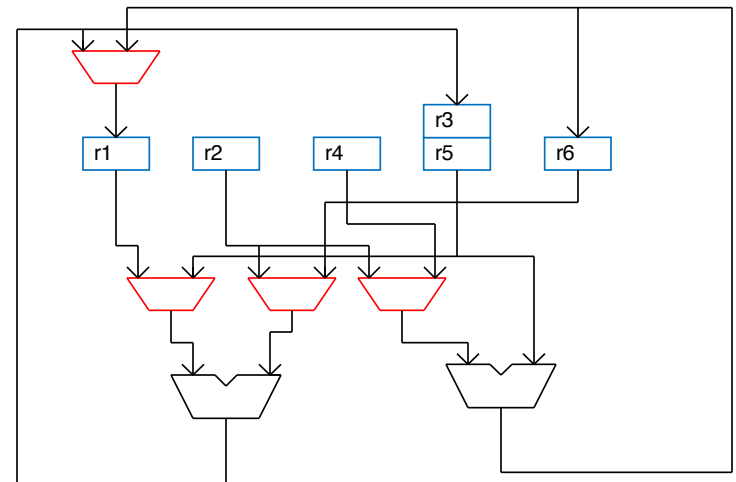


- Program Dependency Graphs – PDG
 - ▶ Data Flow Graph – DFG
 - ▶ Control Dependence Graph - CDG





Module binding
Register binding
Interconnection definition



- ☐ Single cycle
 - ☐ Multi-cycle
 - ☐ Pipelined operator
 - ☐ Multi-function
 - ☐ Variable latency
-
- ☐ XML files to describe operations and resources

```
int laplacian(char *, char *, int, int);
int make_inverse_image(char *, char *, int, int);
int sharpen(char *, char *, int, int);
int sobel(char *, char *, int, int);

int (*pipeline[MAX_DEPTH])(char *, char *, int, int);

void UserApp(char *in, char *out, int x_size, int y_size) {
    // ...
    // Pipeline configuration using function pointers
    add_filter(0, make_inverse_image);
    add_filter(1, sharpen);
    // ...
    // execute is synthesized in hardware
    execute(in, out, x_size, y_size);
}

void execute(char *in, char *out, int x_size, int y_size) {
    int i = 0;
    for (i = 0; i < MAX_PIPELINE_DEPTH; i++) {
        if (pipeline[i] == 0) break;
        // here other hw accelerator are called
        // using function pointers
        int res = pipeline[i](in, out, x_size, y_size);
        if (res != 0) return;
        swap(in, out);
    }
    move_if_odd(i, in, out);
}
```

Adding a memory mapped interface to filters

45

Accelerator
base address



Control register

Input register: in

Input register: out

Input register:
x size


Input register:
y size

Output register

```
void execute(char *in, char *out, int x_size, int y_size) {
    int i = 0;
    for (i = 0; i < MAX_PIPELINE_DEPTH; i++) {
        if (pipeline[i] == 0) break;
        // here other hw accelerator are called
        // using function pointers
        __builtin_indirect_call(
            pipeline[i], 1, in, out, x_size, y_size, &res);
        if (res != 0) return;
        swap(in, out);
    }
    move_if_odd(i, in, out);
}
```

Sequence diagram for function indirect call

47



Call mechanism complexity: $\#cycles = Wl(Np+1) + lhs(Wl+RI)$

Integration of hand-written components in the HLS flow

48

```
#include "module_lib.h"

void my_ip(uint8_t command, uint32_t param1, uint32_t param2) {
    static module1_output_t module1_output;
    static module2_output_t module2_output;
    switch(command) {
        case 0:
            module1(param1, param2 >> 16, &module1_output);
            break;
        case 1:
            module2(param1, &module2_output);
            break;
        case 2:
            printer1(module1_output.output1, module1_output.output2, module1_output.output3,
module1_output.output4);
            break;
        case 3:
            printer2(module2_output.output1, module2_output.output2, module2_output.output3);
            break;
        default:
            break;
    }
}
```


- ❑ Function mapped on IPs has to be declared as extern:
 - ▶ `extern void module1(uint32_t input1, uint16_t input2, module1_output_t *outputs);`
- ❑ C code has to be passed with the following option
 - ▶ `--C-no-parse=module1.c,...`
- ❑ Binding between function **module1** and component **module1** has to be specified with a XML file and passed as an option to bambu
 - ▶ `$ bambu ... module_lib.xml`
- ❑ Check these examples:
 - ▶ `examples/IP_integration`
 - ▶ `examples/breakout`
 - ▶ `examples/pong`
 - ▶ `examples/led_example`

❑ Support of c++ is ongoing:

- ▶ templates
- ▶ C++11 and beyond
- ▶ ac_types from Mentor Graphics could be used
 - ac_int and ac_fixed currently cannot be used in the function interfaces
 - not c++11 friendly

```
#include <algorithm>
int gcd(int x, int y )
{
    if( x < y )
        std::swap( x, y );

    while( y > 0 )
    {
        int f = x % y;
        x = y;
        y = f;
    }
    return x;
}
```

- * euclid.f (FORTRAN 77)
- * Find greatest common divisor using the Euclidean algorithm

```
FUNCTION NGCD (NA, NB)
  IA = NA
  IB = NB
1  IF (IB.NE.0) THEN
    ITEMP = IA
    IA = IB
    IB = MOD (ITEMP, IB)
    GOTO 1
  END IF
  NGCD = IA
  RETURN
END
```

By default parameters are passed
by reference

- * euclid.f (FORTRAN 77)
- * Find greatest common divisor using the Euclidean algorithm

```
FUNCTION NGCD(NA, NB)
  integer, value :: NA ! input
  integer, value :: NB ! input
  IA = NA
  IB = NB
1  IF (IB.NE.0) THEN
    ITEMP = IA
    IA = IB
    IB = MOD(ITEMP, IB)
    GOTO 1
  END IF
  NGCD = IA
  RETURN
END
```

Here parameters are passed
by copy

❑ Fortran support in progress...

❑ #pragma omp simd is supported

```
#pragma omp declare simd
void add (int accelnum, int startidx, int endidx)
{
    int sum = 0;
    int i;
    for (i = 0; i < OPS_PER_ACCEL; i++) {
        sum += array[i+startidx];
    }
    output[accelnum] = sum;
}

main() {
    ...
    #pragma omp simd
    for (i = 0; i < NUM_ACCELS; i++)
    {
        add(i, i * OPS_PER_ACCEL, (i + 1)*OPS_PER_ACCEL);
    }
}
```

❑ OmpMP support in progress...