



# Synthesis and Optimization of Complex Memory Operations

 ***Tutorial @ FPT Conference 2017 – Melbourne – Australia***

**Christian Pilato**

Università della Svizzera italiana (USI)  
Facoltà di Informatica  
[christian.pilato@usi.ch](mailto:christian.pilato@usi.ch)

**Fabrizio Ferrandi**

Politecnico di Milano  
Dipartimento di Elettronica, Informazione  
e Bioingegneria  
[fabrizio.ferrandi@polimi.it](mailto:fabrizio.ferrandi@polimi.it)

- ❑ Algorithms operate on data and data must be stored somewhere
- ❑ Memories are responsible for **70-90% of total cost**
  - ▶ In FPGA:
    - Small/frequently-accessed sets of data in distributed registers (low latency, high resource cost)
    - Medium sets of data in BRAMs (limited number of ports, limited number of resources)
    - Large sets of data in the off-chip memory (e.g., DRAM)

**How to efficiently implement all memory operations?**

## ❑ Software:

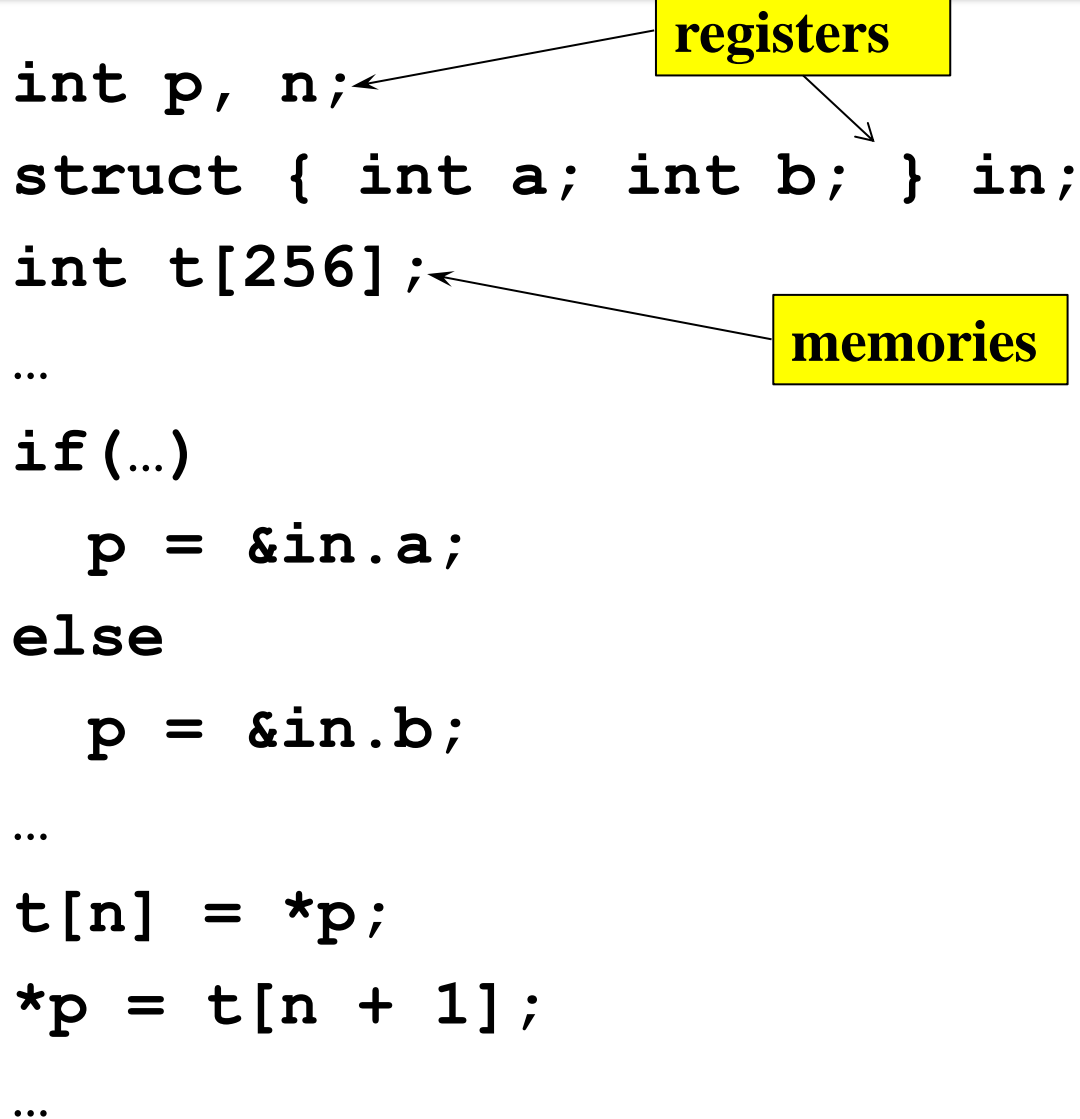
- ▶ *Stack* for storing data determined statically
- ▶ *Heap* for dynamically allocated data
- ▶ *Cache hierarchies* for hiding the latency
- ▶ *Sequential memory operations*

## ❑ Hardware

- ▶ *Heterogeneous and distributed memories*
- ▶ *Abundant hardware parallelism*
- ▶ *No “flexibility” during the execution*

# Example

4



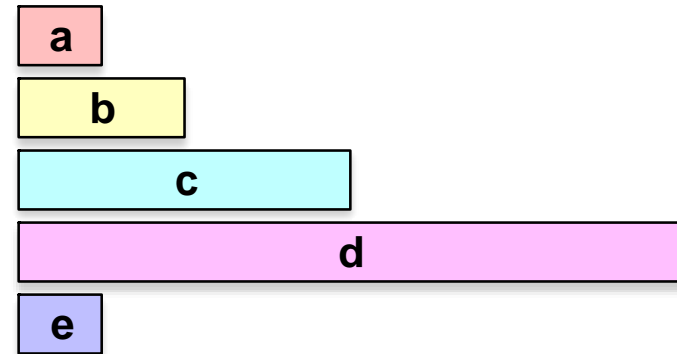
The diagram illustrates the mapping of C code to hardware resources. A yellow box labeled 'registers' has arrows pointing to the variables 'p' and 'n' in the first line of code. Another yellow box labeled 'memories' has an arrow pointing to the array 't' in the third line of code.

```
int p, n;
struct { int a; int b; } in;
int t[256];
...
if (...)
    p = &in.a;
else
    p = &in.b;
...
t[n] = *p;
*p = t[n + 1];
...
```

- ❑ **Arrays of primitive data types** easily implemented with HDL templates (supported by any synthesis tool)
- ❑ There are three types of memory operations to be supported:
  - ▶ Accesses to **complex data structures** (e.g., arrays of complex structures)
  - ▶ **Parameter passing** in function calls
  - ▶ Operation on **pointers** (dynamic resolution, pointer arithmetic, dynamic memory allocation)

**It is relatively easy to implement memory operations, it is NOT easy to implement them efficiently**

```
typedef struct
{
    char a;          // 1 byte
    short b;         // 2 bytes
    long c;          // 4 bytes
    long long d;     // 8 bytes
    char e;          // 1 byte
} mystruct; // tot: 16 bytes
```

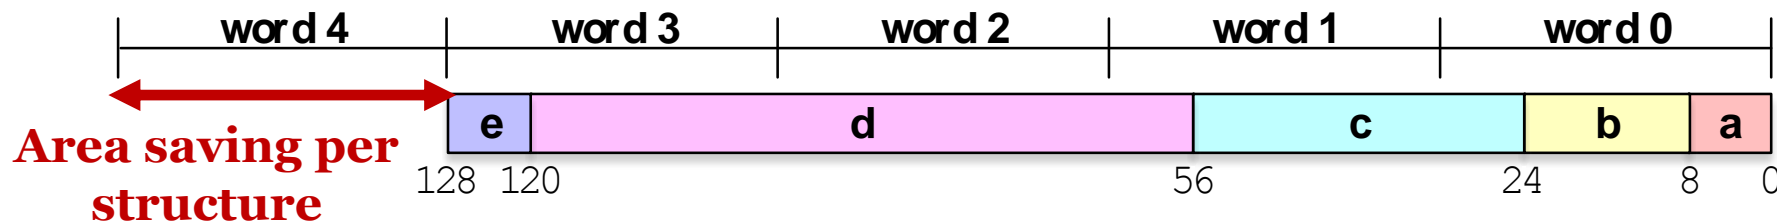
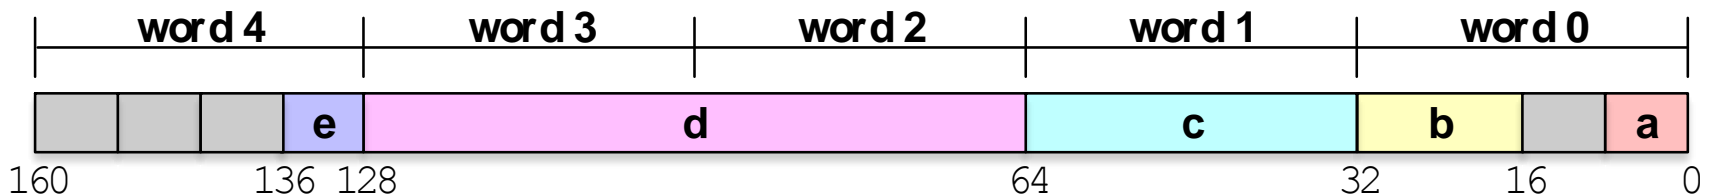


- ❑ How to *efficiently* store an array of **mystruct** structures?
- ❑ How to *efficiently* implement memory operations?

**Reduce the memory footprint without compromising the execution latency**

## ❑ Software annotations to create compact representations

```
typedef struct
{
    char a;          // 1 byte
    short b;         // 2 bytes
    long c;          // 4 bytes
    long long d;     // 8 bytes
    char e;          // 1 byte
} __attribute__((packed)) mystruct;
```



## ❑ Alternative solutions:

- ▶ Vivado HLS decomposes the *array of structures* into *multiple arrays of single fields*
- ▶ LegUp uses *subarrays* as large as the maximum field (support only for aligned accesses)

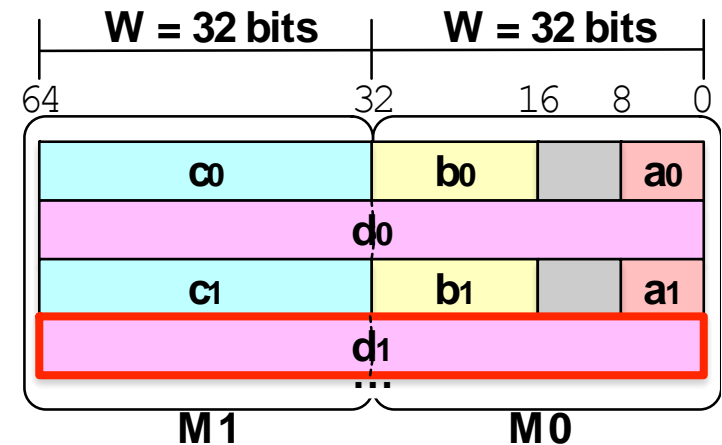
## ❑ We use a **single memory template** to implement both aligned and unaligned accesses

- ▶ *Portable and easy to maintain*
- ▶ *Simple logic* to convert datapath requests into actual memory operations

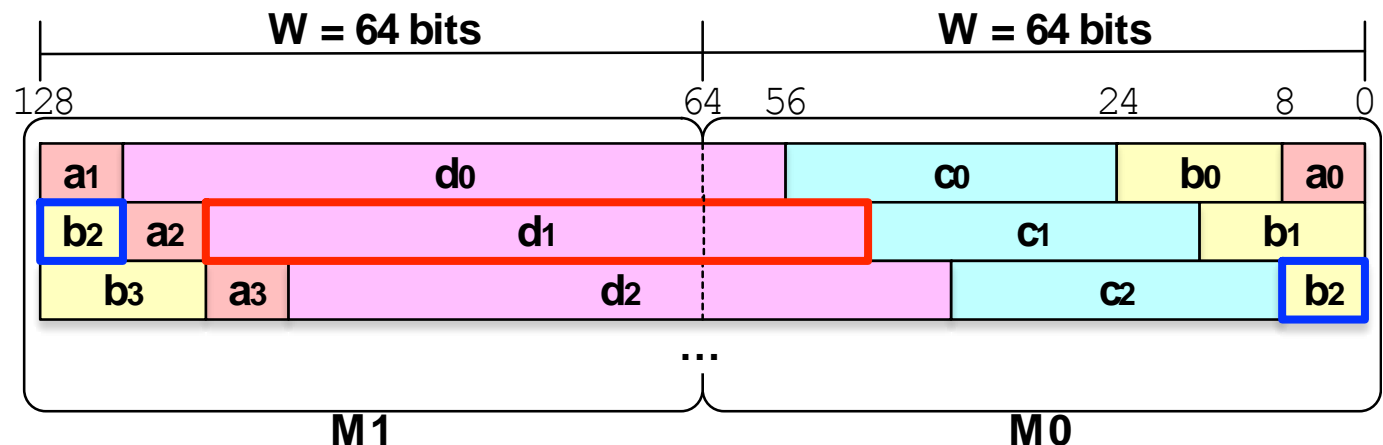


- Two arrays with bitwidth based of the type of accesses
  - At most two (parallel) memory ops for each field

**Aligned structures**  
( $W = SMAX/2$ )

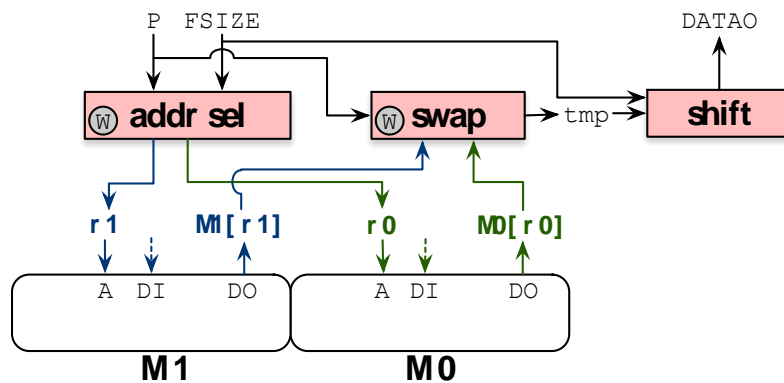


**Unaligned structures**  
( $W = SMAX$ )



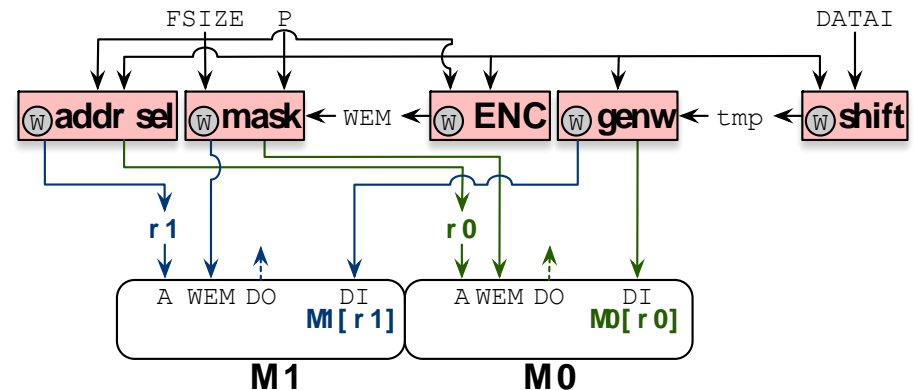
- Unique microarchitecture for both aligned and unaligned accesses

## Read operations




array[0].a -> M0[0][7:0]  
 array[0].b -> M0[0][23:8]  
 array[0].c -> M0[0][55:24]  
 array[0].d -> {M1[0][55:0], M0[0][63:56]}  
 ...  
 array[1].d -> {M1[1][47:0], M0[1][63:48]}  
 ...  
 array[2].b -> {M0[2][7:0], M1[1][63:56]}

## Write operations



array[0].a -> M0[0][7:0]  
 array[0].b -> M0[0][31:16]  
 array[0].c -> M1[0][31:0]  
 array[0].d -> {M1[1][31:0], M0[1][31:0]}  
 array[1].a -> M0[2][7:0]  
 ...  
 array[1].d -> {M1[3][31:0], M0[3][31:0]}

- ❑ Hardware modules are created hierarchically based on the *call graph*
  - ▶ *Software*: parameters on the stack
  - ▶ *Hardware*: actual values provided at the input ports



**When to insert registers to avoid long critical paths?**

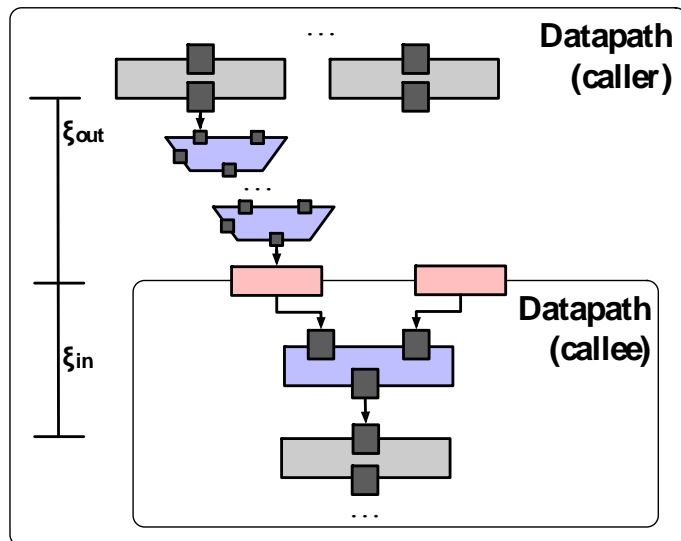
- ❑ We introduce input registers when

$$(\xi_{\text{out}} + \xi_{\text{in}}) > T * \beta$$

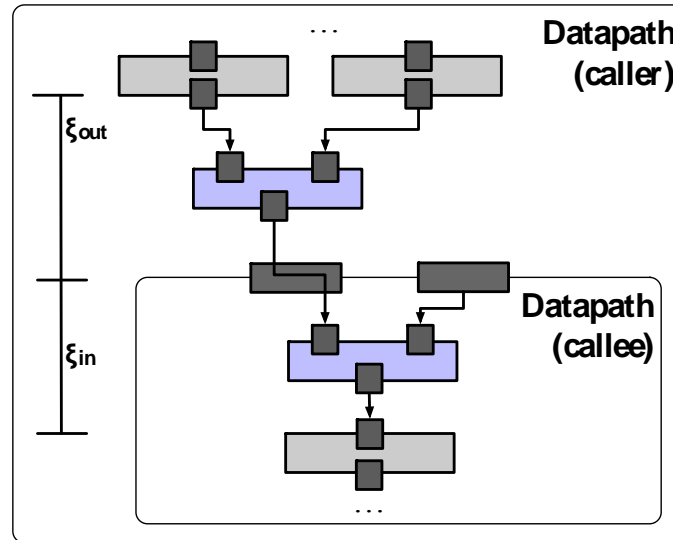
# Some Examples of Input Registering

12

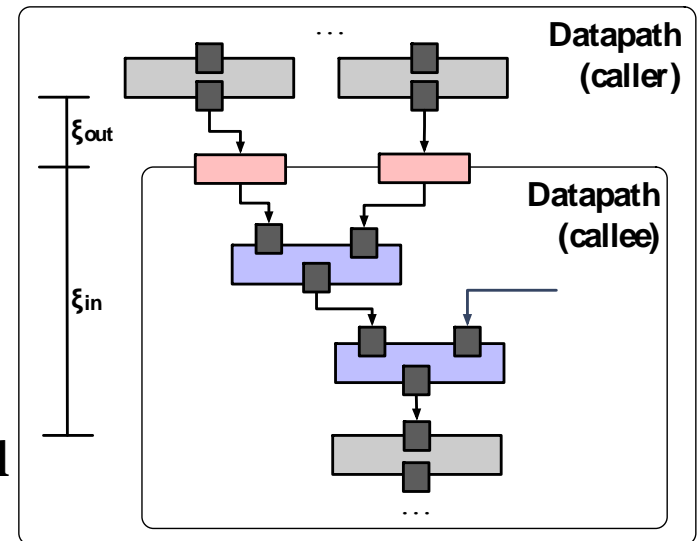
## Long external logic ( $\xi_{out}$ )



## Direct connection



## Long internal logic ( $\xi_{in}$ )



- ❑ We require a **configurable memory space** to store the created data structures
  - ▶ *Memory management* is usually intrinsic within the algorithm
  
- ❑ We leverage **Memmgr**
  - <https://github.com/eliben/code-for-blog/tree/master/2008/memmgr>
  - ▶ library for dynamic memory allocation to a configurable space with a synthesizable **malloc** function
  - ▶ *pre-allocated memory space* with a default size of 384 KB
  - ▶ memory space either in a local memory or in DRAM
  - ▶ memory component to manage *pointer-based requests*

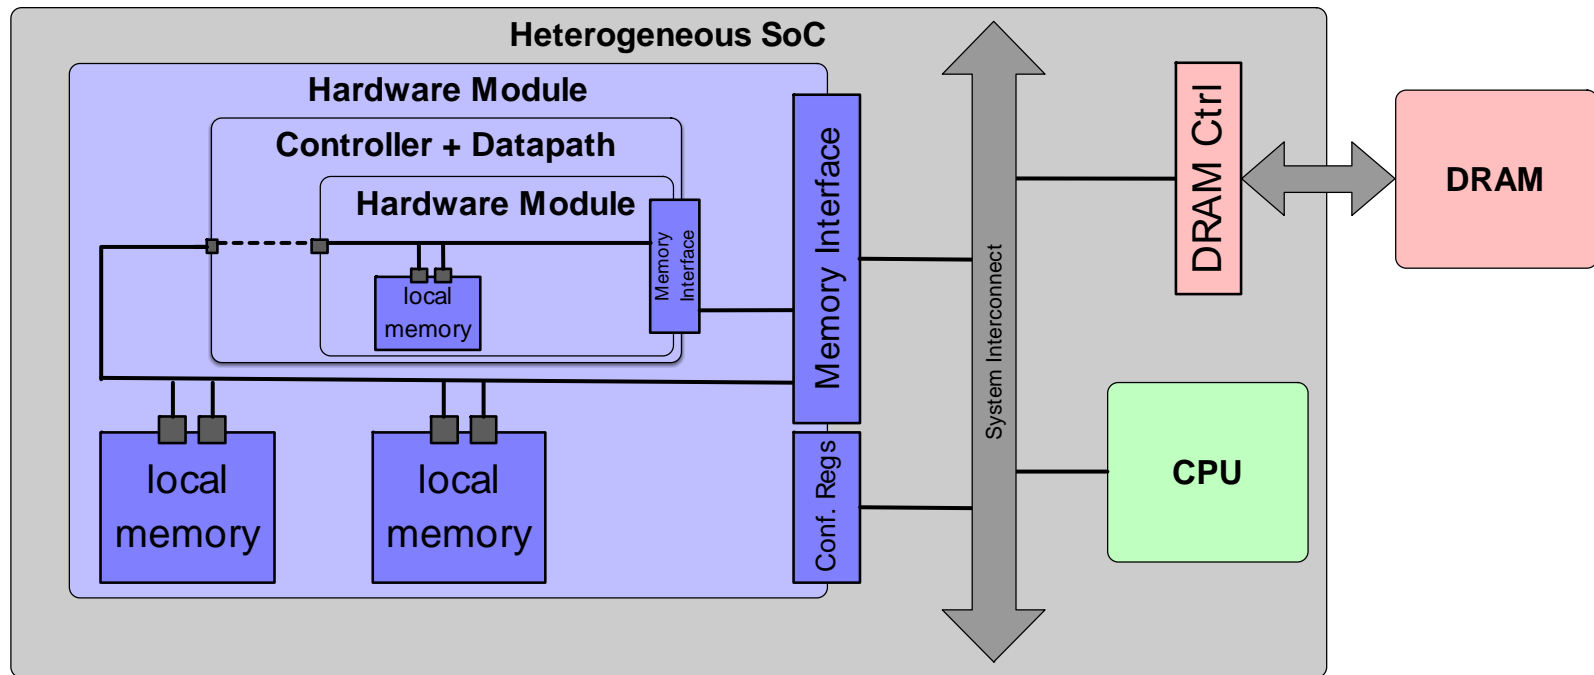
- ❑ Two conflicting goals:
  - ▶ Maintain **flexibility** as in software (operations on pointers, dynamic resolution, etc.)
  - ▶ Increase **efficiency** as in hardware (exploit hardware parallelism whenever possible)

## Definition:

***Points-to set:*** set of data structures that can be potentially accessed by the pointer used in a memory operation

- ❑ This information is obtained during compilation by means of (sophisticated) **alias analysis**

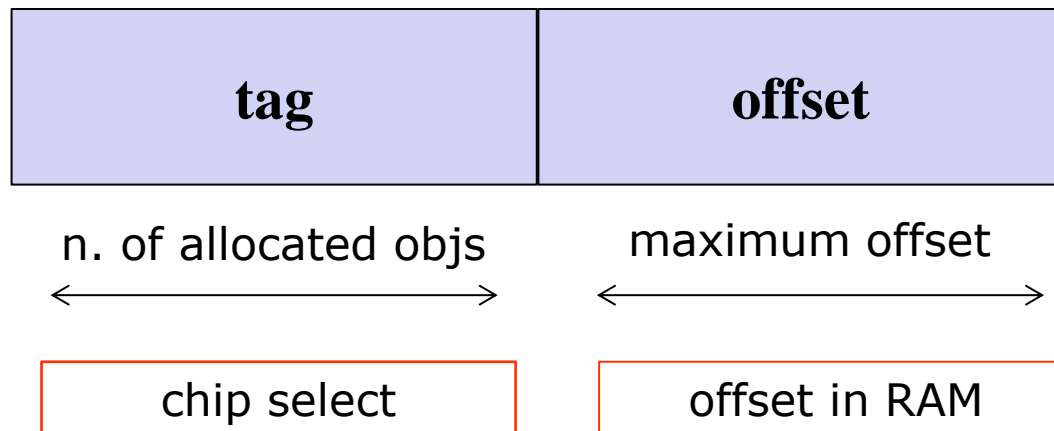
- ❑ **Direct connections** to datapath operators if the operation is completely defined
- ❑ **Internal memory bus** to connect all memory components potentially accessed by unresolved memory operations (*points-to set*)



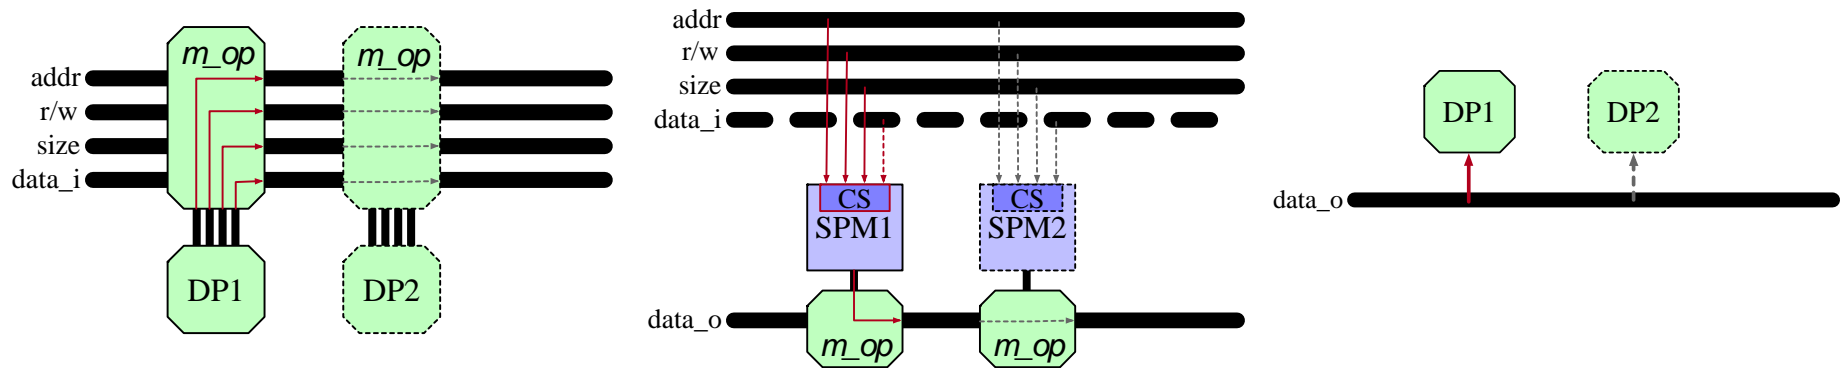
- ❑ **Memory analysis phase** to define an address for any
  - ▶ global scalar/aggregate variables,
  - ▶ local aggregate variables
  - ▶ local scalar variables used as argument of operator &
  
- ❑ Pointers are stored in standard registers
  - ▶ They can be considered as standard variables
  
- ❑ Load and store memory operations can be implemented as standard datapath operations
  - ▶ Connected to dedicated memory resources (memory controllers)



- We envision an **unique address space**, where memory data has an address associated with:
  - ▶ Contiguous or tag-based **address allocation**
  - ▶ Tag-based can simplify the control logic

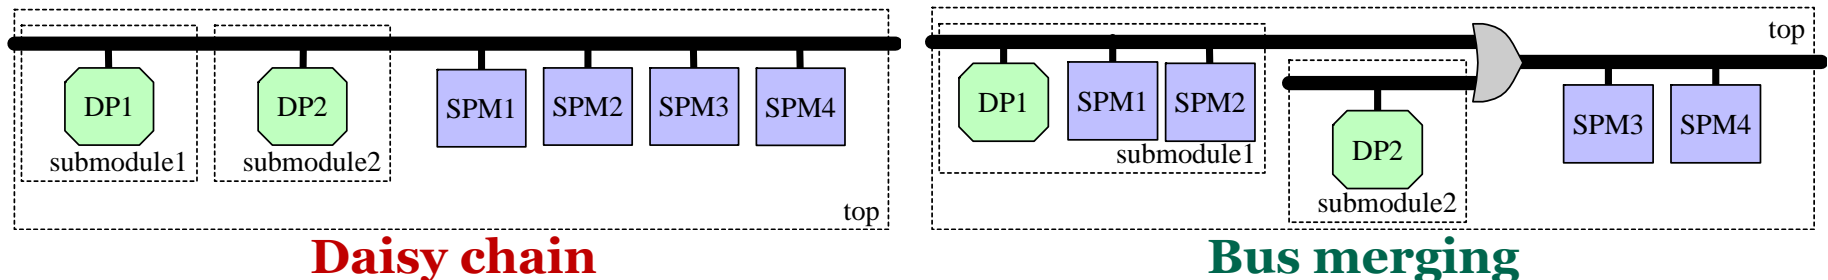


## ❑ Very powerful memory microarchitecture

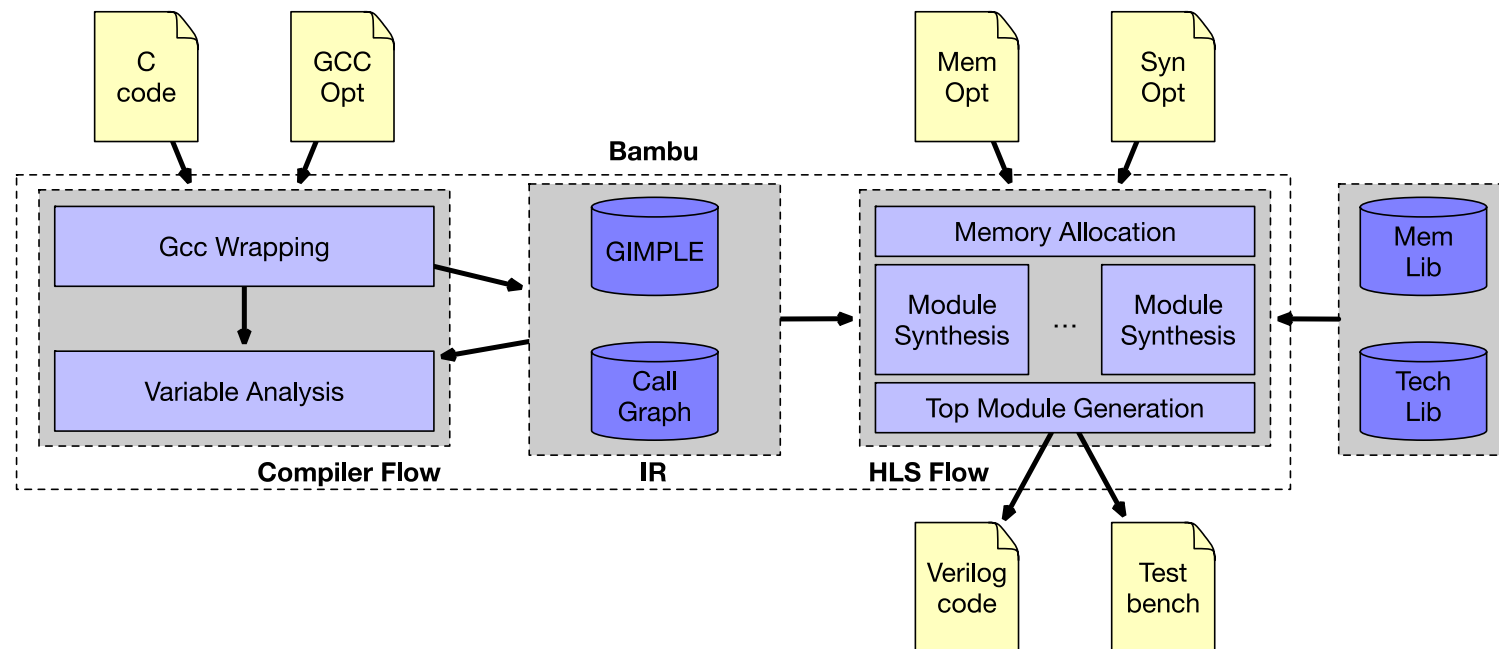


## ❑ Possibility of creating long combinational paths with many memory components connected in chain

- Dominator-based analysis and bus merging (OR gates) for localizing the memories as much as possible



- ❑ **Automatic generation and optimization** fully implemented in Bambu
  - ▶ GCC for extracting **memory-related information**
  - ▶ Extensive set of **memory parameters** to explore many alternative configurations



## ❑ Reference/Naïve approach

- ▶ all data structures with standard memory components, instantiated in the top module and on the internal bus (**maximum flexibility**)

## ❑ Privatization

- ▶ flow-sensitive pointer analysis to reduce the points-to set
- ▶ resolved memory operations directly connected to the memories

## ❑ Creation and duplication of **read-only memories**

## ❑ Conversion into **distributed memories**

- ▶ configurable threshold
- ▶ higher threshold for read-only memories (simpler)



## ❑ **Memory registering**

- ▶ potential timing violations with complex memory operations and high target frequency
- ▶ registers between datapath resources and memory ports

## ❑ **Localization**

- ▶ dominator-based analysis to reduce long combinational paths



- ❑ Memory components can reach 312 MHz (XC7Zo20) and 608 MHz (XC7VX690T)
- ❑ Selected applications from a recent survey on HLS tools (Nane et al. TCAD 2016)
  - ▶ Frequency of 400MHz
  - ▶ Up to 54% of memory reduction



❑ Memory components can reach 600 MHz (Stratix-V)

❑ Selected applications from a recent survey on HLS tools  
(Nane et al. TCAD 2016)

- ▶ Frequency of 400MHz
- ▶ Up to 54% of memory reduction



## Explore different allocation policy

`--memory-allocation-policy=<type>`

Set the policy for memory allocation. Possible values for the <type> argument are the following:

- |                                 |  |
|---------------------------------|--|
| <code>ALL_BRAM</code>           | - all objects that need to be stored in memory are allocated on BRAMs (default)              |
| <code>LSS</code>                | - all local variables, static variables and strings are allocated on BRAMs                   |
| <code>GSS</code>                | - all global variables, static variables and strings are allocated on BRAMs                  |
| <code>NO_BRAM</code>            | - all objects that need to be stored in memory are allocated on an external memory           |
| <code>EXT_PIPELINED_BRAM</code> | - all objects that need to be stored in memory are allocated on an external pipelined memory |



```
$ bambu adpcm.c --memory-allocation-policy=LSS  
--clock-period=15 --simulate -v3
```

**Look for the log section:**

Memory allocation information:

```
Variable external to the top module: test_result - 25438 -  
test_result
```

```
Id: 25438
```

```
Base Address: 1073741824
```

```
Size: 400
```

```
Is a Read Only Memory
```

```
Used &(object)
```

```
Number of functions in which is used: 1
```

```
Maximum number of references per function: 1
```

```
Maximum number of loads per function: 1
```

```
...
```

```
$ bambu adpcm.c --memory-allocation-policy=ALL_BRAM  
--clock-period=15 --simulate -v3
```

```
$ bambu adpcm.c --memory-allocation-policy=LSS  
--clock-period=15 --simulate -v3
```

```
$ bambu adpcm.c --memory-allocation-policy=GSS  
--clock-period=15 --simulate -v3
```

```
$ bambu adpcm.c --memory-allocation-policy=NO_BRAM  
--clock-period=15 --simulate -v3
```

```
$ bambu adpcm.c  
--memory-allocation-policy=EXT_PIPELINED_BRAM  
--clock-period=15 --simulate -v3
```

# Second example – Multi-channel design space exploration

27

```
--channels-type=<type>
```

Set the type of memory connections.

Possible values for <type> are:

- MEM\_ACC\_11 - the accesses to the memory have a single direct connection or a single indirect connection (default)
- MEM\_ACC\_N1 - the accesses to the memory have n parallel direct connections or a single indirect connection
- MEM\_ACC\_NN - the accesses to the memory have n parallel direct connections or n parallel indirect connections

```
--channels-number=<n>
```

Define the number of parallel direct or indirect accesses.

- ❑ When BRAMs are involved only two ports at maximum could be given
- ❑ When option  
--memory-allocation-policy=EXT\_PIPELINED\_BRAM  
is given the number of channels could be greater than 2

```
$ bambu adpcm.c --channels-type=MEM_ACC_NN --memory-  
allocation-policy=EXT_PIPELINED_BRAM --channels-number=4  
--clock-period=15 --simulate -v3
```

Look how long it take the simulation.

Consider `-fwhole-program` option

# Third example – Control the Load/Store latency

29

`--memory-ctrl-type=type`

Define which type of memory controller is used.

Possible values for the <type> argument are the following:

D00 - no extra delay (default)

D10 - 1 clock cycle extra-delay for LOAD, 0 for STORE

D11 - 1 clock cycle extra-delay for LOAD, 1 for STORE

D21 - 2 clock cycle extra-delay for LOAD, 1 for STORE

`--bram-high-latency=[3,4]`

Assume a 'high latency bram'-'faster clock frequency'  
block RAM memory based architectures:

3 => LOAD(II=1,L=3) STORE(1).

4 => LOAD(II=1,L=4) STORE(II=1,L=2).

`--mem-delay-read=value`

Define the external memory latency when LOAD are performed (default 2).

`--mem-delay-write=value`

Define the external memory latency when LOAD are performed (default 1).

```
$ bambu mips.c --memory-ctrl-type=D21 --channels-  
type=MEM_ACC_NN --memory-allocation-policy=EXT_PIPELINED_BRAM  
--channels-number=4  
--clock-period=15 --simulate -v3
```

Look how long it take the simulation.

```
$ bambu mips.c --bram-high-latency=4 --channels-  
type=MEM_ACC_NN --clock-period=15 --simulate -v3
```

Look how long it take the simulation.

# Fourth example – Control asynchronous memories inference

31

`--do-not-use-asynchronous-memories`

Do not add asynchronous memories to the possible set of memories used by bambu during the memory allocation step.

`--distram-threshold=value`

Define the threshold in bitsize used to infer DISTRIBUTED/ASYNCHRONOUS RAMs (default 256).

```
$ bambu mips.c --do-not-use-asynchronous-memories -fwhole-  
program --clock-period=15 --simulate -v3
```

Look how long it take the simulation.

```
$ bambu mips.c --distram-threshold=1024 -fwhole-program --  
clock-period=15 --simulate -v3
```

Look how long it take the simulation.



`--unaligned-access`

Use only memories supporting unaligned accesses.

`--aligned-access`

Assume that all accesses are aligned and so only memories supporting aligned accesses are used.

```
$ bambu mips.c --unaligned-access -fwhole-program --clock-  
period=15 --simulate -v3
```

Look how long it take the simulation.

```
$ bambu mips.c --aligned-access -fwhole-program --clock-  
period=15 --simulate -v3
```

Look how long it take the simulation.

# Sixth example – customize memory layout

35

```
--base-address=address
```

Define the starting address for objects allocated externally to the top module.

```
--initial-internal-address=address
```

Define the starting address for the objects allocated internally to the top module.



```
$ bambu mips.c --base-address=1024 --memory-allocation-  
policy=LSS --clock-period=15 --simulate -v3
```

Look how long it take the simulation.

```
$ bambu mips.c --initial-internal-address=0 --base-  
address=1024 --memory-allocation-policy=LSS --clock-period=15  
--simulate -v3
```

Look how long it take the simulation.

`--sparse-memory[=on/off]`

Control how the memory allocation happens.

on - allocate the data in addresses which reduce the decoding logic (default)

off - allocate the data in a contiguous addresses.

`--serialize-memory-accesses`

Serialize the memory accesses using the GCC virtual use-def chains without taking into account any alias analysis information.

`--do-not-chain-memories`

When enabled LOADs and STOREs will not be chained with other operations.

`--rom-duplication`

Assume that read-only memories can be duplicated in case timing requires.

`--do-not-expose-globals`

All global variables are considered local to the compilation units.

`--data-bus-bitsize=<bitsize>`

Set the bitsize of the external data bus.

`--addr-bus-bitsize=<bitsize>`

Set the bitsize of the external address bus.