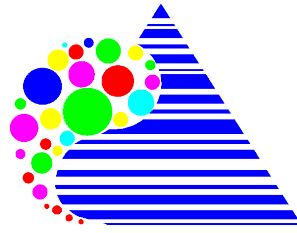
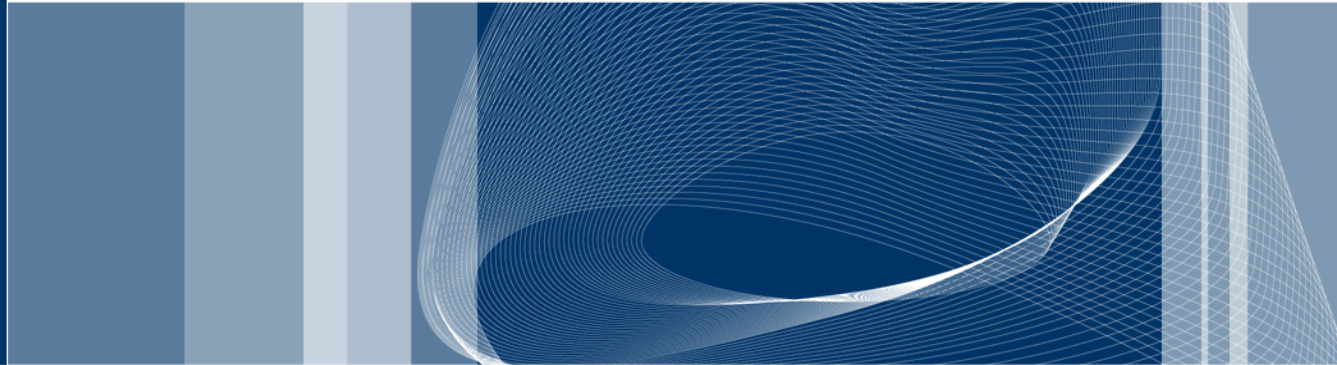


The 28th International Conference on Parallel Architectures and Compilation Techniques

Seattle, WA, USA, September 21st, 2019



 POLITECNICO DI MILANO



PACT19 Tutorial

Bambu: Productive FPGA Programming for Complex Parallel Applications



Fabrizio Ferrandi (Politecnico di Milano), Vito Giovanni Castellana (PNNL), Marco Minutoli (PNNL), Marco Lattuada (Politecnico di Milano), Antonino Tumeo (PNNL)

- ❑ Presentation of bambu
- ❑ Exploiting Vectorization in High Level Synthesis of Nested Irregular Loops
- ❑ Compiler Based Optimizations, Tuning and Customization of Generated Accelerators
- ❑ Target Customization and Tool Integration
- ❑ Enabling the High Level Synthesis of Data Analytics Accelerators

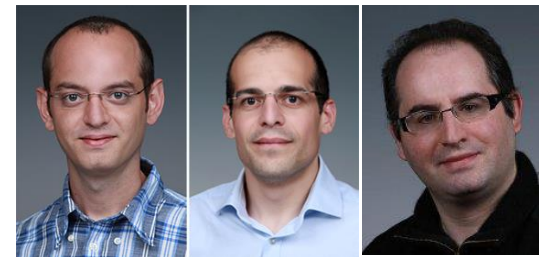
□ Politecnico di Milano

- ▶ Fabrizio Ferrandi
- ▶ Marco Lattuada
- ▶ Christian Pilato
- ▶ Pietro Fezzardi



□ PNNL

- ▶ Vito Giovanni Castellana
- ▶ Marco Minutoli
- ▶ Antonino Tumeo

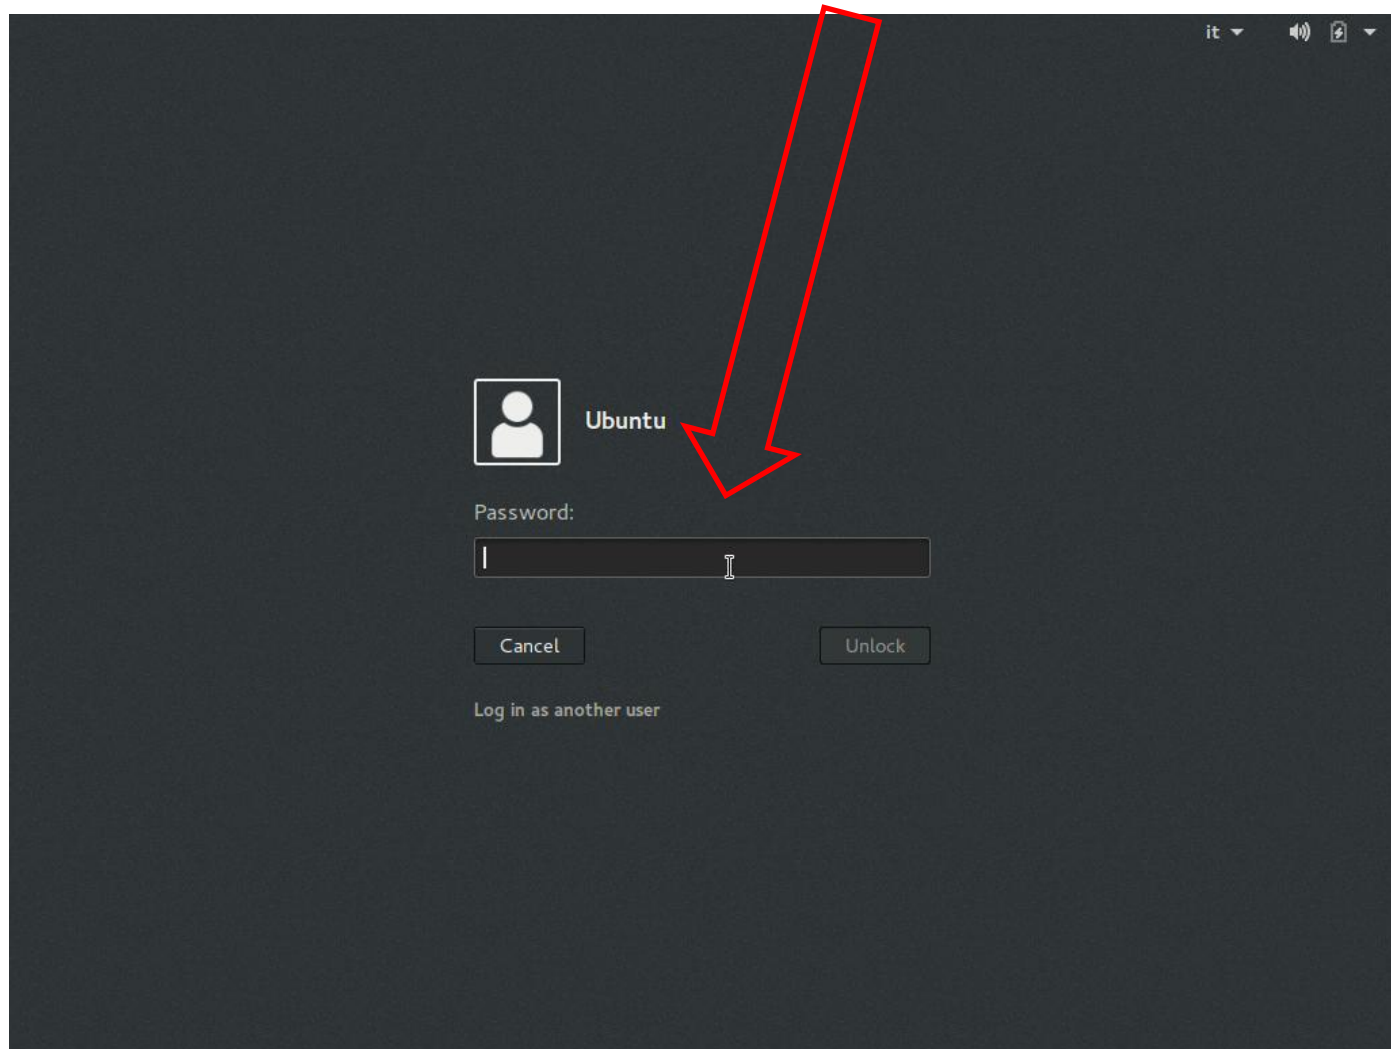


- ❑ Slides and the VM with a pre-configured environment can be downloaded from:

https://panda.dei.polimi.it/?page_id=838

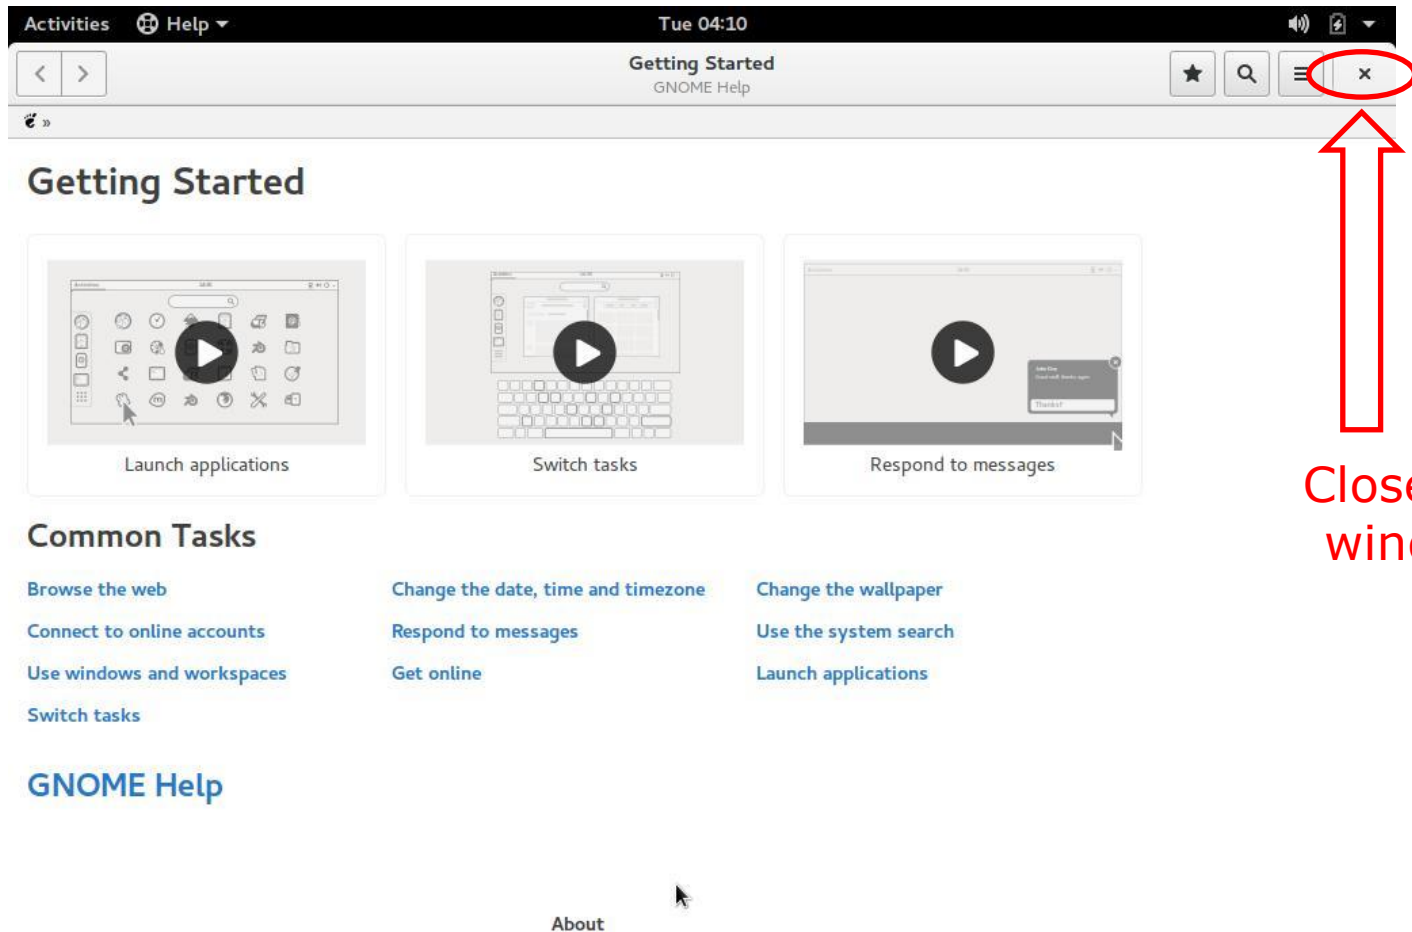
- ❑ There are two .ova files. One for 32bit operating systems and another one for 64bit operating systems.
- ❑ Once downloaded you can import the virtual machine by clicking on the VirtualBox Manager menu:
 - ▶ File
 - Import appliance...

Type as a password: **password**



Starting screen

7



Configure the keyboard

8

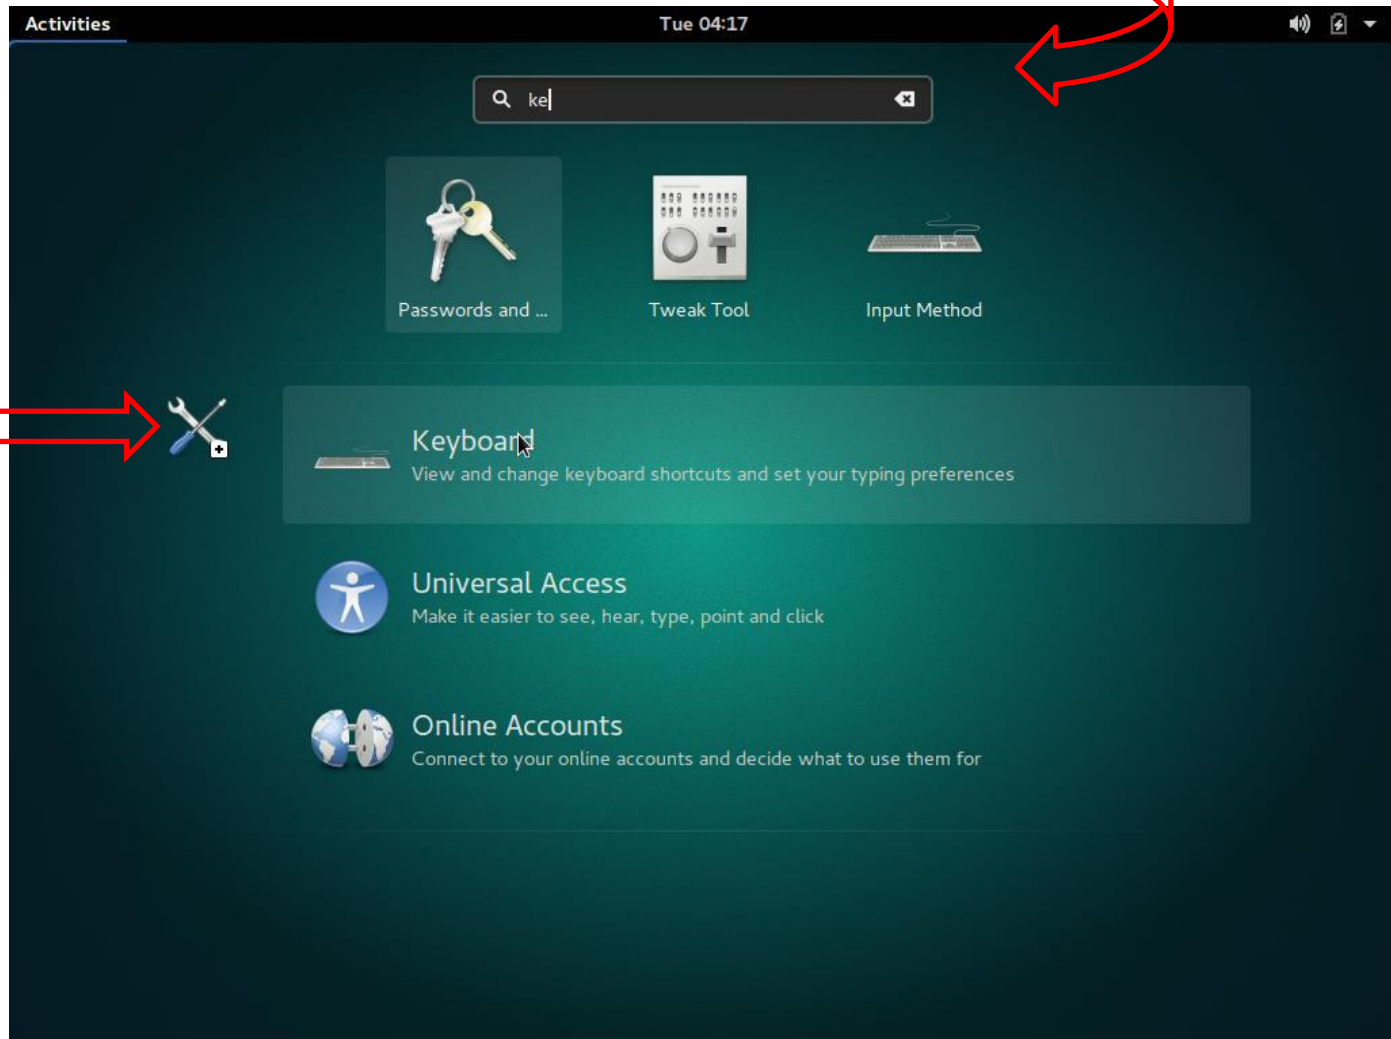
1) Click on **activities**



2) Search for "**keyboard**"

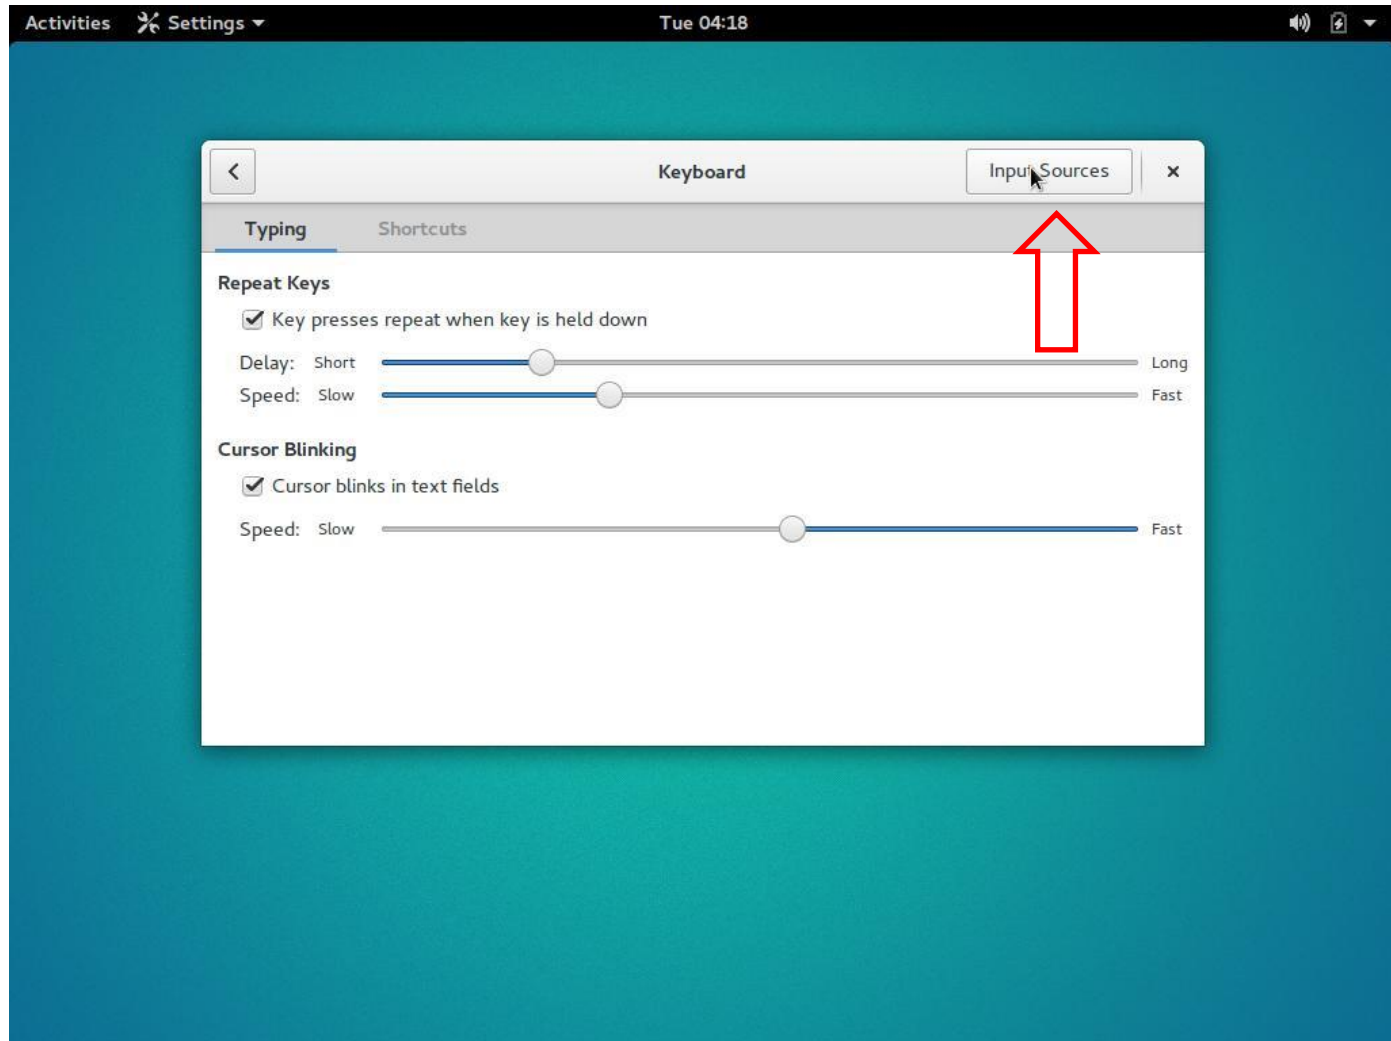


3) Select
keyboard



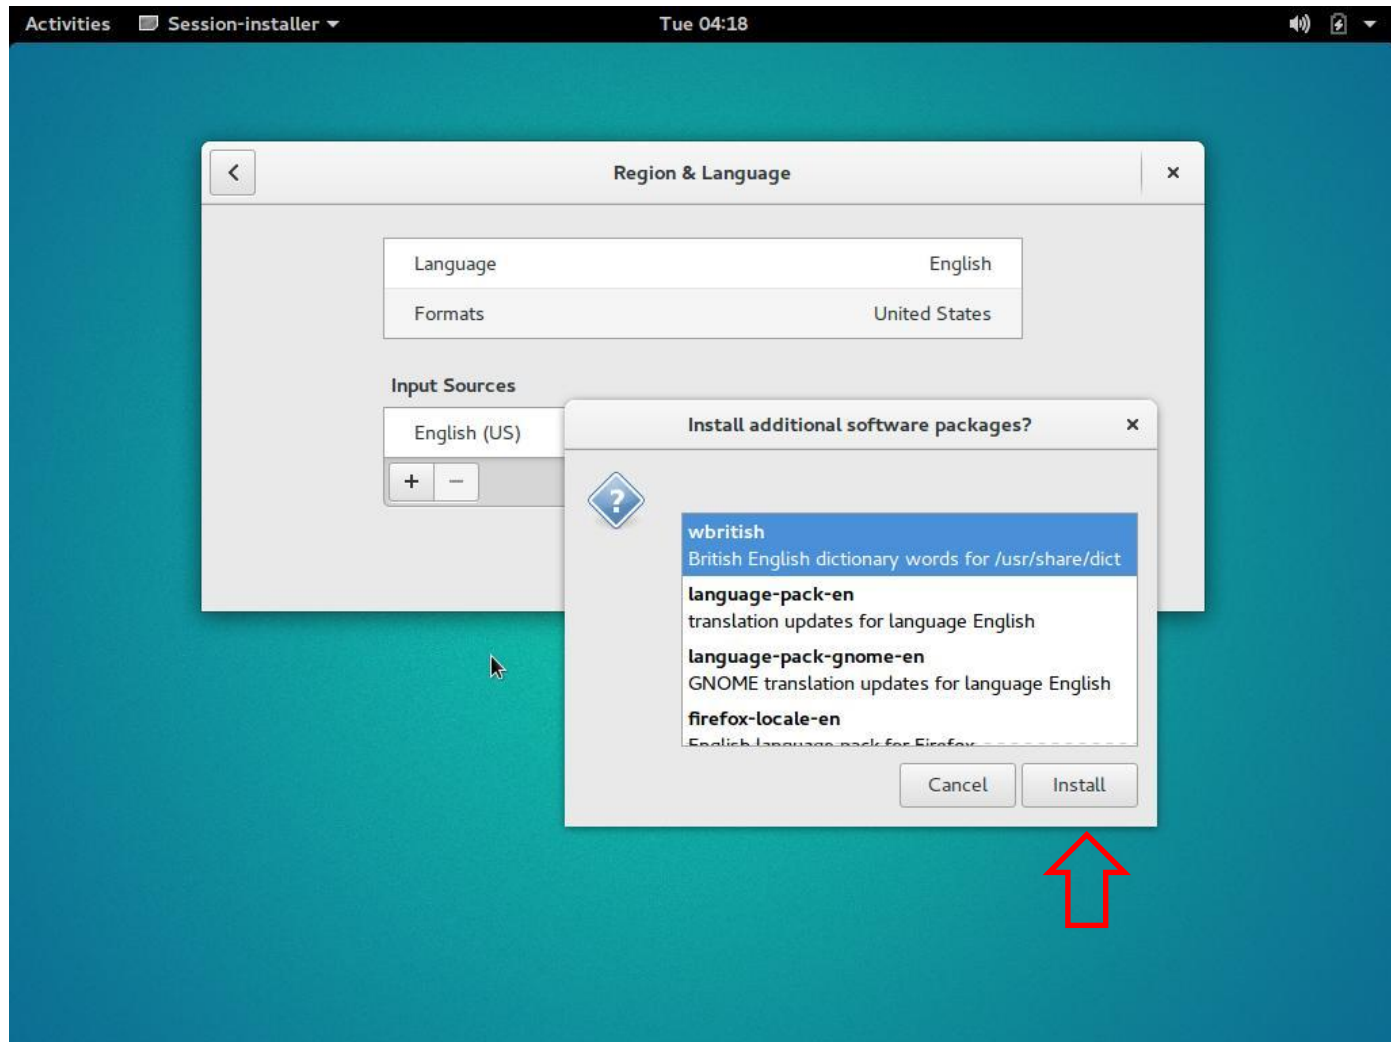
Select the input sources

9



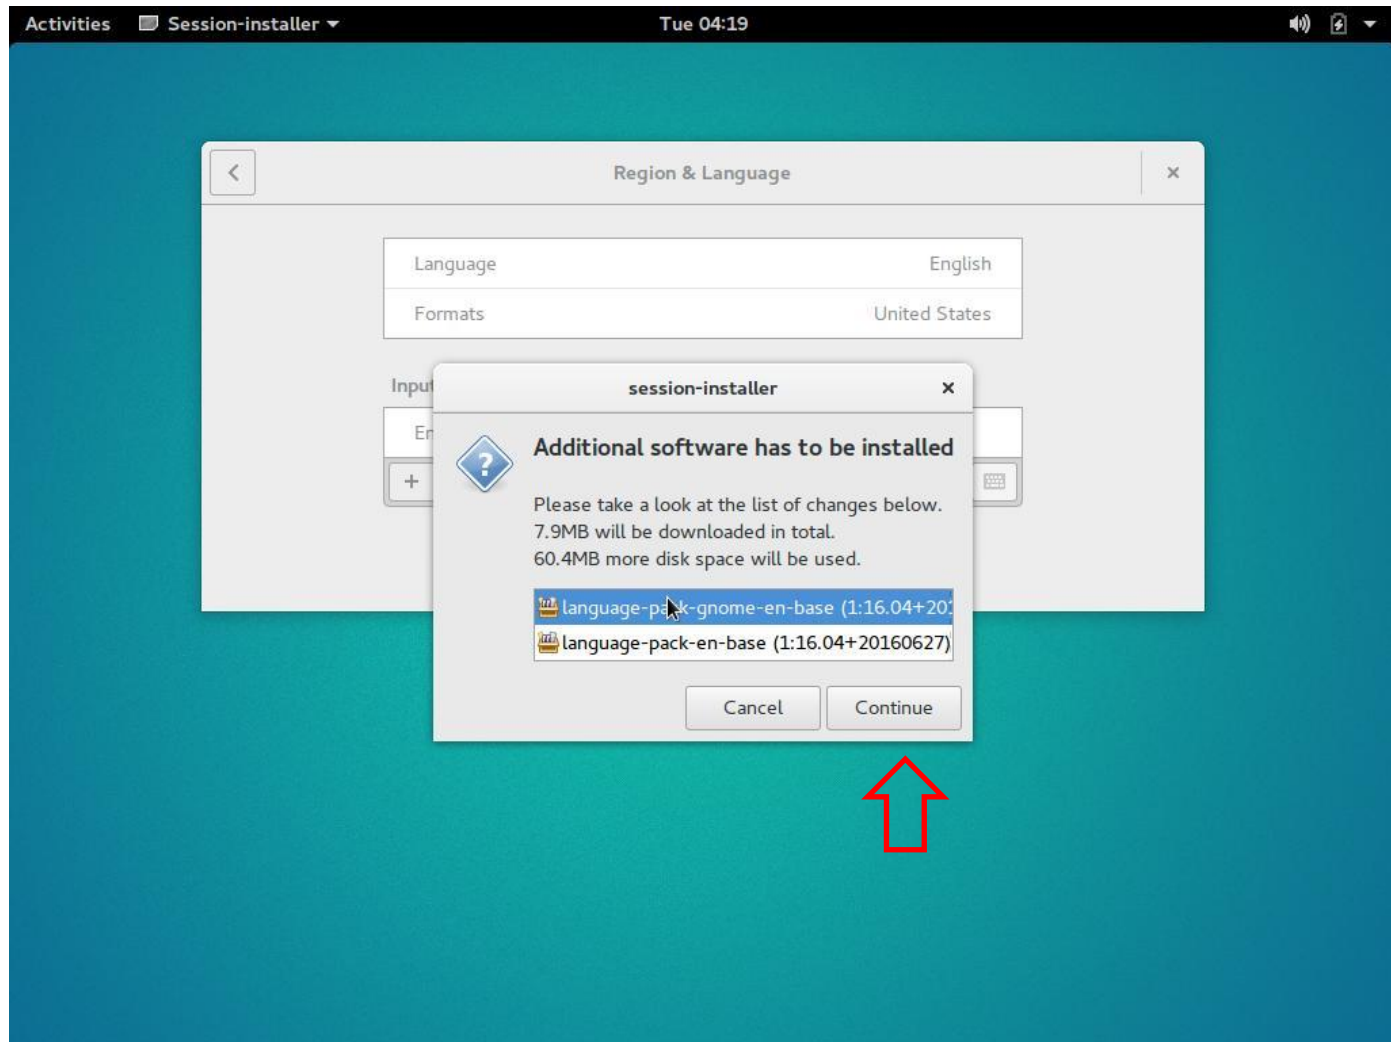
Install additional packages

10



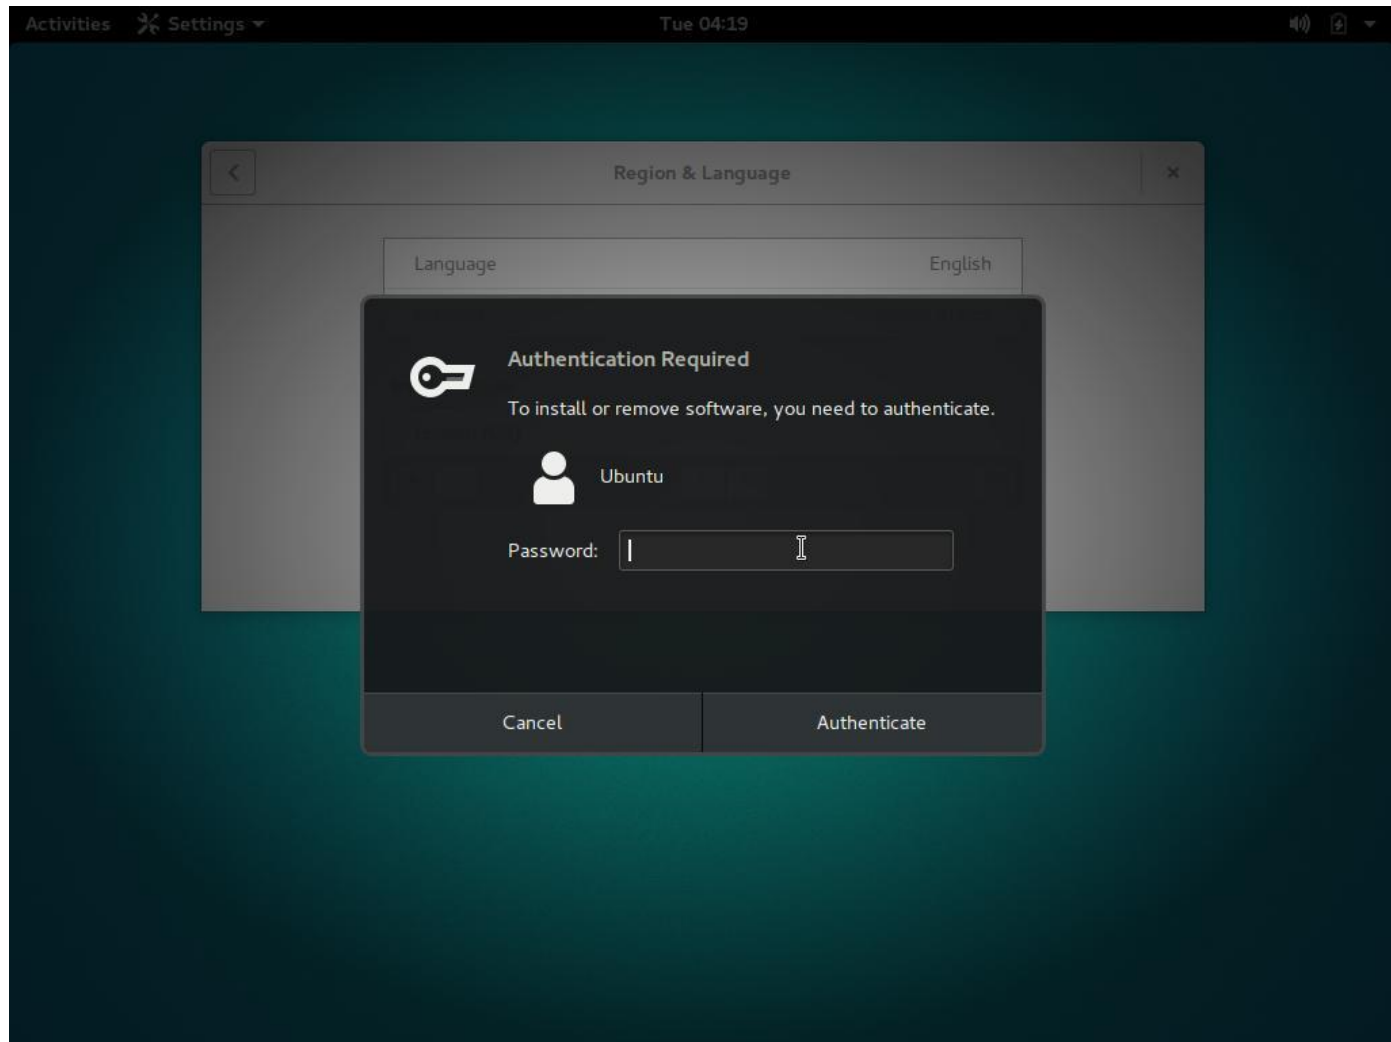
Confirm the installation

11



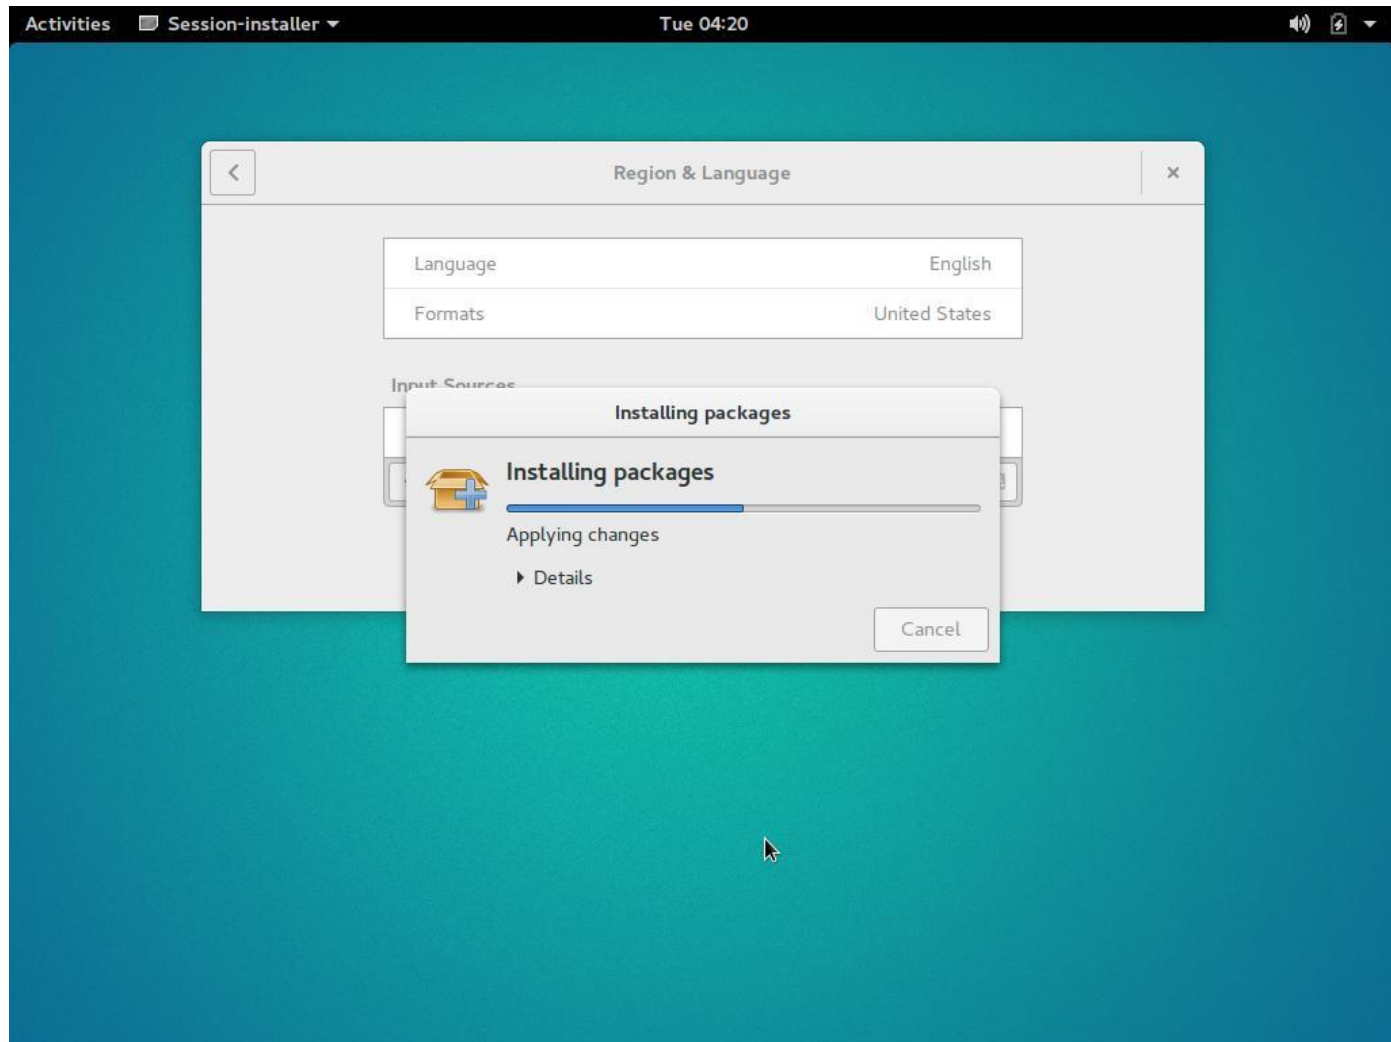
Type as a password: **password**

12



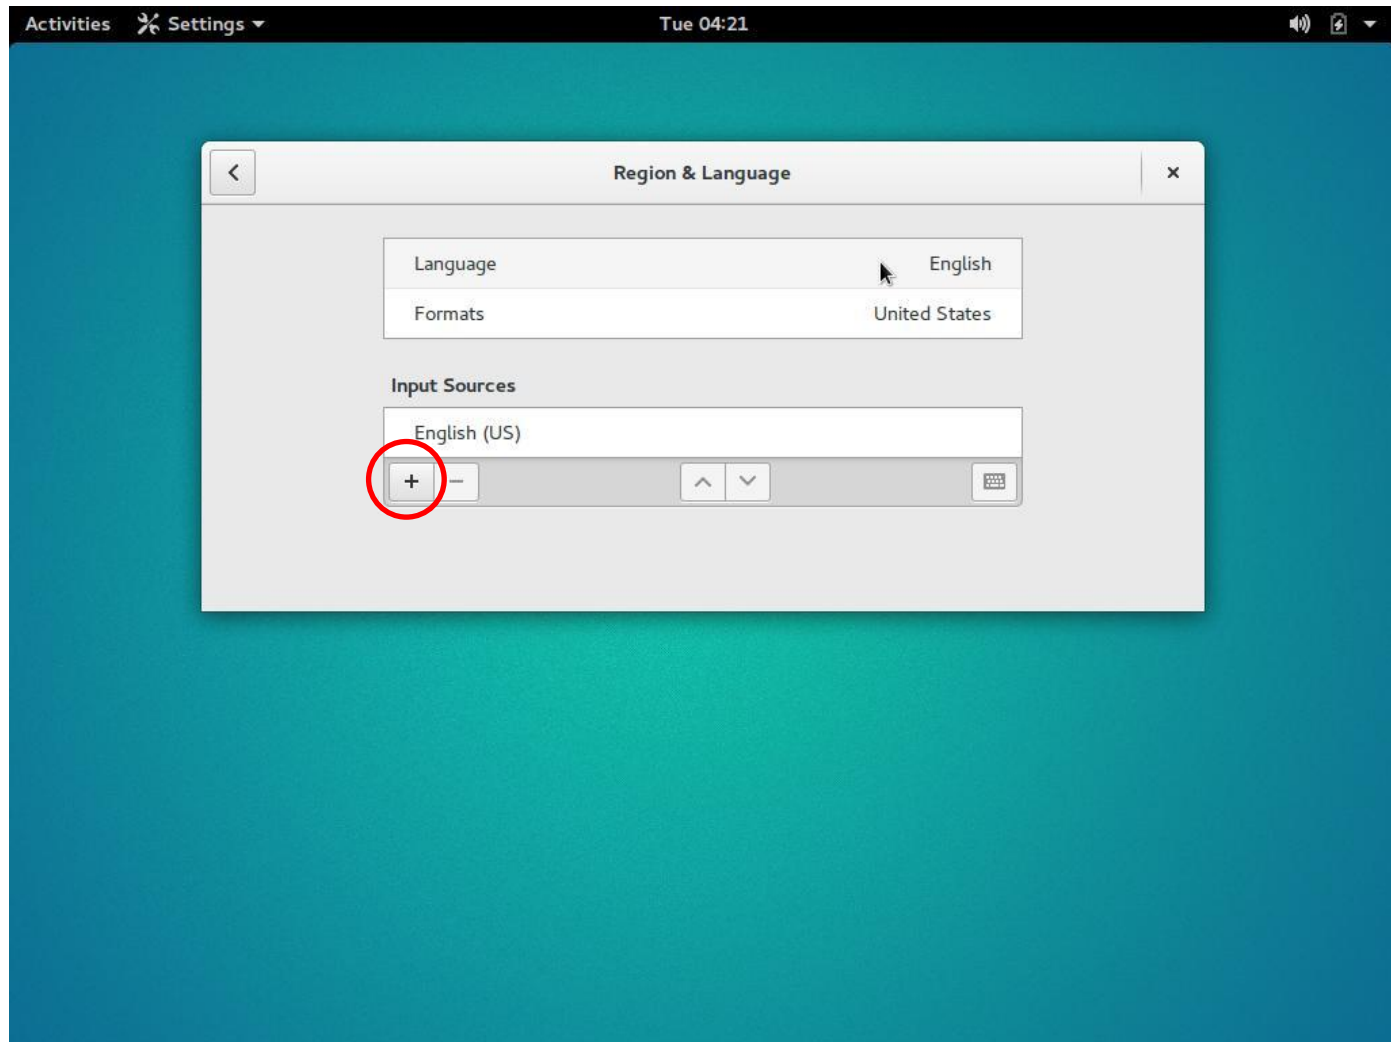
Wait a while

13



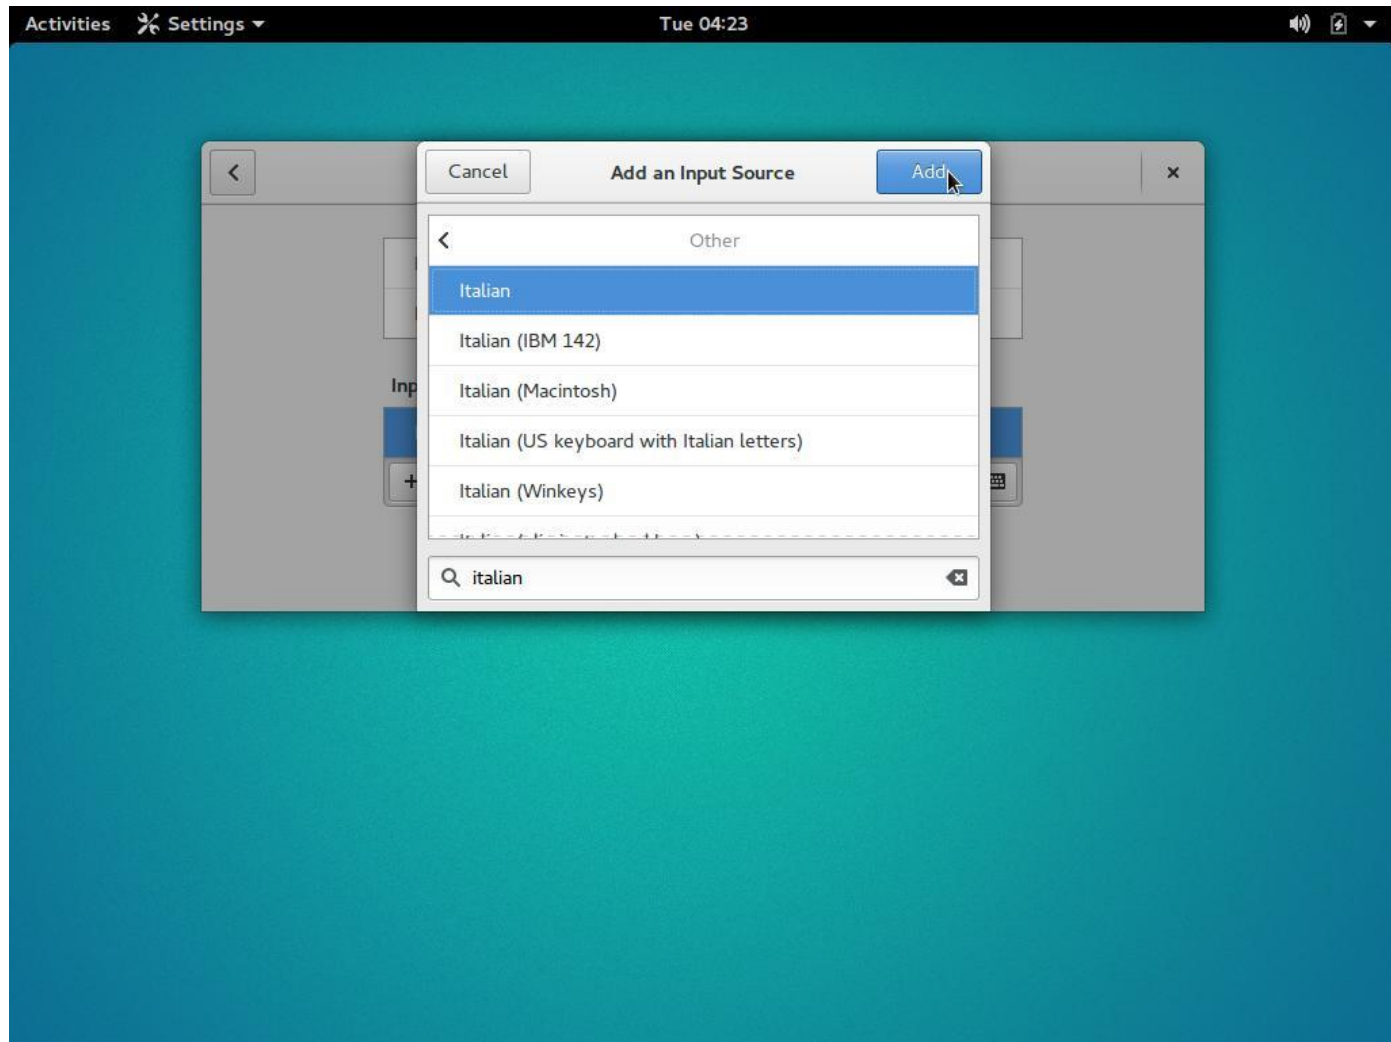
Add a new input source

14



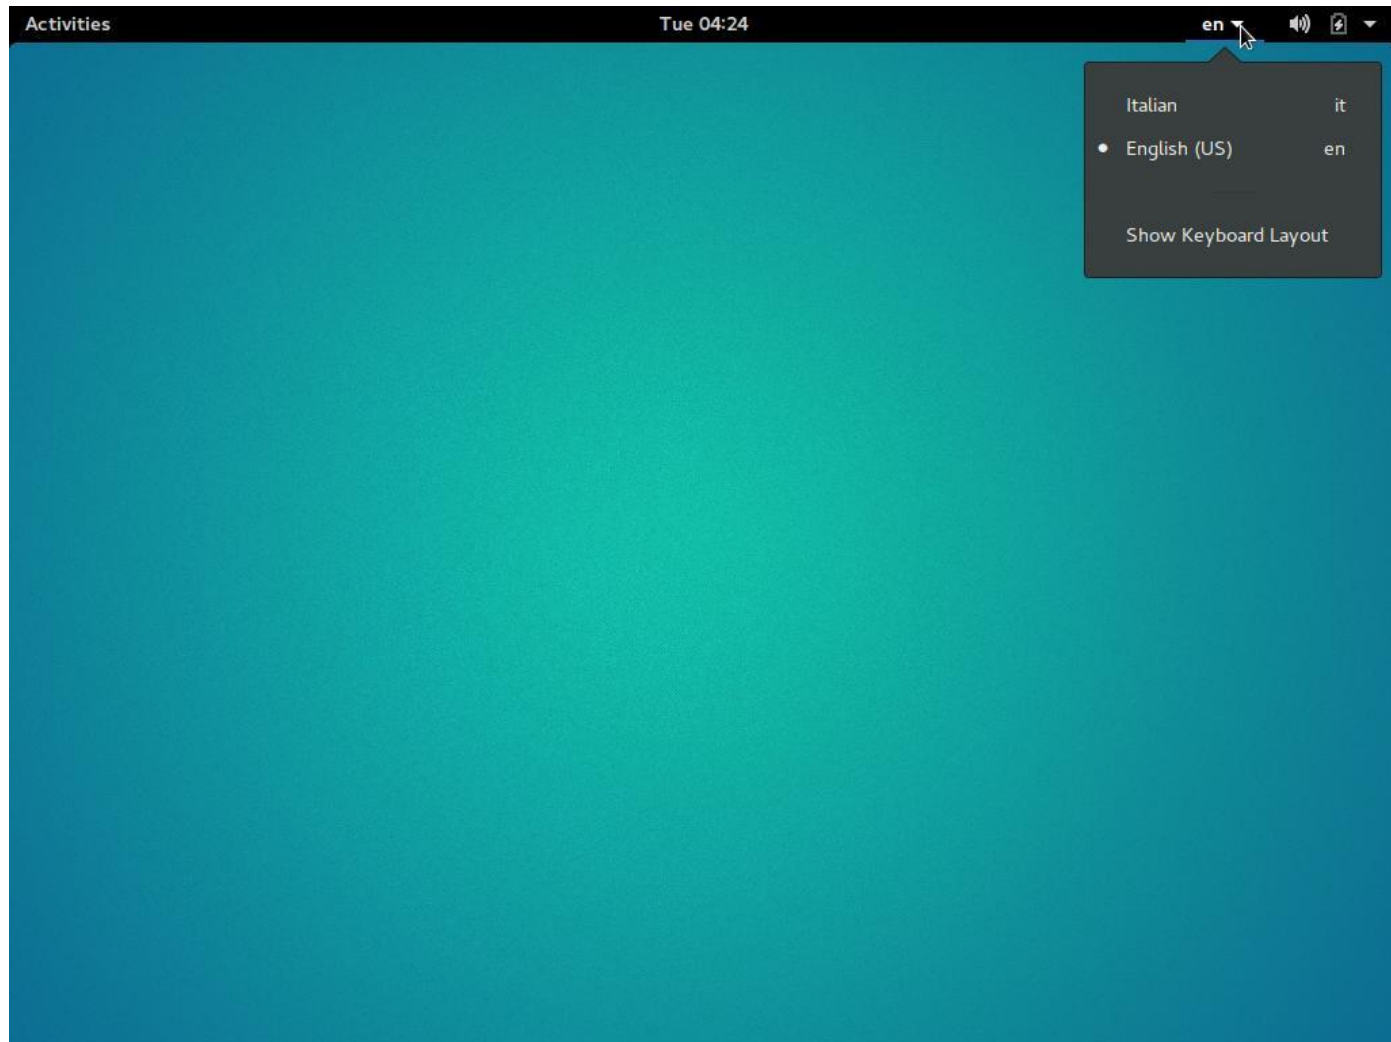
Type the keyboard language

15



Activate the keyboard layout

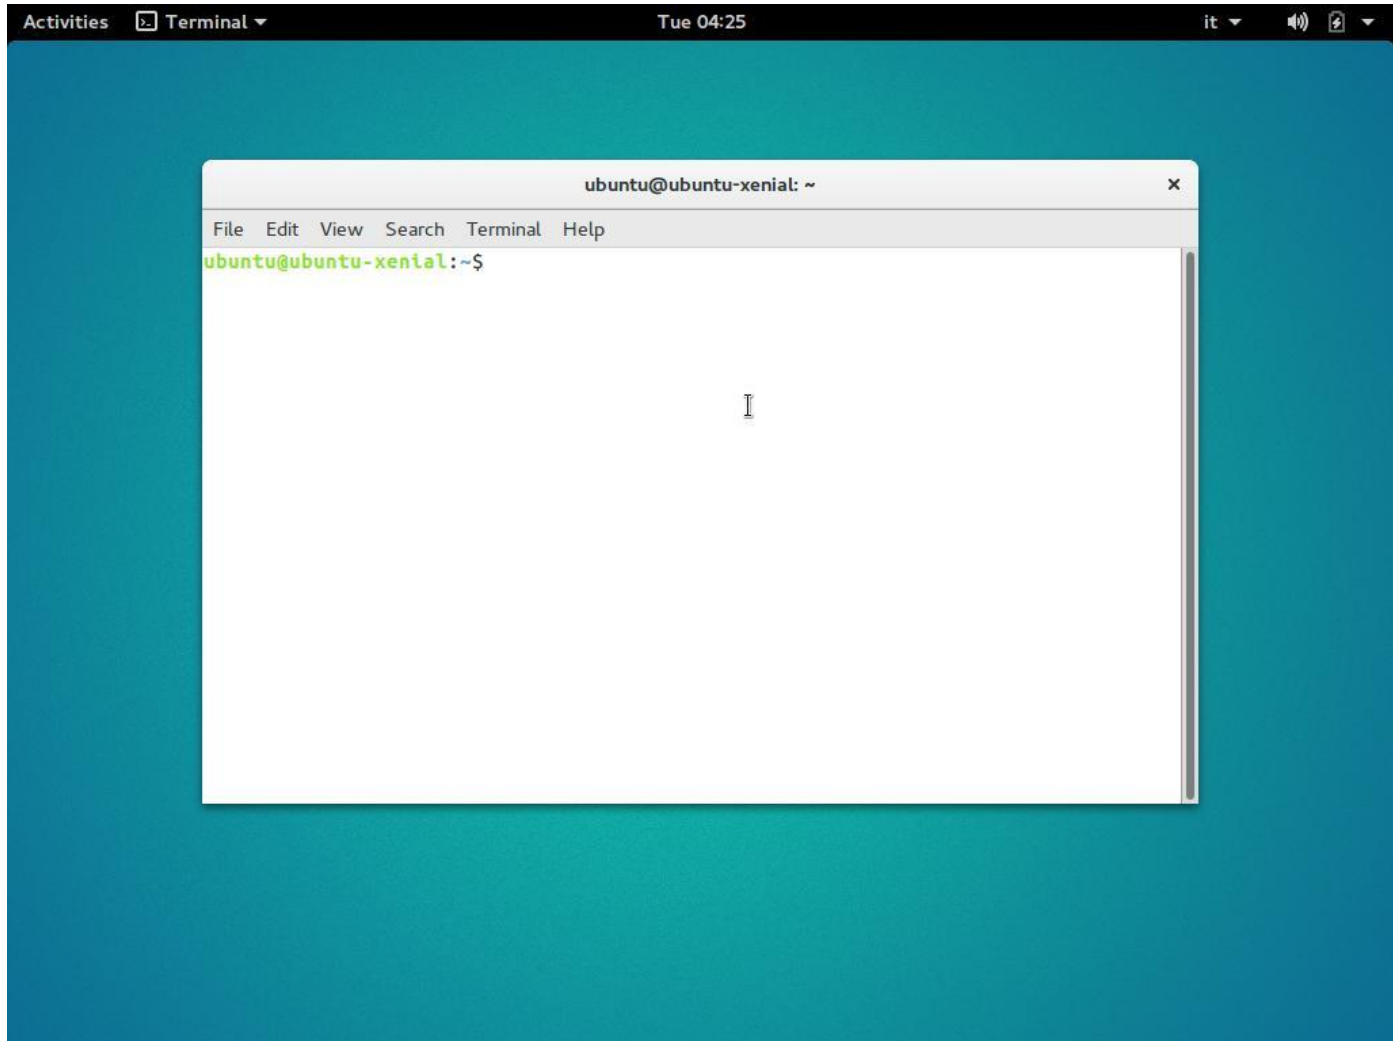
16



Open a terminal

17

Press simultaneously: **ctrl alt t**



Update the git repository

18

```
cd ~/PandA-bambu
```

```
git pull
```

❑ Change the directory by typing:

```
cd ~/PandA-bambu/documentation/tutorial_pact_2019/intro/first
```

❑ Edit the file C by typing:

```
gedit module.c
```

```
unsigned short icrc1(unsigned short crc, unsigned char onech)
{
    int i;
    unsigned short ans=(crc^onech << 8);
    for (i=0;i<8;i++) {
        if (ans & 0x8000)
            ans = (ans <<= 1) ^ 4129;
        else
            ans <<= 1;
    }
    return ans;
}
```

□ Run the script by typing:

```
./bambu.sh
```

```
Summary of resources:
- ASSIGN_SIGNED_FU: 2
- IIconvert_expr_FU: 1
- IUdata_converter_FU: 2
- MUX_GATE: 2
- UIdata_converter_FU: 3
- UUdata_converter_FU: 1
- bit_xor_expr_FU: 1
- constant_value: 5
- ge_expr_FU: 1
- lshift_expr_FU: 1
- ne_expr_FU: 1
- plus_expr_FU: 1
- read_cond_FU: 1
- register_SE: 2
- ui_bit_xor_expr_FU: 1
- ui_cond_expr_FU: 1
- ui_lshift_expr_FU: 1

Start reading vector          1's values from input file.

Value found for input crc: 0000000000001010
Value found for input onech: 00000010
Reading of vector values from input file completed. Simulation started.
Value found for output ex_return_port: 0010101001000010
  return_port = 10818    expected = 10818

Simulation ended after          10 cycles.

Simulation completed with success

- HLS_output//simulation/testbench_icrcl_minimal_interface_tb.v:441: Verilog $finish
1. Simulation completed with SUCCESS; Execution time 10 cycles;
  Total cycles           : 10 cycles
  Number of executions   : 1
  Average execution      : 10 cycles
```

- ❑ Testbench generated automatically
 - ▶ test_icrc1.xml

- ❑ Simulation and synthesis scripts generated automatically:
 - ▶ icrc1/simulate_icrc1_minimal_interface.sh
 - ▶ icrc1/synthesize_Synthesis_icrc1_minimal_interface.sh

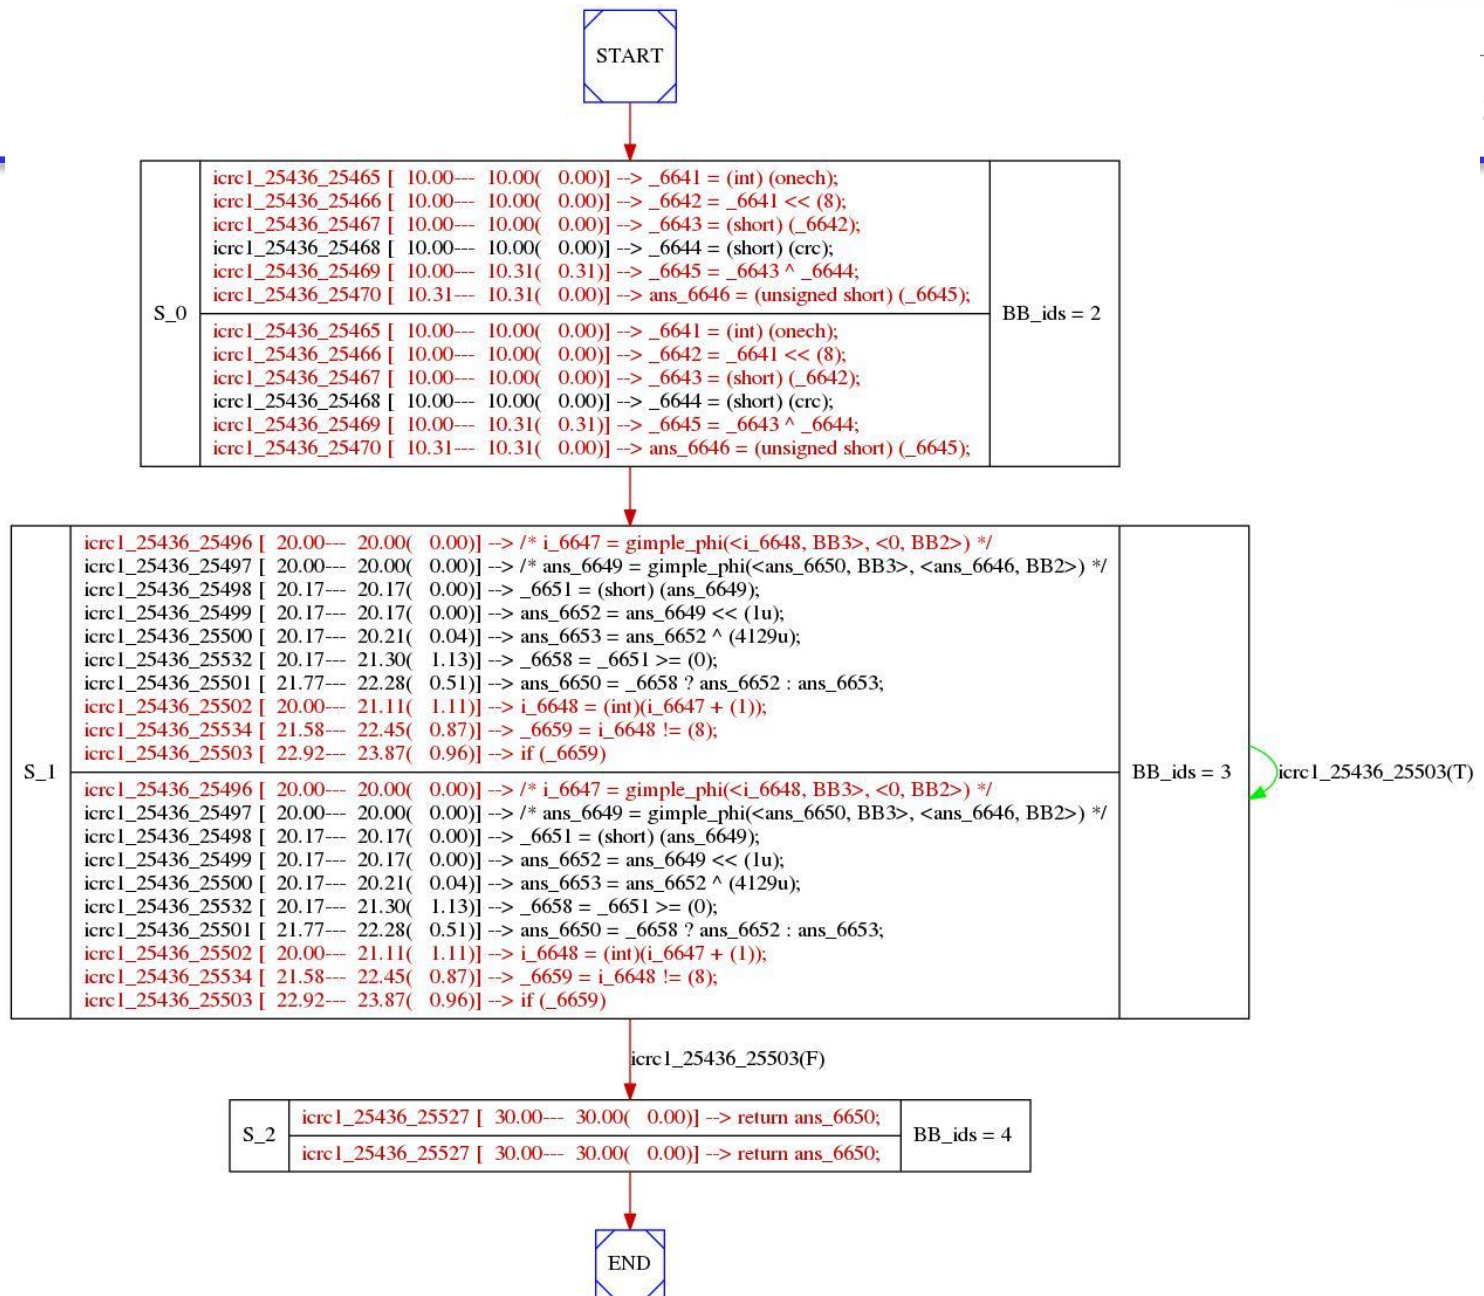
- ❑ Verilog file generated at the end of the HLS step:
 - ▶ icrc1/icrc1.v

- ❑ Change directory to icrc1:

```
cd icrc1
```

- ❑ Display the FSM:

```
xdot HLS_output/dot/icrc1/HLS_STGraph.dot
```

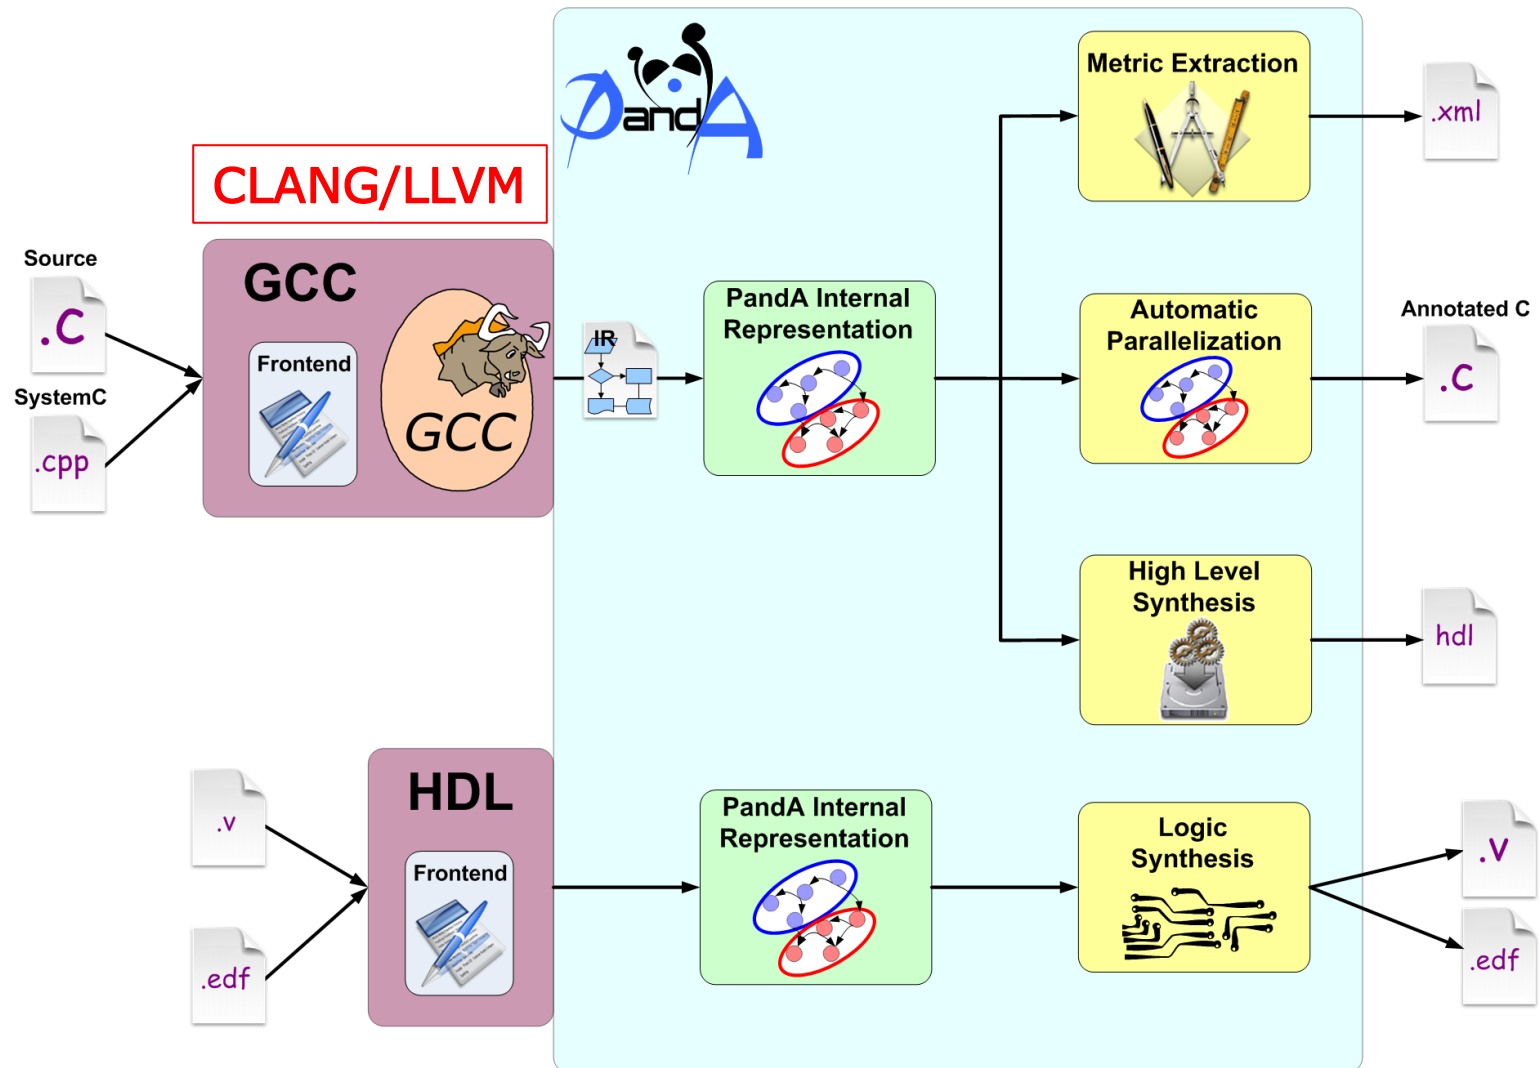


```
module icrc1_minimal_interface(clock, reset, start_port, crc, onech,
done_port, return_port);
    // IN
    input clock;
    input reset;
    input start_port;
    input [15:0] crc;
    input [7:0] onech;
    // OUT
    output done_port;
    output [15:0] return_port;
    // Component and signal declarations

    icrc1 icrc1_i0 (.done_port(done_port), .return_port(return_port),
.clock(clock), .reset(reset), .start_port(start_port), .crc(crc),
.onech(onech));

endmodule
```


- ❑ Panda framework development started on 2004 as a support research infrastructure for PoliMi in the context of ICODES – FP6-IST EU-funded project
 - ▶ Parsing and analysis of TLM 2.0 SystemC descriptions (gcc v.3.5)
- ❑ In the hArtes EU-funded project (2006-2010), it was used to
 - ▶ Analyzing generic C-based application annotated with pragmas (OpenMP)
 - ▶ Extracting parallel tasks
 - ▶ Estimating performance of embedded app
 - ▶ C-to-C rewriting
- ❑ Later, in Synaptic (2009-2013) and in Faster (2011-2014) EU-funded projects, logic- and high-level synthesis has been extended
 - ▶ Bambu (HLS tool) was first released in March 2012.
- ❑ ESA funded many research on code predictability analysis, performance analysis, and integration of HLS in model-based design flows.

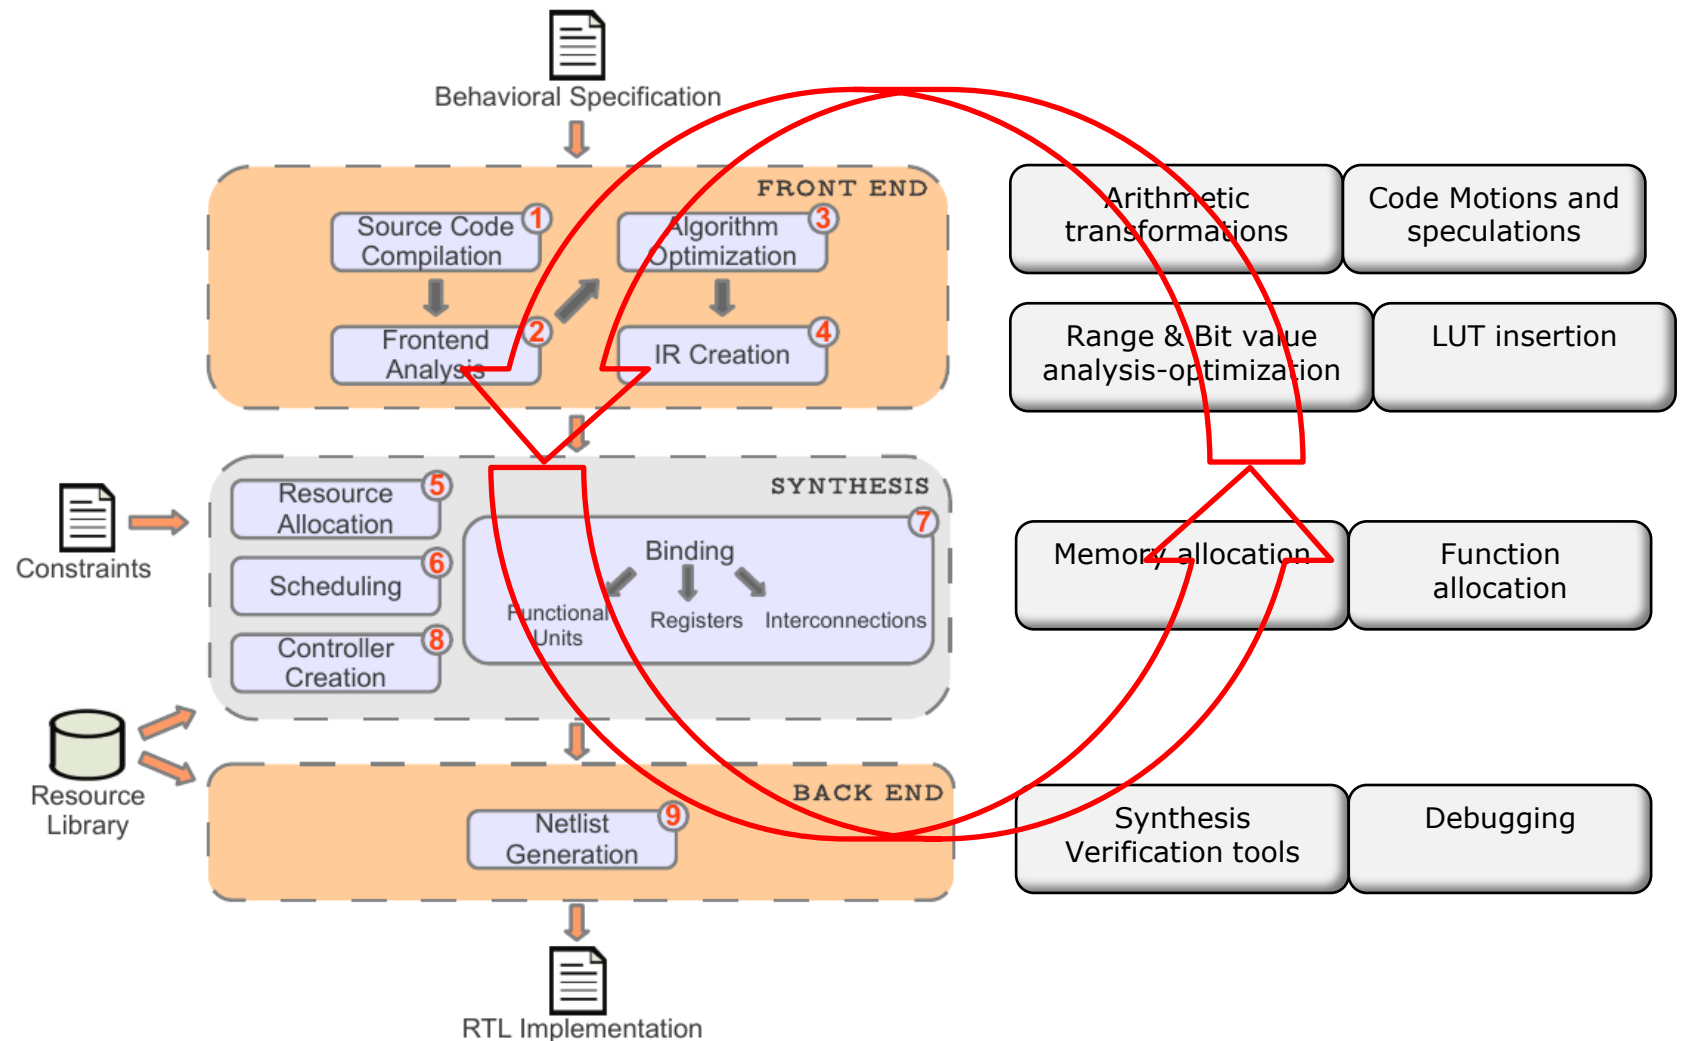




Bambu: an example of modern HLS tools

27

- ❑ HLS tool developed at Politecnico di Milano (Italy) within the Panda framework
 - ▶ Available under GPL v3 at
 - <http://panda.dei.polimi.it/>
 - <https://github.com/ferrandi/PandA-bambu>
- ❑ Example features
 - ▶ Front-end Input: interfacing with GCC/CLANG-LLVM for parsing C code
 - Complete support for ANSI C (except for recursion)
 - Support for pointers, user-defined data types, built-in C functions, etc..
 - Source code optimizations
 - may alias analysis, dead-code elimination, hoisting, loop optimizations, etc...
 - ▶ Target-aware synthesis
 - Characterization of the technology library based on target device
 - ▶ Verification
 - Integrated testbench generation and simulation
 - automated interaction with Iverilog, Verilator, Xilinx Isim, Xilinx Xsim, Mentor Modelsim
 - ▶ Back-end: Automated interaction with commercial synthesis tools
 - FPGA: Xilinx ISE, Xilinx Vivado, Altera Quartus, Lattice Diamond
 - ASIC: Synopsis Design Compiler



Modular Framework based on the specialization of HLS_step

❑ Minimal command

- ▶ `$ bambu filename.c`

❑ Controlling the clock period (100Mhz)

- ▶ `$ bambu filename.c --clock-period=10`

❑ Select the device

- ▶ `$ bambu filename.c -device-name=xc7z020,-1,clg484,VVD`



Subset of synthesizable C (1)

30

- ❑ We support what standard compilers accept as input (CLANG/LLVM and GCC)
- ❑ Supported features:
 - ▶ Expressions of any kind: arithmetic, logical, bitwise, relational, conditional, comma-based expressions.
 - ▶ Types: integers, single- and double-precision floating point, `_Bool` and `Complex`, struct-or-union, bitfields, enum, typedef, pointers and arrays, type qualifiers.
 - ▶ Variable declarations, initialization, storage-specifiers
 - ▶ Functions definition and declaration, extern or static, pointer to functions, parameters passed by copy or reference, tail recursive functions.
 - ▶ Statements and blocks: labeled (`case`), compound, expression, selection (`if`, `switch`), iteration (`while`, `do`, `for`), jump (`goto`, `continue`, `break`, `return`)
 - ▶ All preprocessor directives
 - ▶ Unaligned memory accesses and dynamic pointers resolution
 - ▶ GCC vectorization



Subset not supported

31

- ❑ struct returned by copy
- ❑ Non-tailing recursive functions

- ❑ Search and insertion in a binary tree
- ❑ Change the directory by typing:

```
cd ~/PandA-bambu/documentation/tutorial_pact_2019/intro/second
```

- ❑ Edit the file C by typing:

```
gedit module.c
```

- ▶ Two data structures: stack and binary tree
- ▶ Static memory allocators
- ▶ Tail recursive functions
- ▶ Use of pointer to pointers (some HLSs have problems)

- ❑ `assert, puts, putchar, read, open, close, write, printf, exit, abort`
- ❑ **libc functions:** `bswap32, memcmp, memcpy, memmove, memset, malloc, free, memalign, alloca, with_align, calloc, bcopy, bzero, memchr, memcpy, memchr, rawmemchr, strcpy, strncpy, strcasecmp, strcasestr, strcat, strchr, strchrnul, strcmp, strcpy, strcspn, strdup, strlen, strncasecmp, strncat, strncmp, strncpy, strndup, strnlen, strpbrk, strchr, strsep, strspn, strstr, strtok`
- ❑ **libm functions:** `acos, acosh, asin, asinh, atan, atan2, atanh, cbrt, ceil, cexp, copysign, cos, cosh, drem, erf, exp, exp10, expm1, fabs, fdim, finite, floor, lfloor, fma, fmax, fmin, fmod, fpclassify, frexp, gamma, lgamma, tgamma, hypot, ilogb, infinity, isinf, isnan, j0, j1, jn, ldexp, log, log2, log10, log1p, modf, nan, nearbyint, nextafter, pow, pow10, remainder, remquo, rint, lrint, llrint, round, lround, llround, scalb, scalbln, scalbn, signbit, significand, sin, sincos, sinh, sqrt, tan, tanh, trunc.`

- ❑ Crypto core built composing user defined libraries
- ❑ Change the directory by typing:

```
cd ~/PandA-bambu/documentation/tutorial_pact_2019/intro/third
```

- ❑ Run the script by typing:

```
./multi.sh
```

- ❑ `tree-panda-gcc` could be used to create a custom library that could be deployed by `bambu` by passing `-l` and `-L` options.

Autotools based project described in directory
[examples/crypto_designs/multi-keccak](#)

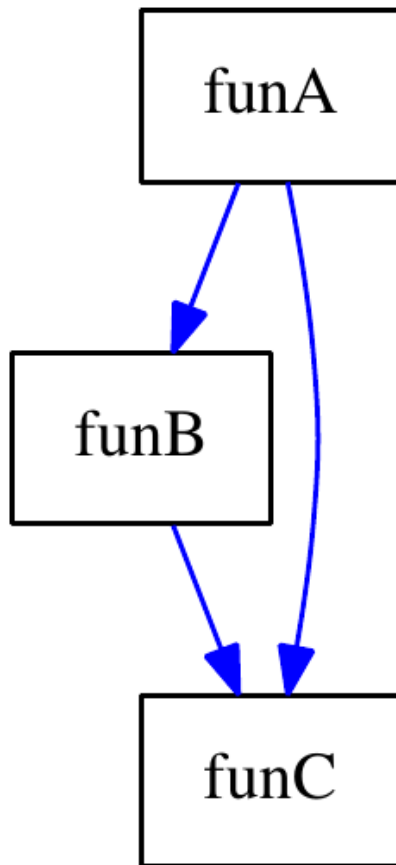
- ❑ `bambu` uses the text-based IR and the builtin linker to build a single in-memory representation and perform HLS

▶ `$ bambu file1.c file2.c --top-fname=fname`

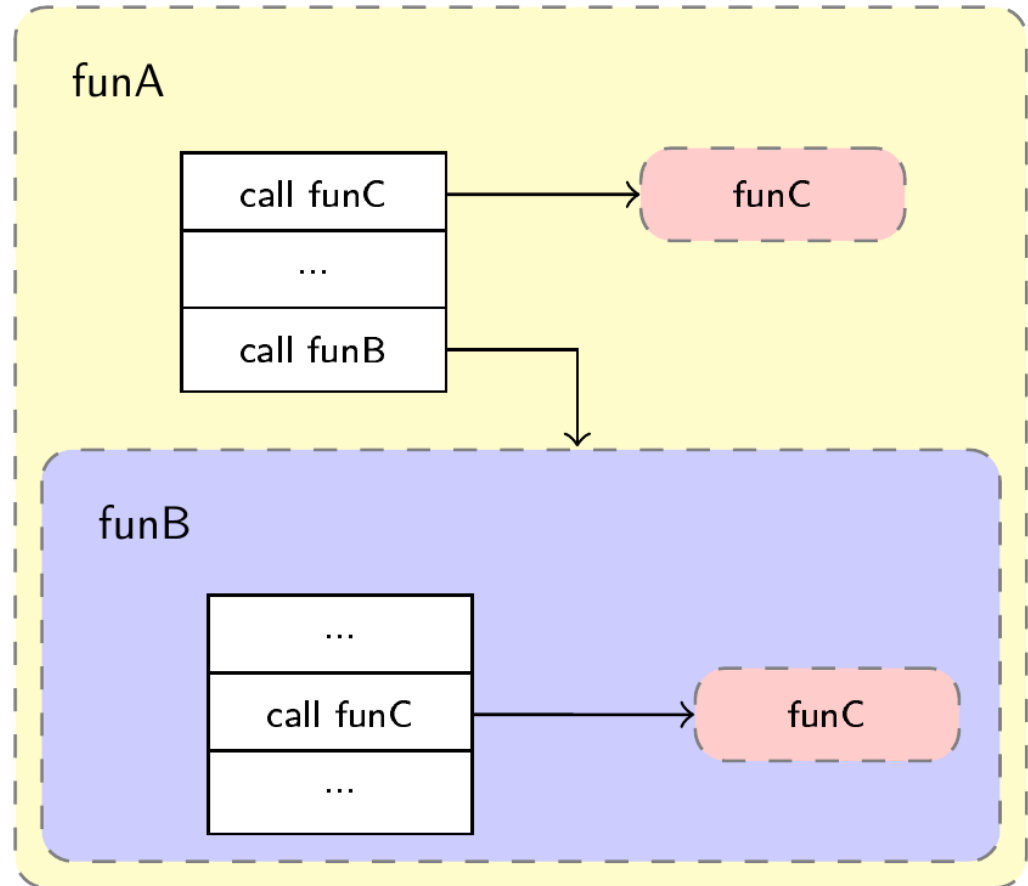
- ❑ In case the option `--top-fname` is not used, `bambu` automatically identifies which function can act as the top one.
- ❑ With clang link time optimization is used.

- ❑ one component per function
 - ▶ function interface
 - ▶ start and done
 - ▶ parameter passing
 - wires
 - memory interaction
 - none (ap_none), acknowledge (ap_ack), valid (ap_vld), ovalid (ap_ovld), handshake (ap_hs), fifo (ap_fifo) and array (ap_memory)
- ❑ hierarchy based on call graph
 - ▶ no-recursion
 - ▶ proxy

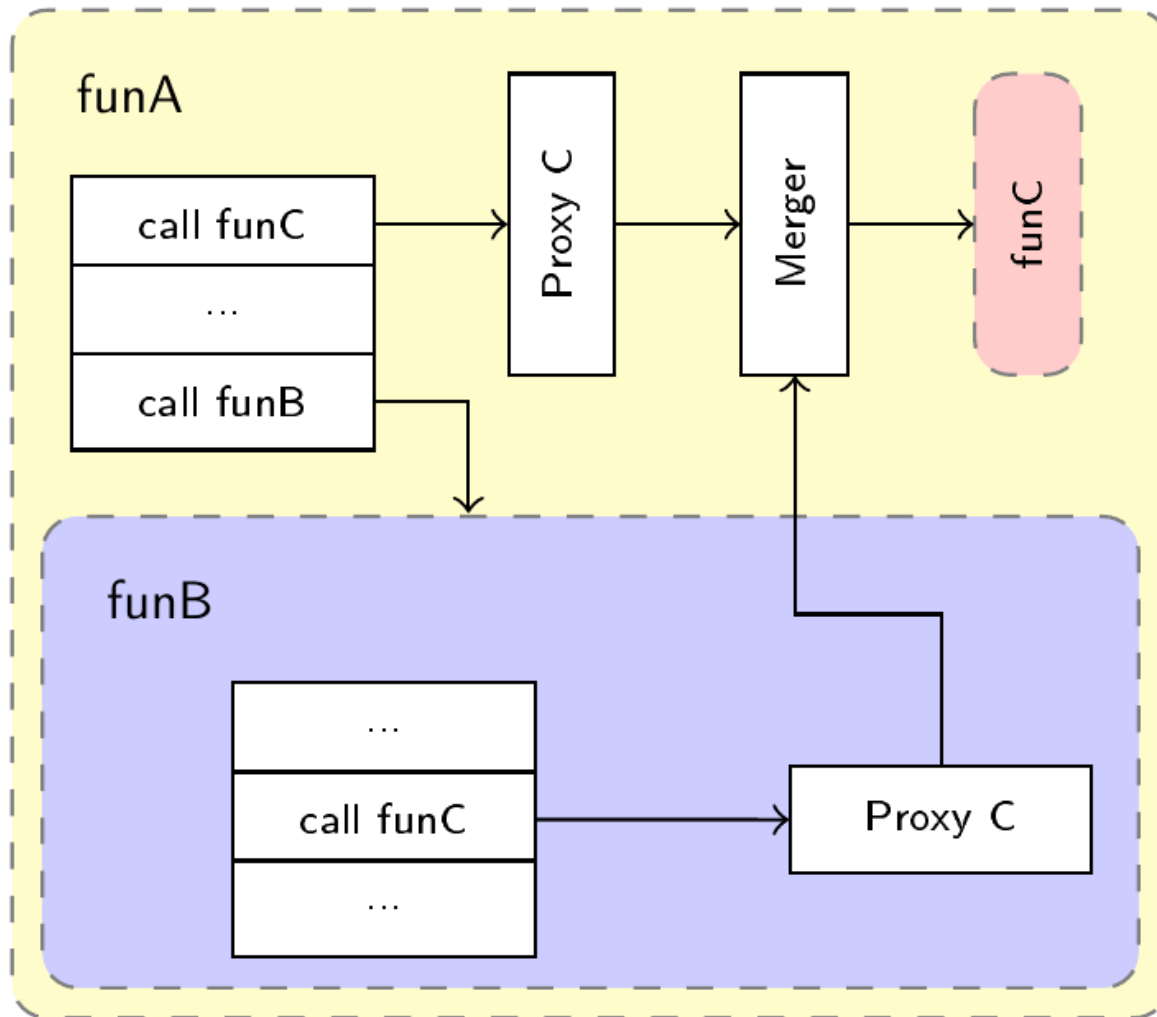
- ❑ One component per function



call graph



RTL hierarchy

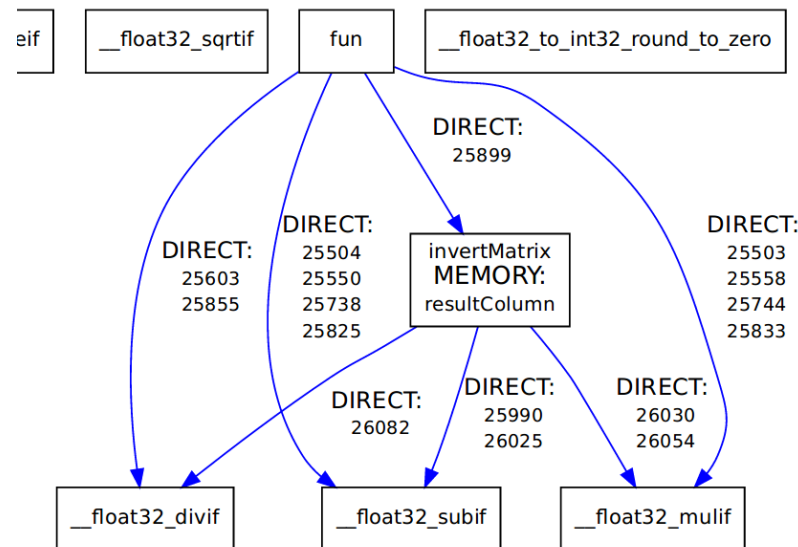


❑ Lu-decomposition with single precision floating point arithmetic

```
cd ~/PandA-bambu/documentation/tutorial_pact_2019/intro/fourth
```

❑ Run the script by typing:

```
./bambu.sh
```



- ❑ bambu assumes that all global variables could be accessed by external CPU/accelerators
- ❑ It is possible to change this default with option:
 - ▶ `--do-not-expose-globals`
 - ▶ All global variables are considered local to the compilation units as if they are declared `static`.

❑ Integration of existing IPs written in Verilog that receives structs passed by pointers

```
cd ~/PandA-bambu/documentation/tutorial_pact_2019/intro/fifth
```

```
#include "module_lib.h"

void my_ip(uint8_t command, uint32_t param1, uint32_t param2) {
    static module1_output_t module1_output;
    static module2_output_t module2_output;
    switch(command) {
        case 0:
            module1(param1, param2 >> 16, &module1_output);
            break;
        case 1:
            module2(param1, &module2_output);
            break;
        case 2:
            printer1(module1_output.output1, module1_output.output2, module1_output.output3,
module1_output.output4);
            break;
        case 3:
            printer2(module2_output.output1, module2_output.output2, module2_output.output3);
            break;
        default:
            break;
    }
}
```

- ❑ Function mapped on IPs has to be declared as extern:
 - ▶ `extern void module1(uint32_t input1, uint16_t input2, module1_output_t *outputs);`
- ❑ C code has to be passed with the following option
 - ▶ `--C-no-parse=module1.c,...`
- ❑ Binding between function **module1** and component **module1** has to be specified with a XML file and passed as an option to bambu
 - ▶ `$ bambu ... module_lib.xml`
- ❑ Check these examples:
 - ▶ `examples/IP_integration`
 - ▶ `examples/breakout`
 - ▶ `examples/pong`
 - ▶ `examples/led_example`

- ❑ Parametric quick sort

```
cd ~/PandA-bambu/documentation/tutorial_pact_2019/intro/sixth
```

- ❑ Quick sort parametric with respect to the comparison function
- ❑ Run the script and check the results

```
int laplacian(char *, char *, int, int);
int make_inverse_image(char *, char *, int, int);
int sharpen(char *, char *, int, int);
int sobel(char *, char *, int, int);

int (*pipeline[MAX_DEPTH])(char *, char *, int, int);

void UserApp(char *in, char *out, int x_size, int y_size) {
    // ...
    // Pipeline configuration using function pointers
    add_filter(0, make_inverse_image);
    add_filter(1, sharpen);
    // ...
    // execute is synthesized in hardware
    execute(in, out, x_size, y_size);
}

void execute(char *in, char *out, int x_size, int y_size) {
    int i = 0;
    for (i = 0; i < MAX_PIPELINE_DEPTH; i++) {
        if (pipeline[i] == 0) break;
        // here other hw accelerator are called
        // using function pointers
        int res = pipeline[i](in, out, x_size, y_size);
        if (res != 0) return;
        swap(in, out);
    }
    move_if_odd(i, in, out);
}
```

Adding a memory mapped interface to filters

45

Accelerator
base address



Control register

Input register: in

Input register: out

Input register:
x size

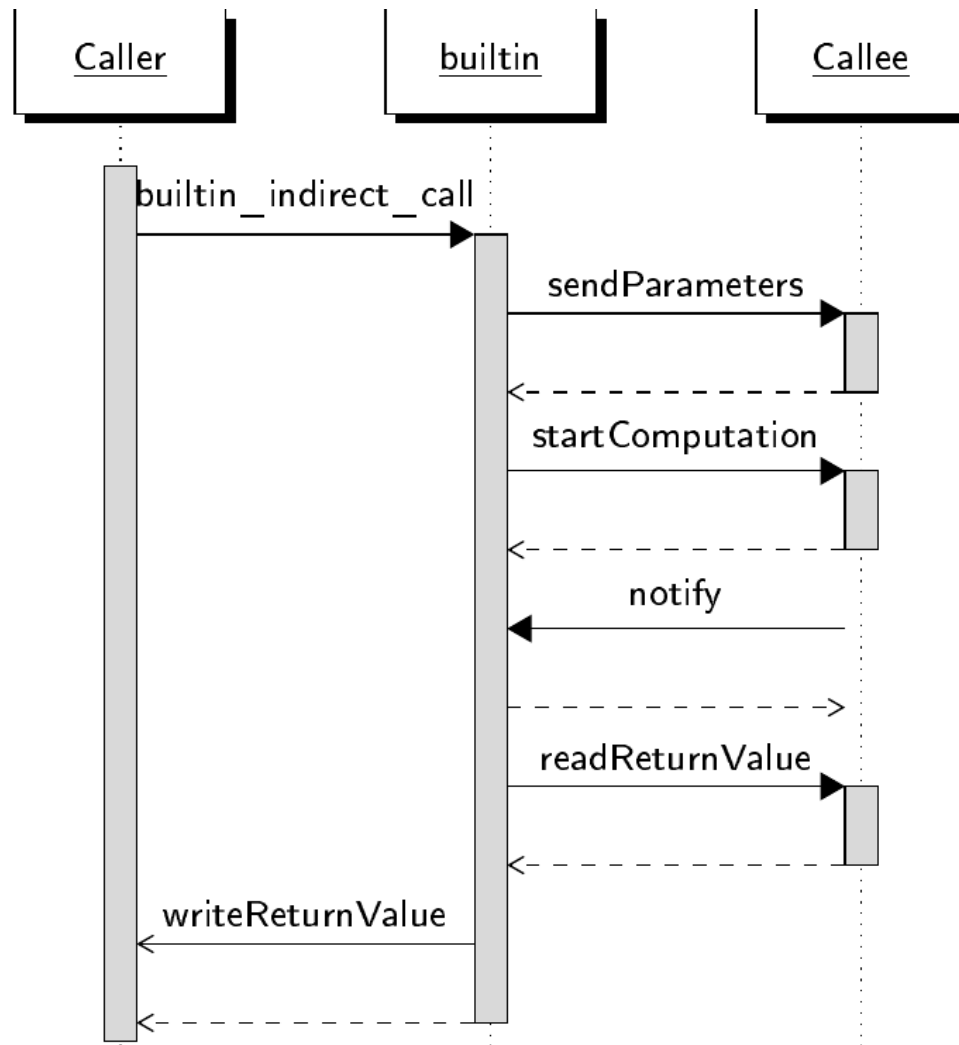
Input register:
y size

Output register

```
void execute(char *in, char *out, int x_size, int y_size) {
    int i = 0;
    for (i = 0; i < MAX_PIPELINE_DEPTH; i++) {
        if (pipeline[i] == 0) break;
        // here other hw accelerator are called
        // using function pointers
        __builtin_indirect_call(
            pipeline[i], 1, in, out, x_size, y_size, &res);
        if (res != 0) return;
        swap(in, out);
    }
    move_if_odd(i, in, out);
}
```

Sequence diagram for function indirect call

47



Call mechanism complexity: $\#cycles = Wl(Np+1) + lrs(Wl+RI)$

- ❑ Support of C++ is ongoing:
 - ▶ templates
 - ▶ C++11 and beyond
 - ▶ ac_types from Mentor Graphics could be used
 - ▶ ap_types from Xilinx support by wrapping ac_types

```
#include <algorithm>
int gcd(int x, int y )
{
    if( x < y )
        std::swap( x, y );

    while( y > 0 )
    {
        int f = x % y;
        x = y;
        y = f;
    }
    return x;
}
```

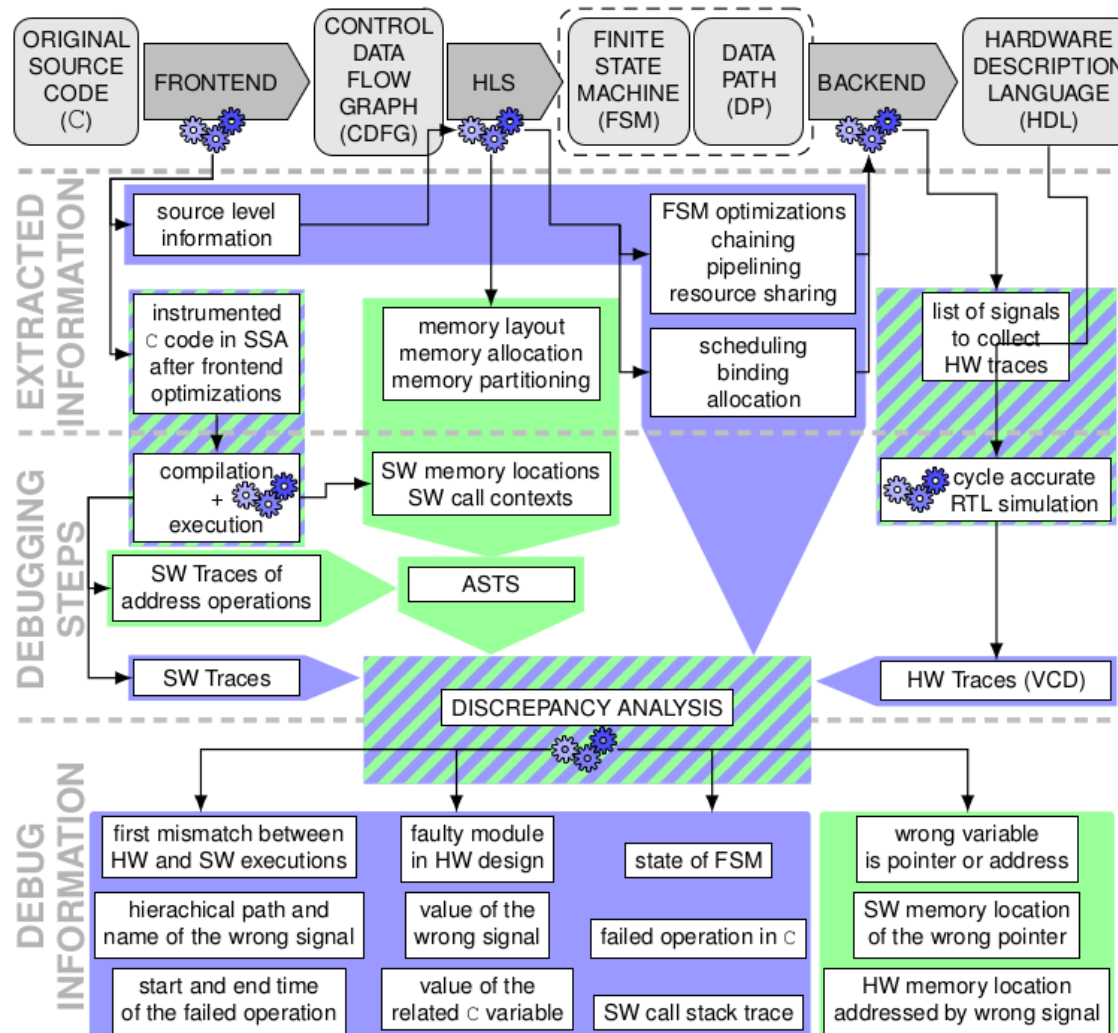

- * euclid.f (FORTRAN 77)
- * Find greatest common divisor using the Euclidean algorithm

```
FUNCTION NGCD (NA, NB)
  IA = NA
  IB = NB
1  IF (IB.NE.0) THEN
    ITEMP = IA
    IA = IB
    IB = MOD (ITEMP, IB)
    GOTO 1
  END IF
  NGCD = IA
  RETURN
END
```

By default parameters are passed
by reference

Discrepancy Analysis Debug Flow

50





Please don't shoot the Panda player
and visit
<http://panda.dei.polimi.it>