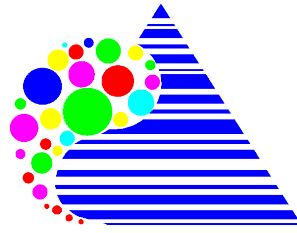


# The 28th International Conference on Parallel Architectures and Compilation Techniques

Seattle, WA, USA, September 21st, 2019



 POLITECNICO DI MILANO



PACT19 Tutorial

## Exploiting Vectorization in High Level Synthesis of Nested Irregular Loops

**Marco Lattuada**

Politecnico di Milano  
Dipartimento di Elettronica, Informazione e Bioingegneria  
[marco.lattuada@polimi.it](mailto:marco.lattuada@polimi.it)

September 21, 2019

- ❑ Introduction
- ❑ Outer Loops Vectorization
- ❑ Proposed Design Flow
- ❑ Examples
- ❑ Experimental Results
- ❑ Conclusions







# DoAll Loops and their Optimizations

3

## □ DoAll loops

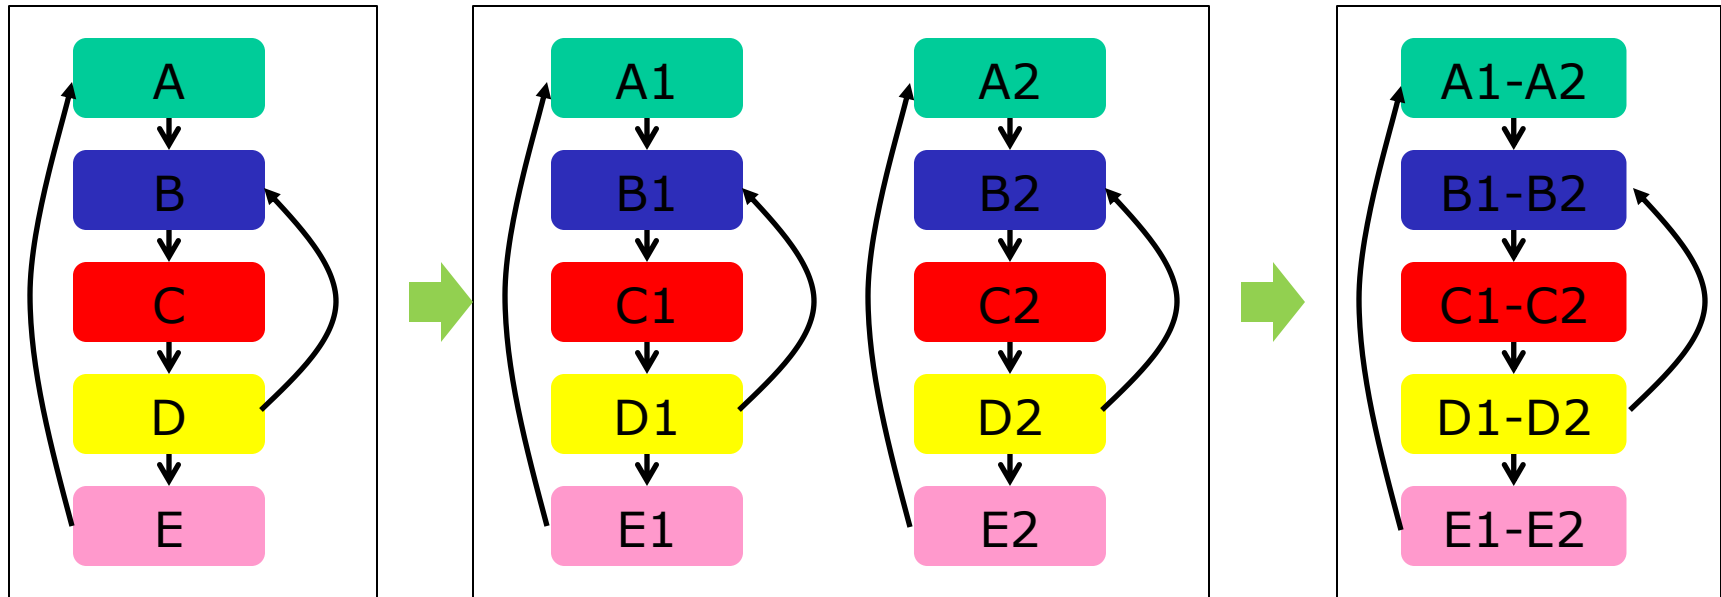
- ▶ All the iterations can be executed in parallel
- ▶ Number of iterations can be unknown at compile time

## □ Optimizations



# Outer Loops Vectorization

4




- ❑ It can be applied even if inner loops cannot be parallelized
- ❑ Operations of outer loops are executed simultaneously
- ❑ Operations of inner loops are parallelized only if they belong to different instances of outer loop



# Outer Loops Vectorization

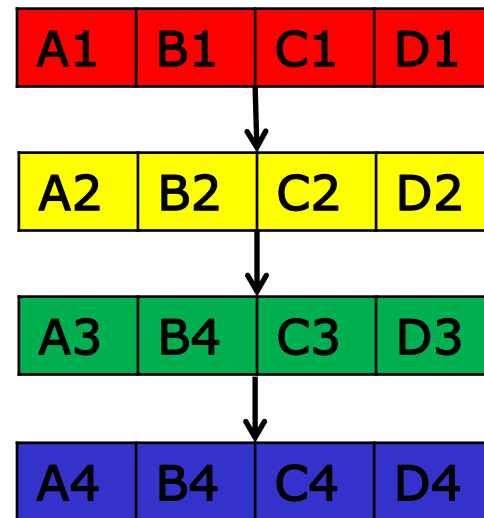
5

```
for(r=0; r<N; r++)  
  for(c=0; c<N; c++)  
    if(c > 0)  
      out[r][c]=out[r][c-1]+in[r][c];  
    else  
      out[r][c]=in[r][c];
```



Inner Loop  
Vectorization

A1	A2	A3	A4
B1	B2	B3	B4
C1	C2	C3	C4
D1	D2	D4	D4



Outer Loop  
Vectorization



- ❑ Outer Loop must be **DoAll loop**
- ❑ Number of iterations must be **multiple of the degree of parallelism**
  - ▶ Ad-hoc management of last iterations would be required
- ❑ Nested loops cannot contain code in mutual exclusions

```
if(condition) {...} else {...}
```

is transformed into

```
if(condition) {...} if(!condition) {...}
```
- ❑ Number of iterations of nested loops **can depend** on a value computed in **outer loop**
  - ▶ Nested loops can be executed arbitrary number of times

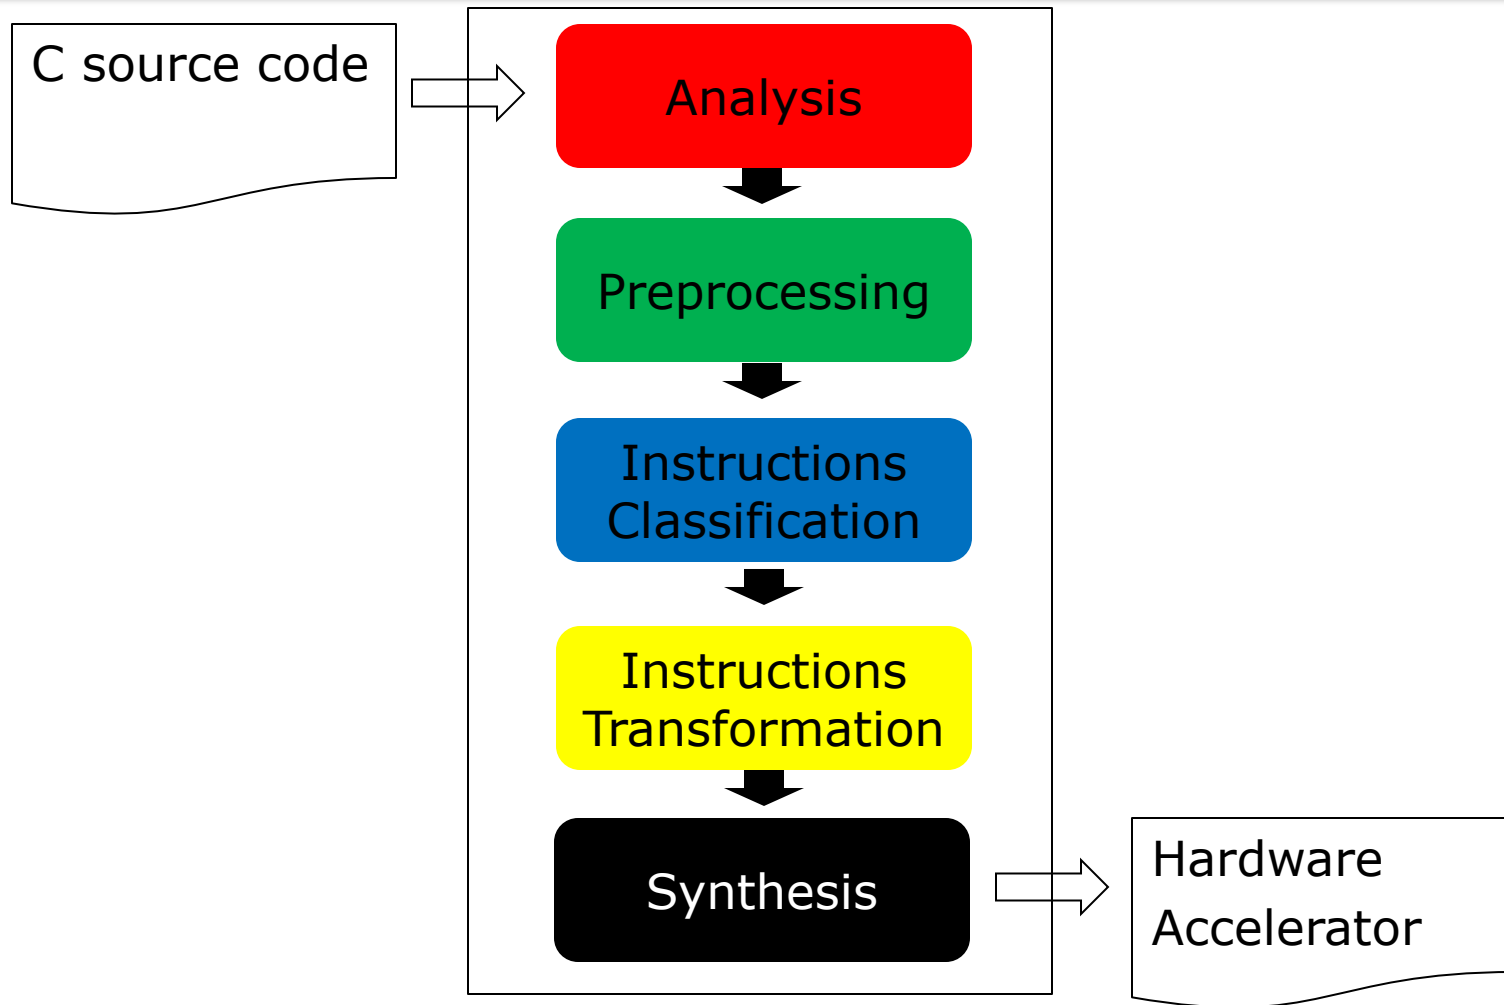


- ❑ Vector functional units can be synthesized for all the computation operations
- ❑ Advantages with respect to complete loop parallelization
  - ▶ **Simpler Finite State Machine**
    - Limits the area overhead
    - Limits the frequency reduction
  - ▶ **Aligned memory accesses**
    - Ad hoc memory allocation
  - ▶ **Vector functional units** can be obtained by
    - **sharing scalar** functional units
    - **replicating** scalar functional units
- ❑ Disadvantages
  - ▶ **Larger memory infrastructure**
  - ▶ **More complex functional units**



# Proposed flow

8





## ❑ Identification of **DoAll loops**

- ▶ Annotations (e.g., openMP pragmas)
- ▶ Compiler analyses
- ▶ Polyhedral analysis
- ▶ Other analyses

## ❑ Code transformation

- ▶ Removal of mutual exclusion code

```
#pragma omp simd
for(i=0; i<16; i++){
    sum=0;
    for(j=0; j<i; j++) {
        if(sum<10) {
            sum=sum+in[i][j];
        } else {
            sum=sum-in2[i][j];
        }
    }
    res[i] = sum/k;
}
```



# Proposed Flow: PreProcessing Transformation

10

- ❑ Removing of mutual exclusion code – potentially worsening performances
- ❑ Decomposition of complex operations
  - ▶ Vectorization can be applied selectively to the different parts of complex operations

```
#pragma omp simd
for(i=0; i<16; i++){
    sum=0;
    for(j=0; j<k; j++) {
        c = sum<10;
        if(c) {
            temp=in[i][j];
            sumS=sum+temp;
        }
        if(!c) {
            temp=in2[i][j];
            sumS=sum-temp;
        }
    }
    res[i] = sum/k;
}
```



# Proposed Flow: Instruction Classification

11

- ❑ **Vector instructions**
- ❑ **Multiscalar instructions**
  - ▶ Cannot implemented as vector or
  - ▶ Too large to be synthesized as vector instruction
- ❑ **DoAll loop instructions**
- ❑ **Nested loop instructions**

```
#pragma omp simd
for(i=0; i<16; i++) {
    sum=0;
    for(j=0; j<k; j++) {
        c = sum<10;
        if(c) {
            temp=in[i][j];
            sumS=sum+temp;
        }
        if(!c) {
            temp=in2[i][j];
            sumS=sum-temp;
        }
    }
    res[i] = sum/k;
}
```



# Proposed Flow: Instructions Transformation

12

## ❑ Vector instructions

- ▶ Transformed in a single vector instruction

## ❑ Multiscalar instructions

- ▶ Transformed in N scalar instructions

## ❑ DoAll loop instructions

- ▶ Increment is fixed

## ❑ Nested loop instructions

- ▶ Operands are fixed

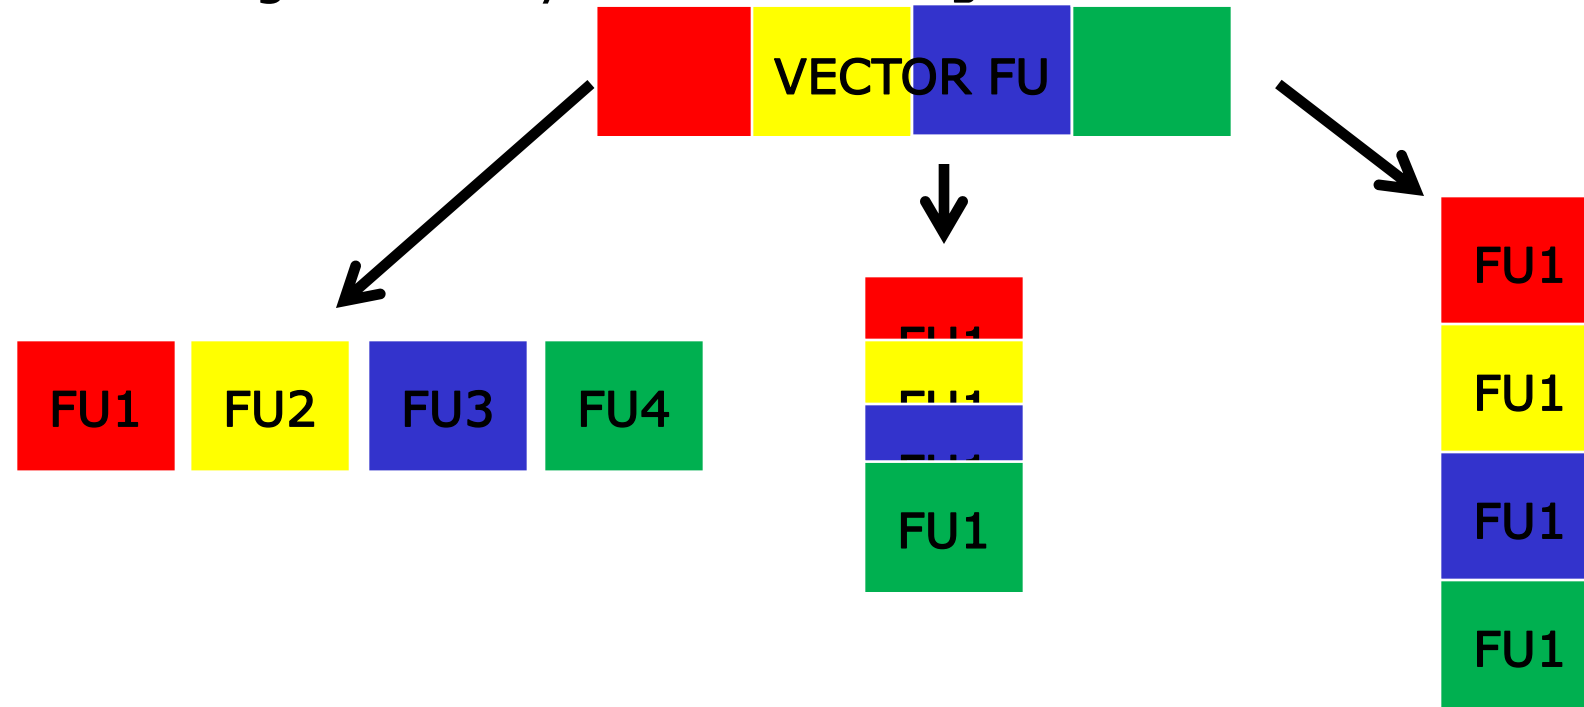
```
#pragma omp simd
for(i={0,1}; i<16; i+={2,2}) {
    sum={0,0};
    for(j=0; j<k; j++) {
        c = sum<{10,10};
        if(c[0] or c[1]) {
            temp[0]=in[i[0]][j];  (c[0])
            temp[1]=in[i[1]][j];  (c[1])
            sumS=sum+temp;  (c[0], c[1])
        }
        if(!(c[0] or c[1])) {
            temp[0]=in2[i[0]][j]; (!c[0])
            temp[1]=in2[i[0]][j]; (!c[1])
            sumS=sum-temp;  (!c[0],!c[1])
        }
    }
}
res[i[0]] = sum[0]/k;
res[i[1]] = sum[1]/k;
}
```



# Proposed Flow: Synthesis

13

- High Level Synthesis starting from transformed IR



Multiple Scalar Functional Units  
+ Very good performances  
+ Shared functional units  
- Larger area

Pipeline Functional Unit  
+ Good performances  
- More Complex design

Scalar Functional Unit  
+ Good area  
- Worst performances



- ❑ Generate an accelerator with vector size equal to one to see the effect of code transformation



```
bambu --compiler=I386_GCC49 --device-  
name=5SGXEA7N2F45C1 --simulate -fwhole-program -fno-  
delete-null-pointer-checks --clock-period=10 --experimental-  
setup=BAMBU-BALANCED-MP -fdisable-tree-cunroll -fdisable-  
tree-ivopts --param max-inline-insns-auto=1000 histogram.c  
-fopenmp-simd=1 --pretty-print=output.c
```



## Second example – Vector Size=4

16

- ❑ Generate an accelerator with vector size equal to 4 to evaluate the speed-up





```
bambu --compiler=I386_GCC49 --device-  
name=5SGXEA7N2F45C1 --simulate -fwhole-program -fno-  
delete-null-pointer-checks --clock-period=10 --experimental-  
setup=BAMBU-BALANCED-MP -fdisable-tree-cunroll -fdisable-  
tree-ivopts --param max-inline-insns-auto=1000 histogram.c -  
fopenmp-simd=4
```



## Third Example – Different Vector Size

18

- Generate accelerators with vector size equal to 2,3,4,8 to evaluate the speedup



```
bambu --compiler=I386_GCC49 --device-  
name=5SGXEA7N2F45C1 --simulate -fwhole-program -fno-  
delete-null-pointer-checks --clock-period=10 --experimental-  
setup=BAMBU-BALANCED-MP -fdisable-tree-cunroll -fdisable-  
tree-ivopts --param max-inline-insns-auto=1000 histogram.c  
-fopenmp-simd=2 --pretty-print=output.c
```

Etc.




- ❑ Parallel benchmarks for High Level Synthesis annotated with `#pragma omp simd`
- ❑ Parallel degrees considered for vectorization: 1 2 4 8
- ❑ Target devices:
  - ▶ Xilinx Zynq-7000 xc7z020
  - ▶ Altera Cyclone II EP2C70F896C6
- ❑ Results:
  - ▶ Max speedup: 7.35x
  - ▶ Max reduction of area-delay product: 40%



# Experimental Evaluation: Case Study : Add

21



- ❑ Vectorization of Outer Loops can be integrated in High Level Synthesis Flow
  - ▶ Synthesis of vector functional unit is required
- ❑ Memory accesses are the bottleneck which prevents obtaining maximum speedup
  - ▶ Different memory allocation is required

